

Final Project Submission

Please fill out:

- Student name: Heath Rittler
- Student pace: Self paced
- Scheduled project review date/time: 10/28/2022
- Instructor name: Mark Barbour
- Blog post URL: <https://medium.com/@heathlikethecandybar/>

Introduction

Business Case/ Summary

My business stakeholder is the owner of a local real estate agency. The agency focuses on home improvement recommendations to homeowners that will optimize the buying and selling of their homes. I will be deciphering which home improvements lead to the most value of a home. The output of my analysis will show 3 features that will impact the value of a home. Each one of these features should be in control of the homeowner, meaning they will be able to make those changes if they so desire before listing their home.

Core Field Names and Definitions from Data Source

The dataset comes from King County, in Washington state. The data in its raw form consists of 21 columns, and 21,597 records before any cleaning or feature engineering. More information on the columns and definitions can be found in the repository here (data/column_names.md). Any further information regarding the dataset can be found at the King County website (<https://info.kingcounty.gov/assessor/esales/Glossary.aspx?type=r>)

- id - Unique identifier for a house
- date - Date house was sold
- price - Sale price (prediction target)
- bedrooms - Number of bedrooms
- bathrooms - Number of bathrooms
- sqft_living - Square footage of living space in the home

- `sqft_lot` - Square footage of the lot
- `floors` - Number of floors (levels) in house
- `waterfront` - Whether the house is on a waterfront
 - Includes Duwamish, Elliott Bay, Puget Sound, Lake Union, Ship Canal, Lake Washington, Lake Sammamish, other lake, and river/slough waterfronts
- `view` - Quality of view from house
 - Includes views of Mt. Rainier, Olympics, Cascades, Territorial, Seattle Skyline, Puget Sound, Lake Washington, Lake Sammamish, small lake / river / creek, and other
- `condition` - How good the overall condition of the house is. Related to maintenance of house.
 - Additional details below
- `grade` - Overall grade of the house. Related to the construction and design of the house.
 - See the [King County Assessor Website](#) for further explanation of each building grade code
- `sqft_above` - Square footage of house apart from basement
- `sqft_basement` - Square footage of the basement
- `yr_built` - Year when house was built
- `yr_renovated` - Year when house was renovated
- `zipcode` - ZIP Code used by the United States Postal Service
- `lat` - Latitude coordinate
- `long` - Longitude coordinate
- `sqft_living15` - The square footage of interior housing living space for the nearest 15 neighbors
- `sqft_lot15` - The square footage of the land lots of the nearest 15 neighbors

Additional Details - Building Condition Definitions

Relative to age and grade. Coded 1-5.

- 1 = Poor-Worn out Repair and overhaul needed on painted surfaces, roofing, plumbing, heating and numerous functional inadequacies. Excessive deferred maintenance and abuse, limited value-in-use, approaching abandonment or major reconstruction; reuse or change in occupancy is imminent. Effective age is near the end of the scale regardless of the actual chronological age.
- 2 = Fair-Badly worn Much repair needed. Many items need refinishing or overhauling, deferred maintenance obvious, inadequate building utility and systems all shortening the life expectancy and increasing the effective age.

- 3 = Average Some evidence of deferred maintenance and normal obsolescence with age in that a few minor repairs are needed, along with some refinishing. All major components still functional and contributing toward an extended life expectancy. Effective age and utility is standard for like properties of its class and usage.
- 4 = Good No obvious maintenance required but neither is everything new. Appearance and utility are above the standard and the overall effective age will be lower than the typical property.
- 5 = Very Good All items well maintained, many having been overhauled and repaired as they have shown signs of wear, increasing the life expectancy and lowering the effective age with little deterioration or obsolescence evident with a high degree of utility.

Residential Building Grades

- Grades 1 – 3 Falls short of minimum building standards. Normally cabin or inferior structure.
- Grade 4 Generally older low quality construction. Does not meet code.
- Grade 5 Lower construction costs and workmanship. Small, simple design.
- Grade 6 Lowest grade currently meeting building codes. Low quality materials, simple designs.
- Grade 7 Average grade of construction and design. Commonly seen in plats and older subdivisions.
- Grade 8 Just above average in construction and design. Usually better materials in both the exterior and interior finishes.
- Grade 9 Better architectural design, with extra exterior and interior design and quality.
- Grade 10 Homes of this quality generally have high quality features. Finish work is better, and more design quality is seen in the floor plans and larger square footage.
- Grade 11 Custom design and higher quality finish work, with added amenities of solid woods, bathroom fixtures and more luxurious options.
- Grade 12 Custom design and excellent builders. All materials are of the highest quality and all conveniences are present.
- Grade 13 Generally custom designed and built. Approaching the Mansion level. Large amount of highest quality cabinet work, wood trim and marble; large entries.

Data Load, Cleaning, & Exploratory

Package load

```
In [1]: # Importing packages for analysis
```

```
import pandas as pd
import numpy as np
from sklearn.preprocessing import LabelEncoder
from matplotlib import pyplot as plt
import matplotlib.ticker as mtick
import seaborn as sns
import statsmodels.api as sm
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
lr = LinearRegression()
plt.style.use('seaborn-talk')
import sklearn.metrics as metrics

%matplotlib inline
```

In [2]: # Choosing standard colors for project

```
pal = ['#003f5c',
       '#374c80',
       '#7a5195',
       '#bc5090',
       '#ef5675',
       '#ff764a',
       '#ffa600']

color_codes = ['darkblue', 'indigo', 'purple', 'pink', 'peach', 'orange', 'yellow']

c = lambda x: pal[color_codes.index(x)]

sns.palplot(sns.color_palette(pal))
```



Ooo that's pretty. Let's take a look now at our data -- coming from a csv.

Load

In [3]: # Here we go! Read csv, look at initial shape

```
df = pd.read_csv('data/kc_house_data.csv')
```

```
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21597 entries, 0 to 21596
Data columns (total 21 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   id                21597 non-null   int64  
 1   date              21597 non-null   object  
 2   price              21597 non-null   float64 
 3   bedrooms           21597 non-null   int64  
 4   bathrooms          21597 non-null   float64 
 5   sqft_living        21597 non-null   int64  
 6   sqft_lot            21597 non-null   int64  
 7   floors              21597 non-null   float64 
 8   waterfront          19221 non-null   object  
 9   view               21534 non-null   object  
 10  condition           21597 non-null   object  
 11  grade               21597 non-null   object  
 12  sqft_above          21597 non-null   int64  
 13  sqft_basement       21597 non-null   object  
 14  yr_built            21597 non-null   int64  
 15  yr_renovated        17755 non-null   float64 
 16  zipcode             21597 non-null   int64  
 17  lat                 21597 non-null   float64 
 18  long                21597 non-null   float64 
 19  sqft_living15       21597 non-null   int64  
 20  sqft_lot15          21597 non-null   int64  
dtypes: float64(6), int64(9), object(6)
memory usage: 3.5+ MB
```

Because my stakeholder and business proposition is recommending enhancements for homeowners, I am going to remove the features that are not impactable by the homeowner. (i.e. Homeowner can't change the fact that they are or are not on waterfront)

```
In [4]: # Removing waterfront, view, lat, long, zipcode from dataframe
```

```
df_clean = df.drop(['waterfront', 'view', 'lat', 'long', 'zipcode'], axis=1)
```

```
df_clean.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21597 entries, 0 to 21596
Data columns (total 16 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   id                21597 non-null   int64
```

```
1   date          21597 non-null  object
2   price         21597 non-null  float64
3   bedrooms      21597 non-null  int64
4   bathrooms     21597 non-null  float64
5   sqft_living   21597 non-null  int64
6   sqft_lot      21597 non-null  int64
7   floors        21597 non-null  float64
8   condition     21597 non-null  object
9   grade         21597 non-null  object
10  sqft_above    21597 non-null  int64
11  sqft_basement 21597 non-null  object
12  yr_built     21597 non-null  int64
13  yr_renovated 17755 non-null  float64
14  sqft_living15 21597 non-null  int64
15  sqft_lot15   21597 non-null  int64
dtypes: float64(4), int64(8), object(4)
memory usage: 2.6+ MB
```

```
In [5]: # Digging into object/ string fields to understand how we will transform. First
# suspect is sqft_basement, would assume numeric values and 0 for no basement.

df_clean['sqft_basement'].value_counts()
```

```
Out[5]: 0.0      12826
?          454
600.0     217
500.0     209
700.0     208
...
784.0      1
792.0      1
2810.0     1
556.0      1
225.0      1
Name: sqft_basement, Length: 304, dtype: int64
```

```
In [6]: # Changing basement sqft column. Adjusting ? to 0.0, and changing data type
# to numeric in order to feed model.

df_clean['sqft_basement'] = df_clean['sqft_basement'].replace(['?'], '0.0')

df_clean['sqft_basement'].astype(float)

df_clean['sqft_basement'].value_counts();
```

```
In [7]: # Digging into yr_renovated
```

```
df_clean['yr_renovated'].value_counts();
```

```
In [8]: # Not a ton of values, so going to add a boolean feature = is_renovated
```

```
conditions = [df_clean.loc[:, 'yr_renovated'] > 0,
              df_clean.loc[:, 'yr_renovated'] == 0]

values = [1, 0]

df_clean.loc[:, 'is_renovated'] = np.select(conditions, values, default=0)

df_clean['is_renovated'].value_counts();
```

```
In [9]: # Adding another column with the count of years between yr_built and renovation 'yr_frm_btr'
```

```
conditions = [
    (df_clean['yr_renovated'] == 0),
    (df_clean['yr_renovated'] > 0)
]

# Create a list of the values we want to assign for each condition
values = [0, (df_clean['yr_renovated'] - df_clean['yr_built'])]

# Create a new column and use np.select to assign values to it using our lists as arguments
df['yr_frm_btr'] = np.select(conditions, values);
```

```
In [10]: # Dropping the yr_renovated column. Don't necessarily need it right now
```

```
df_clean = df_clean.drop('yr_renovated', axis = 1)
```

```
In [11]: # Another look again at info
```

```
df_clean.info();
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21597 entries, 0 to 21596
Data columns (total 16 columns):
 #   Column            Non-Null Count  Dtype  
 ---  -- 
 0   id                21597 non-null   int64  
 1   date               21597 non-null   object  
 2   price              21597 non-null   float64 
 3   bedrooms           21597 non-null   int64  
 4   bathrooms          21597 non-null   float64 
 5   sqft_living        21597 non-null   int64  
 6   sqft_lot            21597 non-null   int64  
 7   floors              21597 non-null   float64 
 8   condition           21597 non-null   object  
 9   grade               21597 non-null   float64 
 10  roof_style          21597 non-null   object  
 11  exterior_color      21597 non-null   object  
 12  exterior_condition  21597 non-null   object  
 13  year_built          21597 non-null   int64  
 14  year_renovated      21597 non-null   int64  
 15  yr_frm_btr          21597 non-null   float64
```

```
6    sqft_lot      21597 non-null  int64
7    floors        21597 non-null  float64
8    condition     21597 non-null  object
9    grade         21597 non-null  object
10   sqft_above    21597 non-null  int64
11   sqft_basement 21597 non-null  object
12   yr_built      21597 non-null  int64
13   sqft_living15 21597 non-null  int64
14   sqft_lot15    21597 non-null  int64
15   is_renovated  21597 non-null  int64
dtypes: float64(3), int64(9), object(4)
memory usage: 2.6+ MB
```

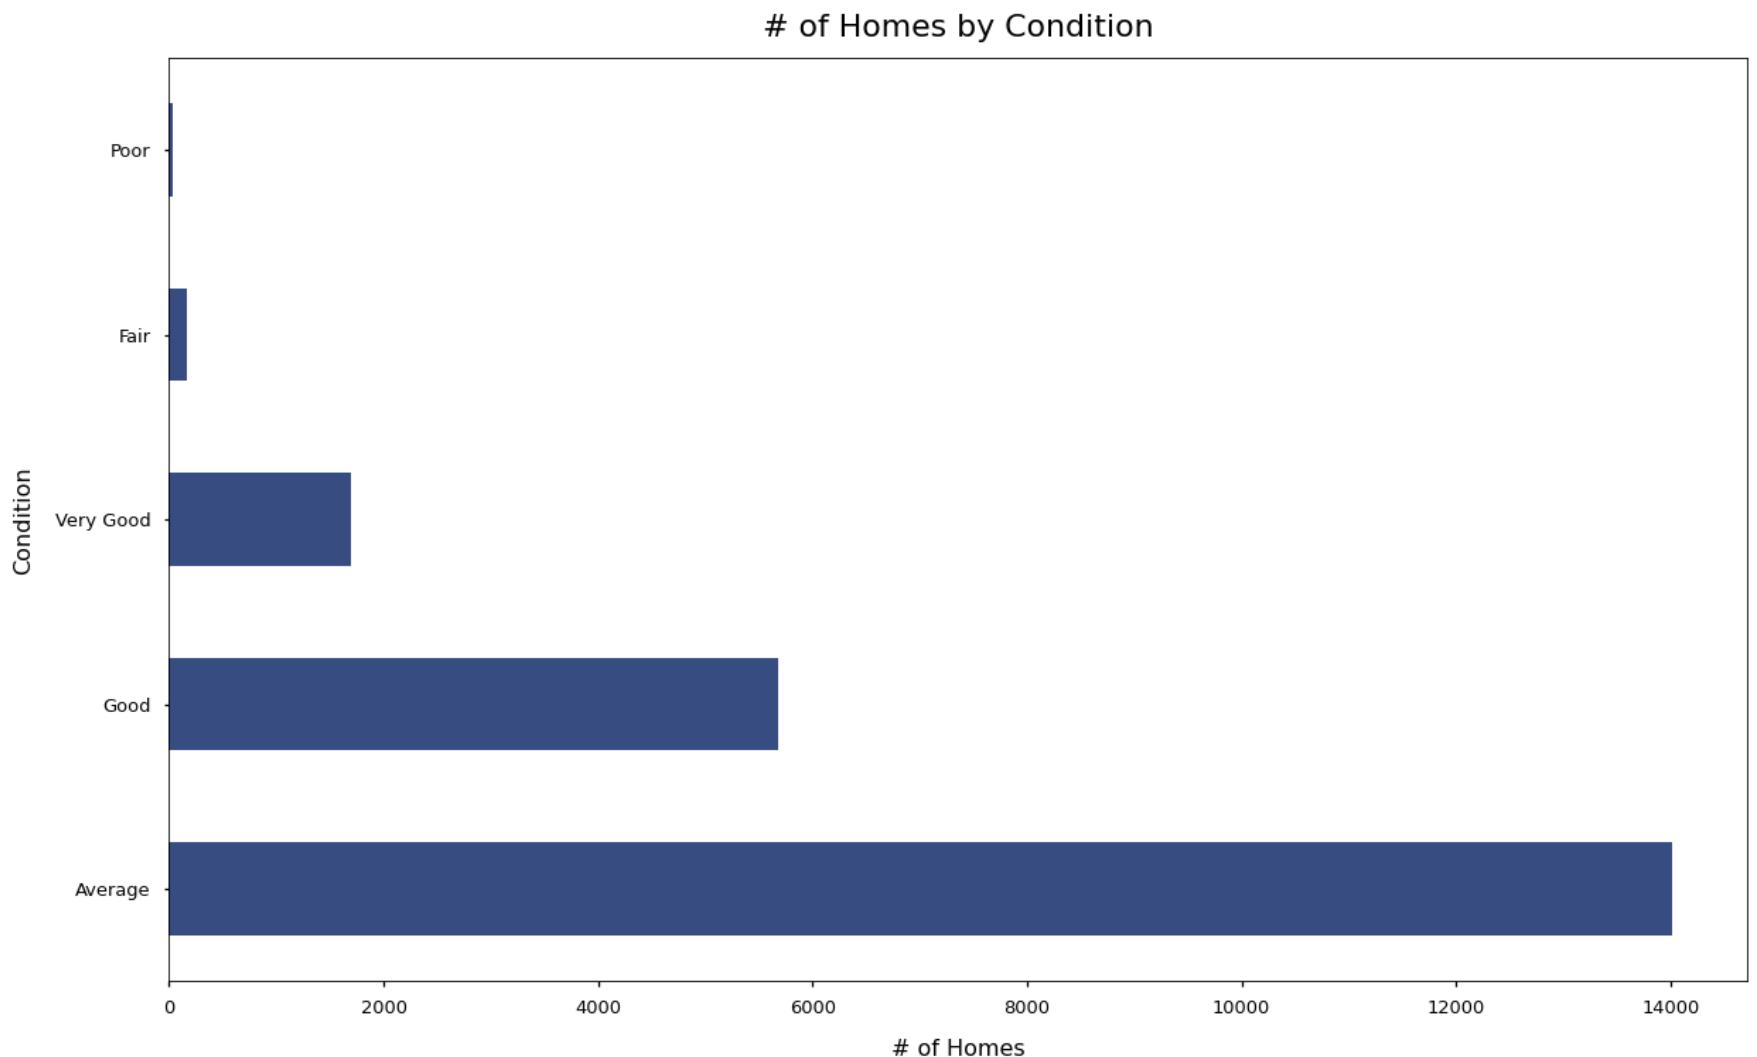
```
In [12]: # Checking out condition; will need to one hot encode these columns

fig, ax = plt.subplots(figsize = (20,12))

df_clean['condition'].value_counts().plot(kind='barh', color=c("indigo"))

ax.set_title('# of Homes by Condition', pad = 15, fontsize = 22)
ax.set_xlabel('# of Homes', labelpad = 15, fontsize = 16)
ax.set_ylabel('Condition', labelpad = 15, fontsize = 16)

plt.show();
```



This breakdown is a little disappointing. There really isn't a great distribution of condition. Everything is listed as good, and above. Going to look at grade to see if that has a better distribution before encoding.

```
In [13]: # Checking out grade; going to strip string so we can make column numeric.  
df_clean['grade'].value_counts()
```

```
Out[13]: 7 Average      8974  
8 Good        6065  
9 Better       2615  
6 Low Average  2038
```

```
10 Very Good      1134
11 Excellent     399
5 Fair           242
12 Luxury         89
4 Low            27
13 Mansion        13
3 Poor            1
Name: grade, dtype: int64
```

```
In [14]: # Split column and add new columns to df

gd_sp = df_clean['grade'].str.split(' ', n = 1, expand = True)

# Add column names
gd_sp.columns = ['grade_num', "grade_cat"]

# Concat back to df_clean dataframe
df_clean = pd.concat([df_clean, gd_sp], axis = 1)

# Update grade_num column to int datatype
df_clean['grade_num'] = df_clean['grade_num'].astype(int)
```

```
In [15]: # Dropping original column from dataframe

df_clean = df_clean.drop(['grade'], axis = 1)

# Checking df_clean to see where we are at.
df_clean.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21597 entries, 0 to 21596
Data columns (total 17 columns):
 #   Column           Non-Null Count  Dtype  
 ---  -- 
 0   id               21597 non-null   int64  
 1   date              21597 non-null   object 
 2   price             21597 non-null   float64 
 3   bedrooms          21597 non-null   int64  
 4   bathrooms          21597 non-null   float64 
 5   sqft_living       21597 non-null   int64  
 6   sqft_lot           21597 non-null   int64  
 7   floors             21597 non-null   float64 
 8   condition          21597 non-null   object 
 9   sqft_above          21597 non-null   int64  
 10  sqft_basement      21597 non-null   object 
 11  yr_built           21597 non-null   int64  
 12  sqft_living15      21597 non-null   int64
```

```
13    sqft_lot15      21597 non-null  int64
14    is_renovated    21597 non-null  int64
15    grade_num        21597 non-null  int64
16    grade_cat        21597 non-null  object
dtypes: float64(3), int64(10), object(4)
memory usage: 2.8+ MB
```

```
In [16]: # Adding ratio of above ground square footage to square footage of living area.
# Trying to understand if there is more living square footage below ground,
# is that attractive to a home buyer.

df_clean.loc[:, 'sqft_a/l'] = (
    df_clean['sqft_above'] / df_clean['sqft_living']
)

df_clean.head()
```

```
Out[16]:
```

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	condition	sqft_above	sqft_basement	yr_bu
0	7129300520	10/13/2014	2219000.0	3	1.00	1180	5650	1.0	Average	1180	0.0	19
1	6414100192	12/9/2014	5380000.0	3	2.25	2570	7242	2.0	Average	2170	400.0	19
2	5631500400	2/25/2015	1800000.0	2	1.00	770	10000	1.0	Average	770	0.0	19
3	2487200875	12/9/2014	6040000.0	4	3.00	1960	5000	1.0	Very Good	1050	910.0	19
4	1954400510	2/18/2015	5100000.0	3	2.00	1680	8080	1.0	Average	1680	0.0	19

```
In [17]: # Building ratio for living square footage to # of bedrooms.

df_clean.loc[:, 'sqft_l/b'] = (
    round(df_clean['sqft_living'] / df_clean['bedrooms'], 2)
)

df_clean.head()
```

```
Out[17]:
```

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	condition	sqft_above	sqft_basement	yr_bu
0	7129300520	10/13/2014	2219000.0	3	1.00	1180	5650	1.0	Average	1180	0.0	19
1	6414100192	12/9/2014	5380000.0	3	2.25	2570	7242	2.0	Average	2170	400.0	19
2	5631500400	2/25/2015	1800000.0	2	1.00	770	10000	1.0	Average	770	0.0	19

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	condition	sqft_above	sqft_basement	yr_buil
3	2487200875	12/9/2014	604000.0	4	3.00	1960	5000	1.0	Very Good	1050	910.0	19
4	1954400510	2/18/2015	510000.0	3	2.00	1680	8080	1.0	Average	1680	0.0	19

```
In [18]: # Look for missing data/ na's

df_clean.isna().any()
```

```
Out[18]: id      False
date     False
price    False
bedrooms False
bathrooms False
sqft_living False
sqft_lot  False
floors   False
condition False
sqft_above False
sqft_basement False
yr_built  False
sqft_living15 False
sqft_lot15  False
is_renovated False
grade_num  False
grade_cat  False
sqft_a/l   False
sqft_l/b   False
dtype: bool
```

```
In [19]: # Taking another peek to see what is left to clean or transform

df_clean.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21597 entries, 0 to 21596
Data columns (total 19 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   id                21597 non-null   int64  
 1   date              21597 non-null   object 
 2   price             21597 non-null   float64
 3   bedrooms          21597 non-null   int64  
 4   bathrooms         21597 non-null   float64
 5   sqft_living       21597 non-null   int64
```

```
6    sqft_lot          21597 non-null  int64
7    floors            21597 non-null  float64
8    condition         21597 non-null  object
9    sqft_above        21597 non-null  int64
10   sqft_basement     21597 non-null  object
11   yr_built          21597 non-null  int64
12   sqft_living15    21597 non-null  int64
13   sqft_lot15        21597 non-null  int64
14   is_renovated     21597 non-null  int64
15   grade_num         21597 non-null  int64
16   grade_cat          21597 non-null  object
17   sqft_a/1           21597 non-null  float64
18   sqft_l/b           21597 non-null  float64
dtypes: float64(5), int64(10), object(4)
memory usage: 3.1+ MB
```

In [20]: `# OHE categoricals final categoricals`

```
categoricals = ['condition', 'grade_cat']

dummies = pd.get_dummies(df_clean[categoricals],
                         prefix=categoricals,
                         drop_first=True)

df_preprocessed = df_clean.drop(categoricals, axis=1)

df_clean = pd.concat([df_preprocessed, dummies], axis=1)

df_clean.head()
```

Out[20]:

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	sqft_above	sqft_basement	...	grade_cat_B
0	7129300520	10/13/2014	221900.0	3	1.00	1180	5650	1.0	1180		0.0	...
1	6414100192	12/9/2014	538000.0	3	2.25	2570	7242	2.0	2170		400.0	...
2	5631500400	2/25/2015	180000.0	2	1.00	770	10000	1.0	770		0.0	...
3	2487200875	12/9/2014	604000.0	4	3.00	1960	5000	1.0	1050		910.0	...
4	1954400510	2/18/2015	510000.0	3	2.00	1680	8080	1.0	1680		0.0	...

5 rows × 31 columns

In [21]: `# Cleaning final sqft_basement column, and dropping the date/ id column for now`

```
df_clean = df_clean.drop('date', axis=1)
df_clean = df_clean.drop('id', axis=1)
df_clean['sqft_basement'] = df_clean['sqft_basement'].astype(float)
df_clean.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21597 entries, 0 to 21596
Data columns (total 29 columns):
 #   Column            Non-Null Count  Dtype  
---  -- 
 0   price              21597 non-null   float64
 1   bedrooms           21597 non-null   int64  
 2   bathrooms          21597 non-null   float64
 3   sqft_living        21597 non-null   int64  
 4   sqft_lot            21597 non-null   int64  
 5   floors              21597 non-null   float64
 6   sqft_above          21597 non-null   int64  
 7   sqft_basement       21597 non-null   float64
 8   yr_built            21597 non-null   int64  
 9   sqft_living15       21597 non-null   int64  
 10  sqft_lot15          21597 non-null   int64  
 11  is_renovated        21597 non-null   int64  
 12  grade_num           21597 non-null   int64  
 13  sqft_a/l            21597 non-null   float64
 14  sqft_l/b            21597 non-null   float64
 15  condition_Fair      21597 non-null   uint8  
 16  condition_Good       21597 non-null   uint8  
 17  condition_Poor       21597 non-null   uint8  
 18  condition_Very Good 21597 non-null   uint8  
 19  grade_cat_Better     21597 non-null   uint8  
 20  grade_cat_Excellent  21597 non-null   uint8  
 21  grade_cat_Fair       21597 non-null   uint8  
 22  grade_cat_Good       21597 non-null   uint8  
 23  grade_cat_Low         21597 non-null   uint8  
 24  grade_cat_Low Average 21597 non-null   uint8  
 25  grade_cat_Luxury     21597 non-null   uint8  
 26  grade_cat_Mansion    21597 non-null   uint8  
 27  grade_cat_Poor       21597 non-null   uint8  
 28  grade_cat_Very Good  21597 non-null   uint8  
dtypes: float64(6), int64(9), uint8(14)
memory usage: 2.8 MB
```

In [22]: *# Early correlation matrix to understand relationships with what we currently have.*

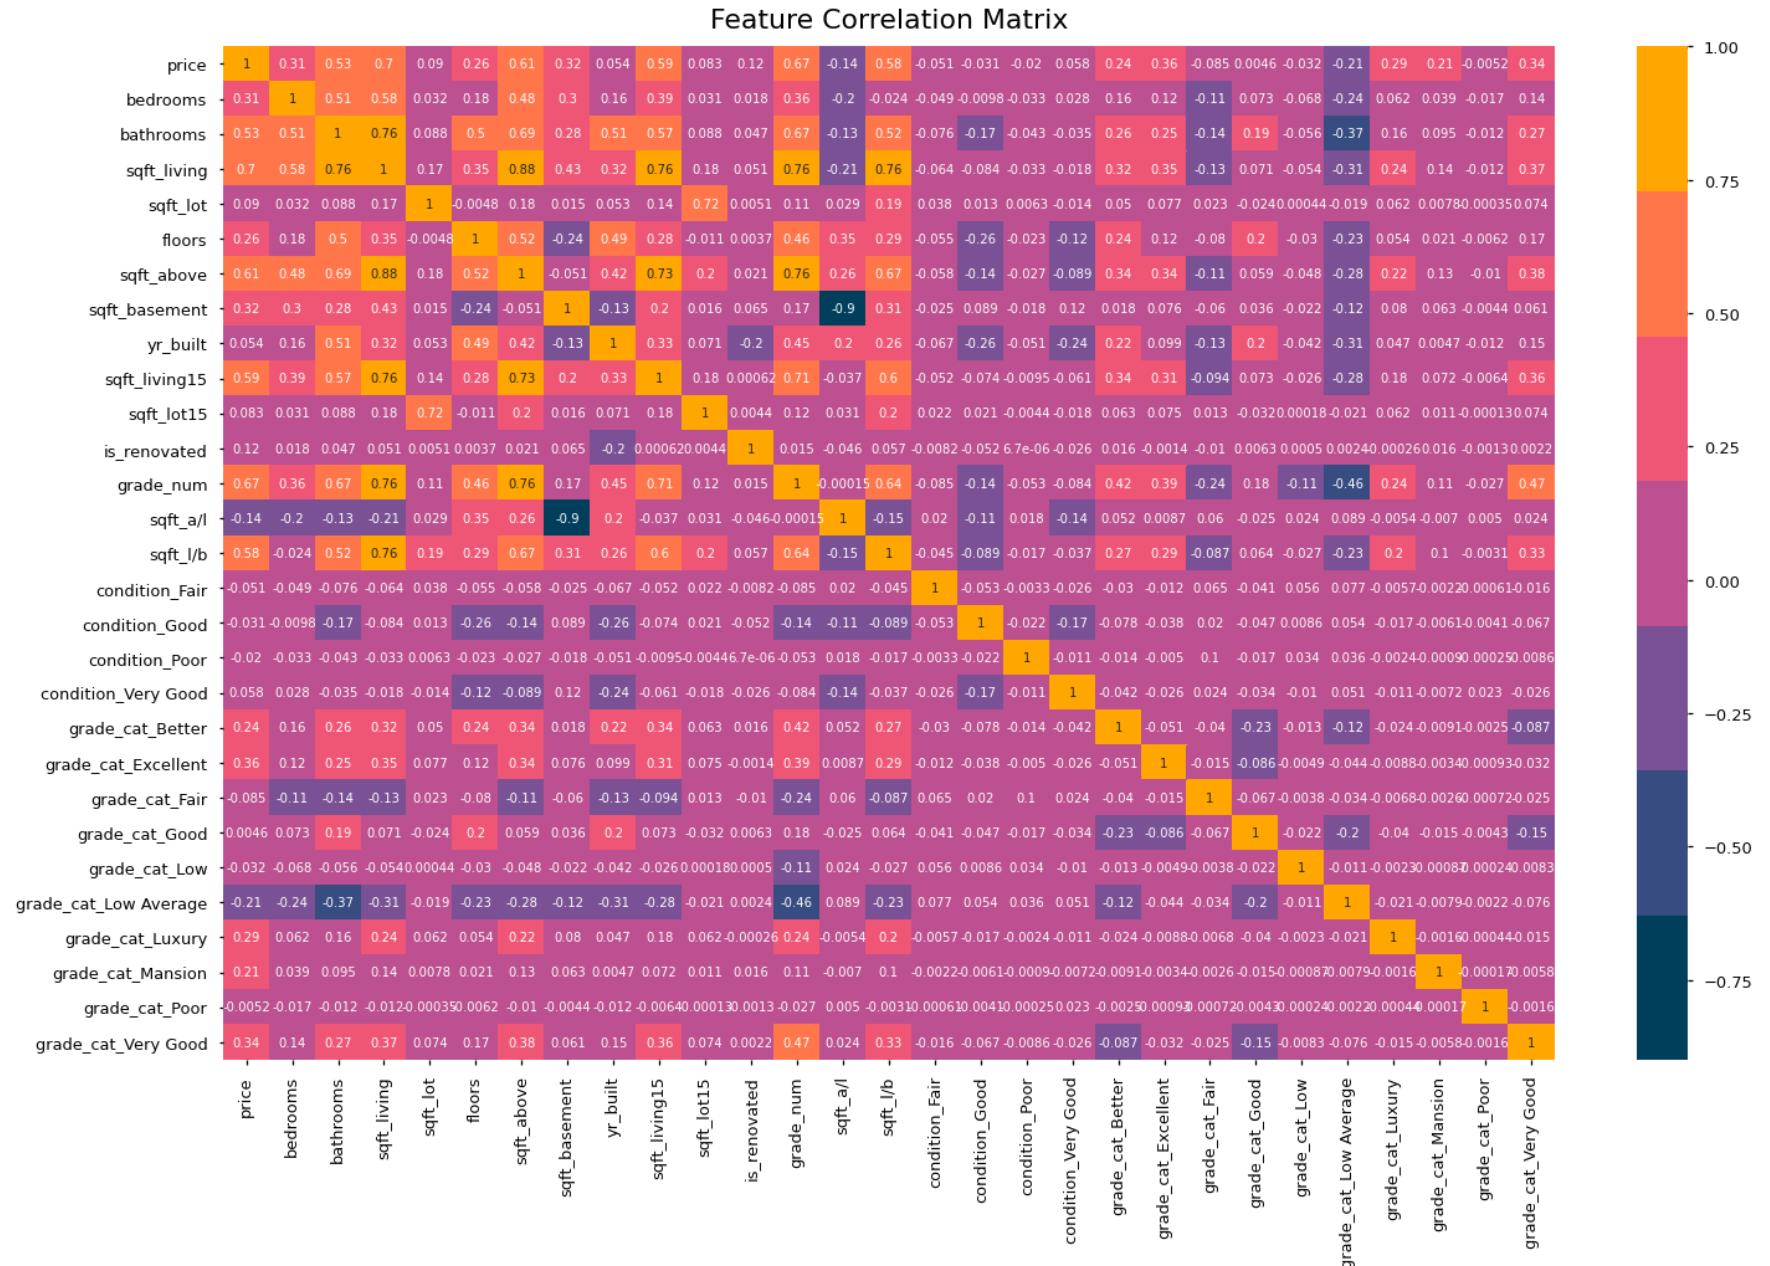
```
fig, ax = plt.subplots(figsize=(24,15))
corrM = df_clean.corr()

ax.set_title('Feature Correlation Matrix', pad=15, fontsize=22)
```

```

sns.heatmap(corrM, annot=True, cmap=pa1)
plt.show()

```



Looks like bathrooms, sqft_living, sqft_above, sqft_lot15, and condition fair, all have decent (>.50) r² with our dependent variable

price. We will see later if the p-value of these independent variables are strong enough for us to continue to evaluate these variables as having the ability to predict price. Unfortunately, we cannot do anything about the sqft_lot15 independent variable, so we will most likely be dropping this feature.

Since we have grade_num, I am going to remove the grade_cat breakouts. Doesn't seem like they are helping much anyways. We can always add them back in, but fewer features may make more sense going into our baseline.

```
In [23]: # Removing the grade category OHE features.
```

```
df_clean = df_clean.drop(df_clean.iloc[:,15:], axis=1)  
df_clean
```

```
Out[23]:
```

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	sqft_above	sqft_basement	yr_built	sqft_living15	sqft_lot15	i
0	221900.0	3	1.00	1180	5650	1.0	1180	0.0	1955	1340	5650	
1	538000.0	3	2.25	2570	7242	2.0	2170	400.0	1951	1690	7639	
2	180000.0	2	1.00	770	10000	1.0	770	0.0	1933	2720	8062	
3	604000.0	4	3.00	1960	5000	1.0	1050	910.0	1965	1360	5000	
4	510000.0	3	2.00	1680	8080	1.0	1680	0.0	1987	1800	7503	
...	
21592	360000.0	3	2.50	1530	1131	3.0	1530	0.0	2009	1530	1509	
21593	400000.0	4	2.50	2310	5813	2.0	2310	0.0	2014	1830	7200	
21594	402101.0	2	0.75	1020	1350	2.0	1020	0.0	2009	1020	2007	
21595	400000.0	3	2.50	1600	2388	2.0	1600	0.0	2004	1410	1287	
21596	325000.0	2	0.75	1020	1076	2.0	1020	0.0	2008	1020	1357	

21597 rows × 15 columns

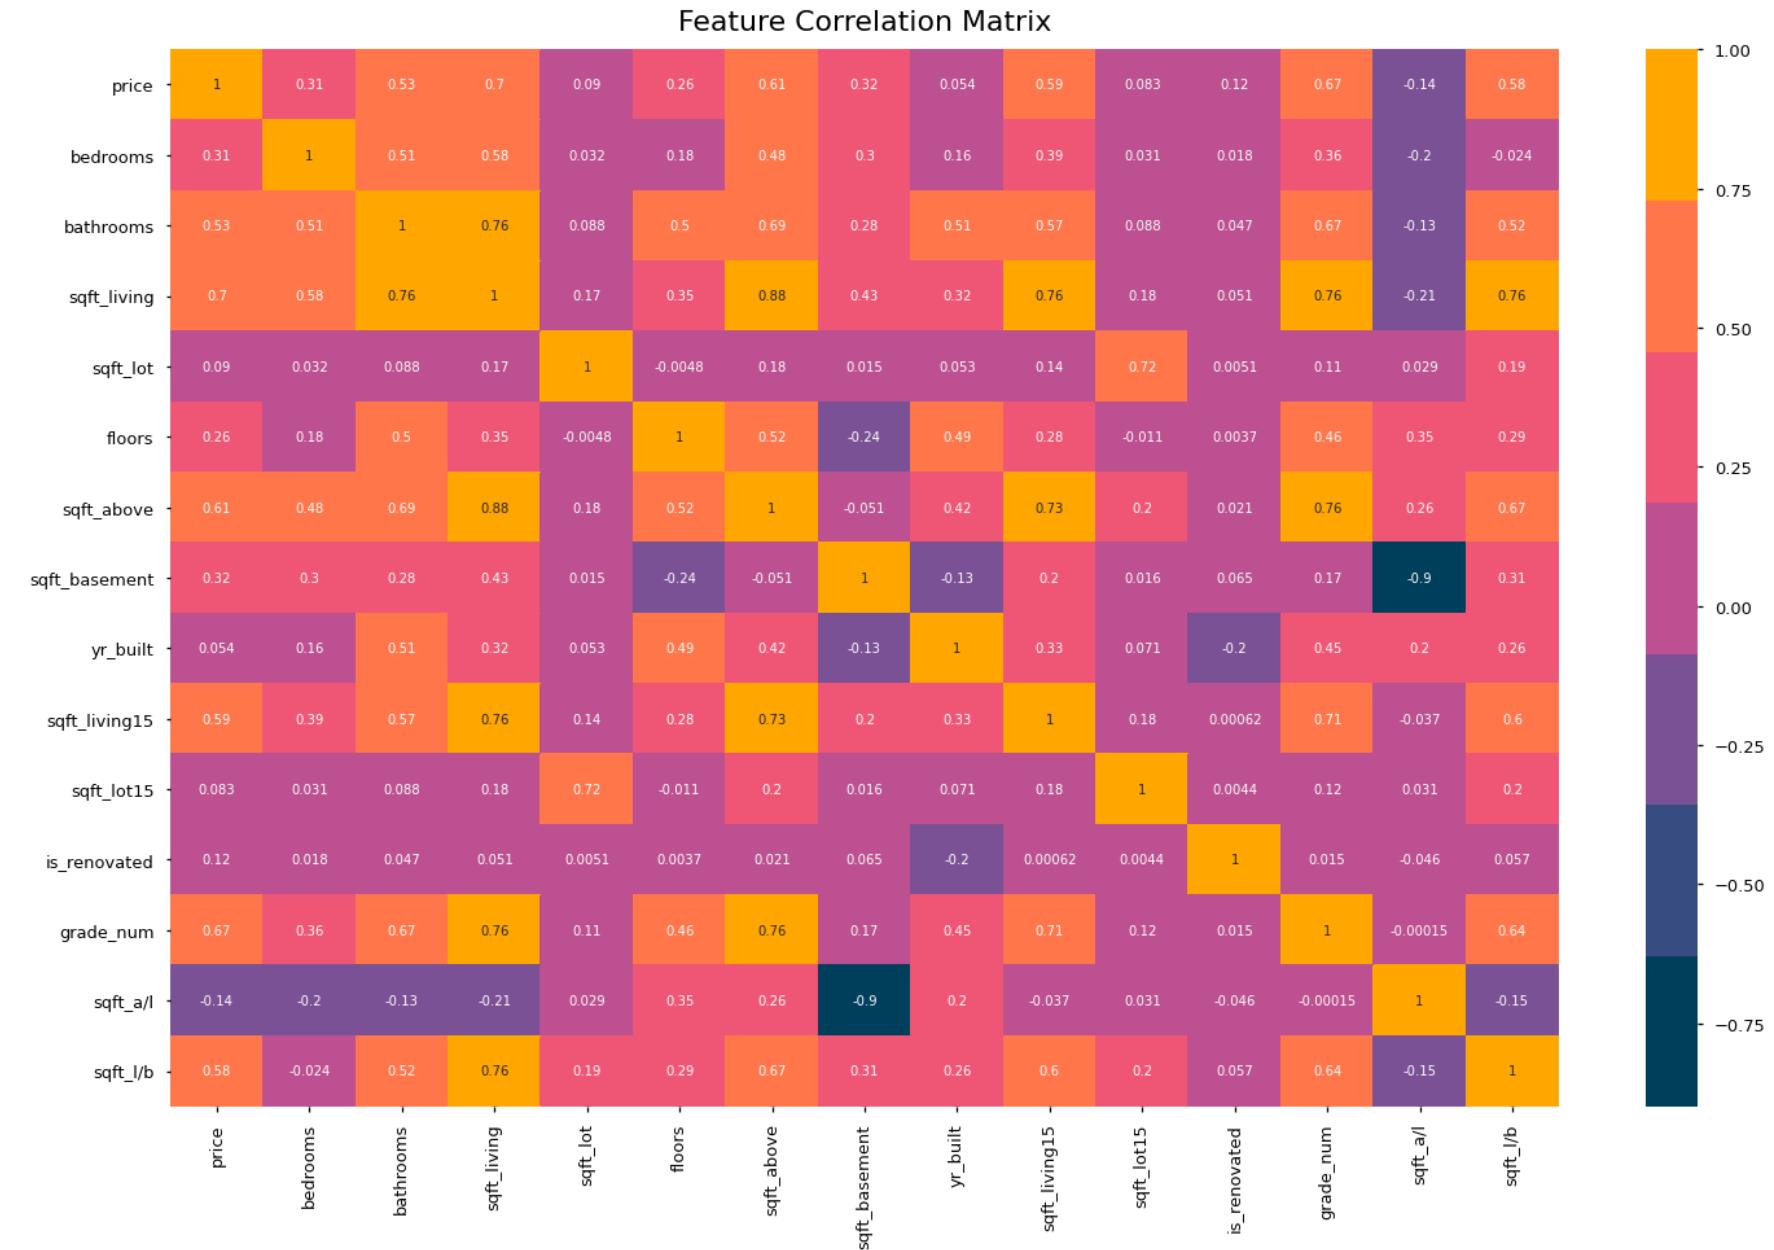
```
In [24]: # Early correlation matrix to understand relationships with what we currently have.
```

```
fig, ax = plt.subplots(figsize=(24,15))  
corrM = df_clean.corr()  
  
ax.set_title('Feature Correlation Matrix', pad=15, fontsize=22)
```

```

sns.heatmap(corrM, annot=True, cmap=pa1)
plt.show()

```



Exploratory Data Analysis

I am going to first focus on 3 variables that have traditionally been included when determining the price of a house, the number of bedrooms, the number of bathrooms, and the square footage of the living area. I am going to dig into these individually, and then look at the broader data before running any regressions.

```
In [25]: # Look at initial distributions of columns  
df_clean.describe()
```

```
Out[25]:
```

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	sqft_above	sqft_basement	yr_built
count	2.159700e+04	21597.000000	21597.000000	21597.000000	2.159700e+04	21597.000000	21597.000000	21597.000000	21597.00
mean	5.402966e+05	3.373200	2.115826	2080.321850	1.509941e+04	1.494096	1788.596842	285.716581	1970.99
std	3.673681e+05	0.926299	0.768984	918.106125	4.141264e+04	0.539683	827.759761	439.819830	29.37
min	7.800000e+04	1.000000	0.500000	370.000000	5.200000e+02	1.000000	370.000000	0.000000	1900.00
25%	3.220000e+05	3.000000	1.750000	1430.000000	5.040000e+03	1.000000	1190.000000	0.000000	1951.00
50%	4.500000e+05	3.000000	2.250000	1910.000000	7.618000e+03	1.500000	1560.000000	0.000000	1975.00
75%	6.450000e+05	4.000000	2.500000	2550.000000	1.068500e+04	2.000000	2210.000000	550.000000	1997.00
max	7.700000e+06	33.000000	8.000000	13540.000000	1.651359e+06	3.500000	9410.000000	4820.000000	2015.00

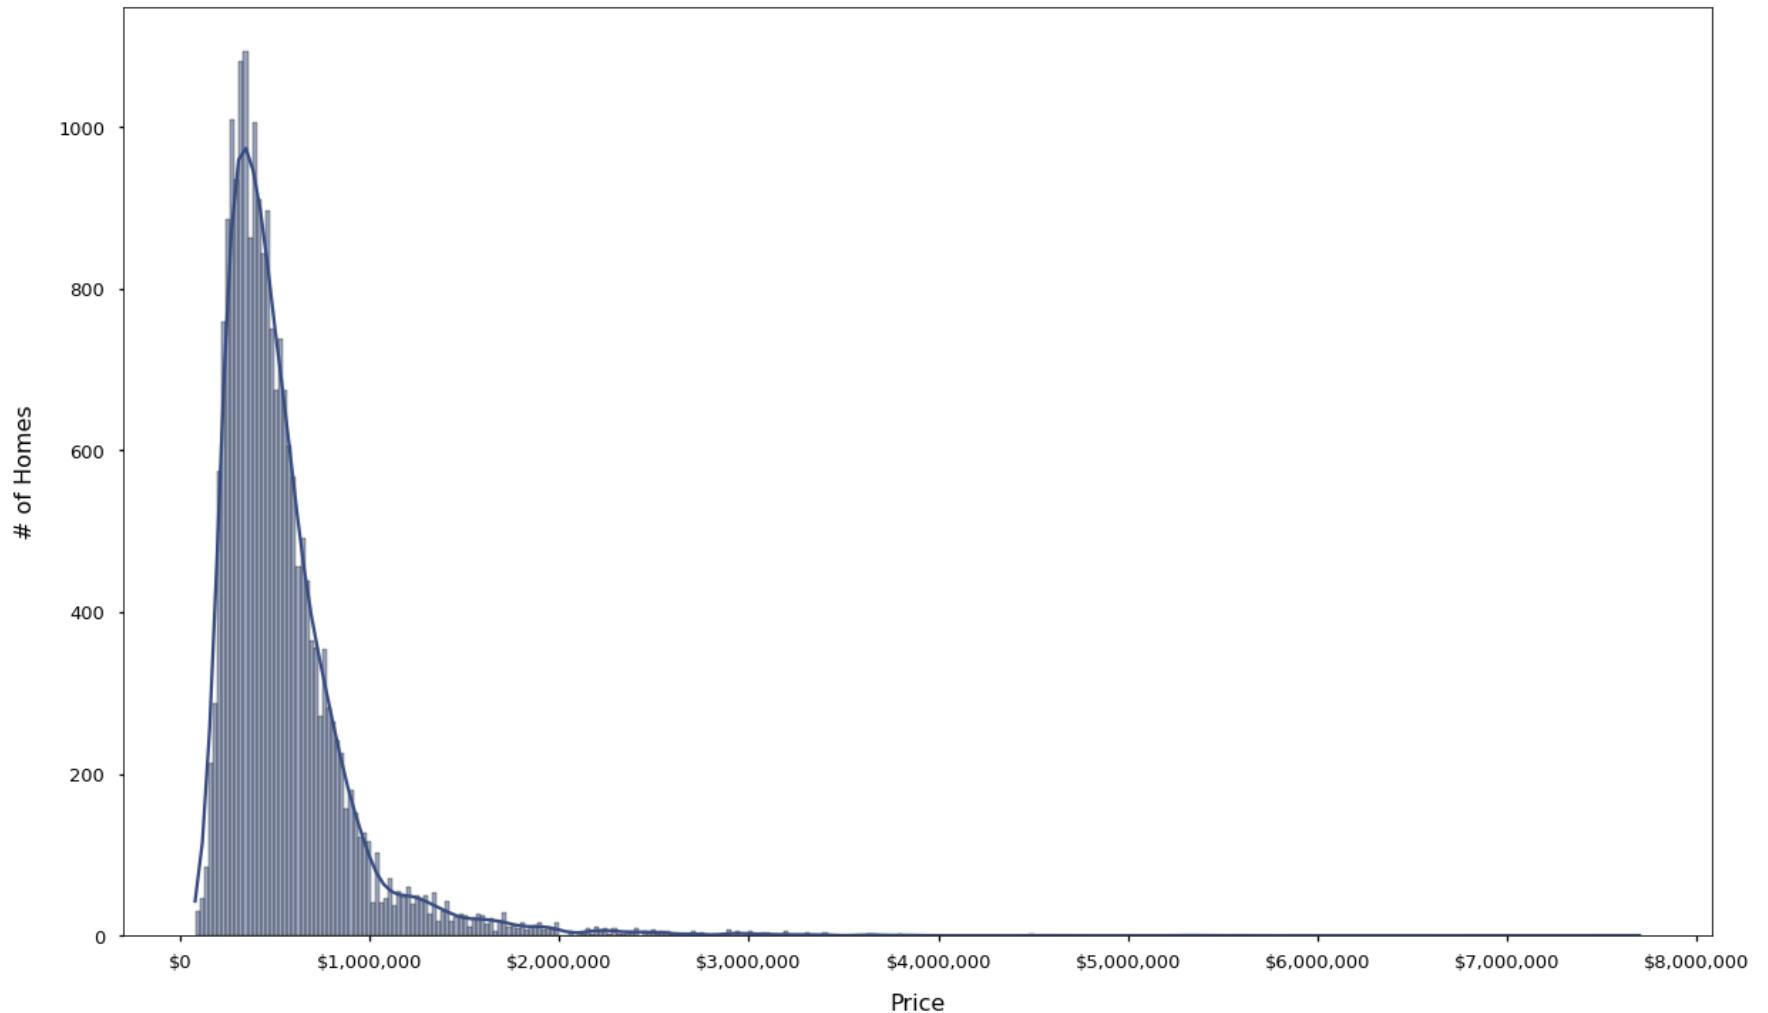
Average price is 540k, and median price is 450k, suggesting that there are some outliers driving the average higher than the median number. This is also realized with bedrooms, sqft_living. Which makes sense if we have some higher cost houses. Conventional thought, I would expect those to be similarly distributed as the price because that is what I think typically drives the price of a house. Another interesting component is yr_built. With how we think about houses evolving (i.e. # of bathrooms, square footage, etc) from 1900-2015, I am assuming we will see some interesting things with price as it is associated with homes over time.

```
In [26]: # Quick look at price, our dependent variable  
  
fig, ax = plt.subplots(figsize=(20,12))  
  
sns.histplot(data=df_clean,  
              x='price',  
              kde=True,  
              color=c("indigo"))  
  
ax.set_title('# of Homes by Price', pad=15, fontsize=22)  
ax.set_xlabel('Price', labelpad=15, fontsize=16)  
ax.set_ylabel('# of Homes', labelpad=15, fontsize=16)
```

```
fmt = '${x:,.0f}'
tick = mtick.StrMethodFormatter(fmt)
ax.xaxis.set_major_formatter(tick)

plt.show();
```

of Homes by Price



Distribution seems to be right skewed. Meaning that our mean is above our Median and being inflated by outliers (as mentioned above). However not the end of the world at this point. Going to continue to evaluate.

```
In [27]: # Quick look at price, our dependent variable transformed as a log value
```

```
fig, ax = plt.subplots(figsize=(20,12))

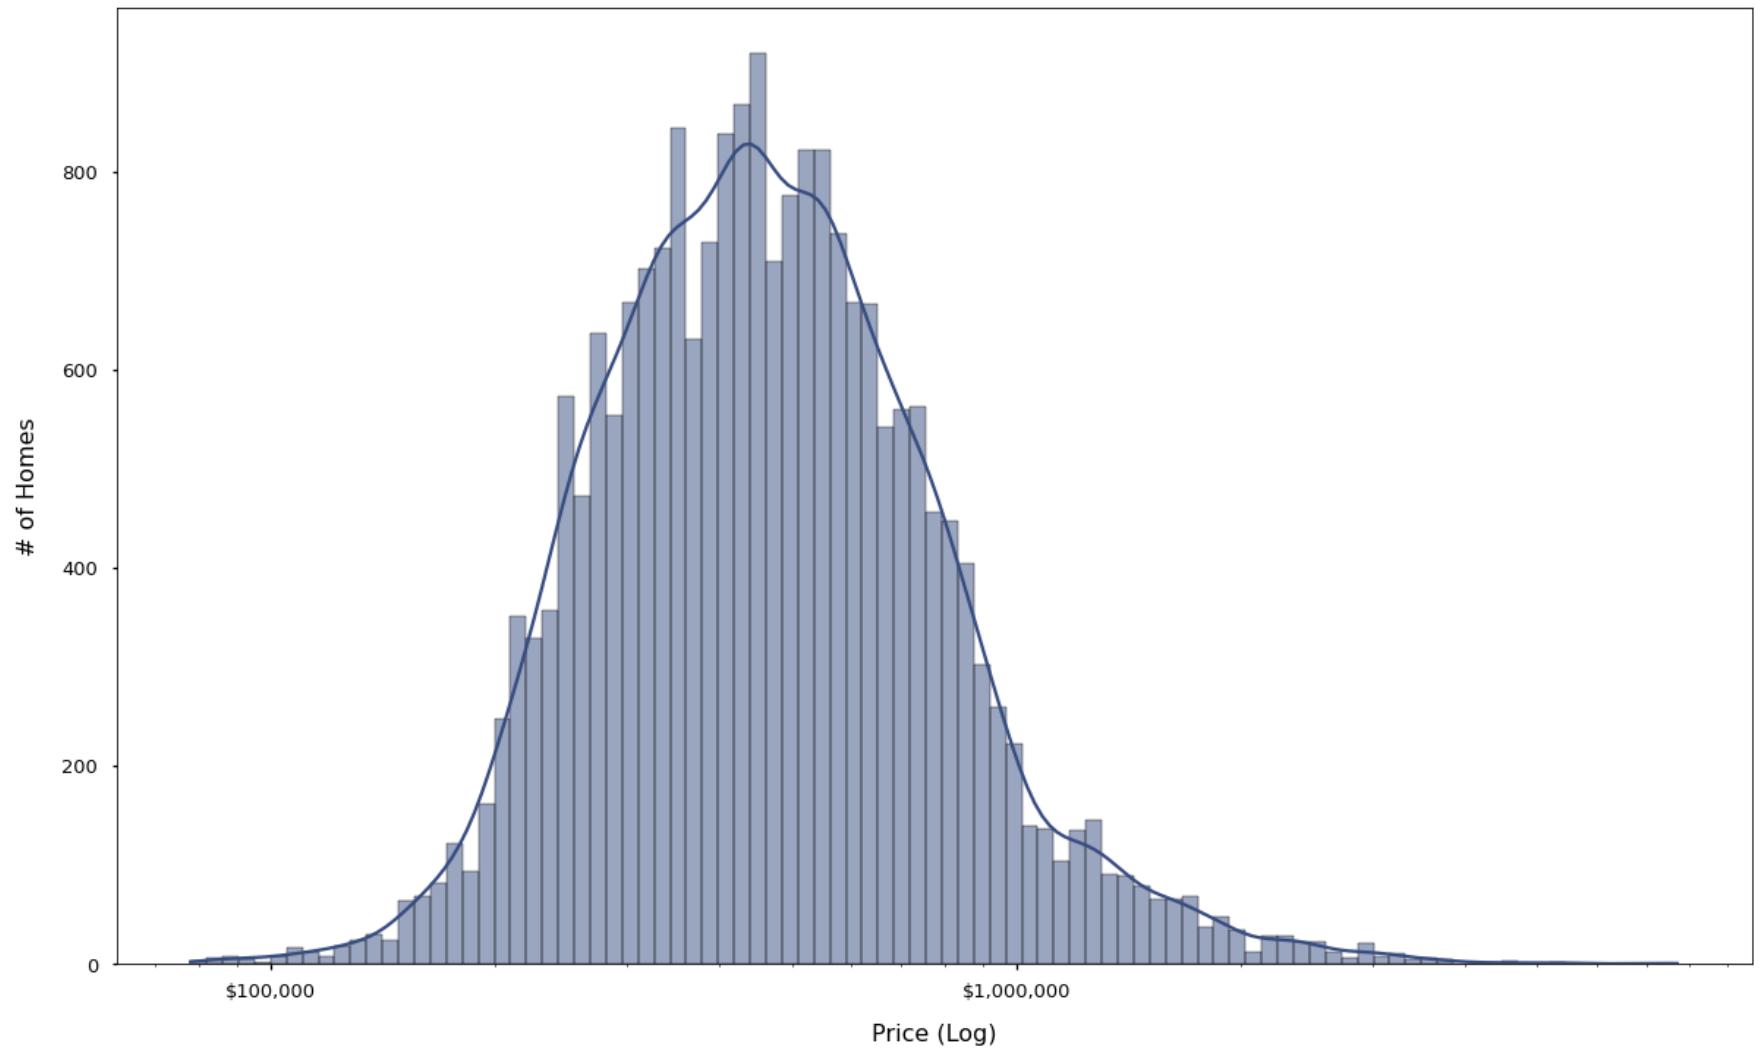
sns.histplot(data=df_clean,
              x='price',
              kde=True,
              log_scale=True,
              color=c("indigo"))

ax.set_title('# of Homes by Price (Log)', pad=15, fontsize=22)
ax.set_xlabel('Price (Log)', labelpad=15, fontsize=16)
ax.set_ylabel('# of Homes', labelpad=15, fontsize=16)

fmt = '${x:,.0f}'
tick = mtick.StrMethodFormatter(fmt)
ax.xaxis.set_major_formatter(tick)

plt.show();
```

of Homes by Price (Log)



This looks much better. I am going to keep my original values for the first values, but then will most likely change to a log scale in order to distribute the data normally.

```
In [28]: # Quick look at price distribution as it relates to a box plot/ distribution.  
fig, ax = plt.subplots(figsize=(20,6))  
ax.set_title('Price Distribution')
```

```

bp = sns.boxplot(data=df_clean,
                  x='price',
                  palette=pal)

fmt = '${x:,.0f}'
tick = mtick.StrMethodFormatter(fmt)
ax.xaxis.set_major_formatter(tick)

bp.set_xlabel('Price');

```



Going to check to see how many records we would drop if we remove outliers outside of IQR.

In [29]:

```

# Developing variables for IQR removal, and record count.

price_Q1 = df_clean['price'].quantile(0.25)
price_Q3 = df_clean['price'].quantile(0.75)

IQR = price_Q3 - price_Q1

lower_bnd = price_Q1 - 1.5 * IQR
upper_bnd = price_Q3 + 1.5 * IQR

```

In [30]:

```

# Creating variable for lower bound

outliers_low = df_clean['price'] > lower_bnd

```

```
In [31]: # Creating variable for lower bound
```

```
outliers_up = df_clean['price'] < upper_bnd
```

```
In [32]: print("% of Records Retained After Oulier Removal:",
           "{:.1%}".format(len(df_clean['price'][~outliers_low | outliers_up])/ len(df_clean)))
      )
```

```
% of Records Retained After Oulier Removal: 94.6%
```

The removal of outliers really doesn't impact our sample size. Which is good. We will most likely remove these outliers later as we iterate on our regression models.

```
In [33]: # Taking a look at the new distribution less outliers
```

```
price_check = df_clean['price'][~outliers_low | outliers_up]

fig, ax = plt.subplots(figsize=(20,12))

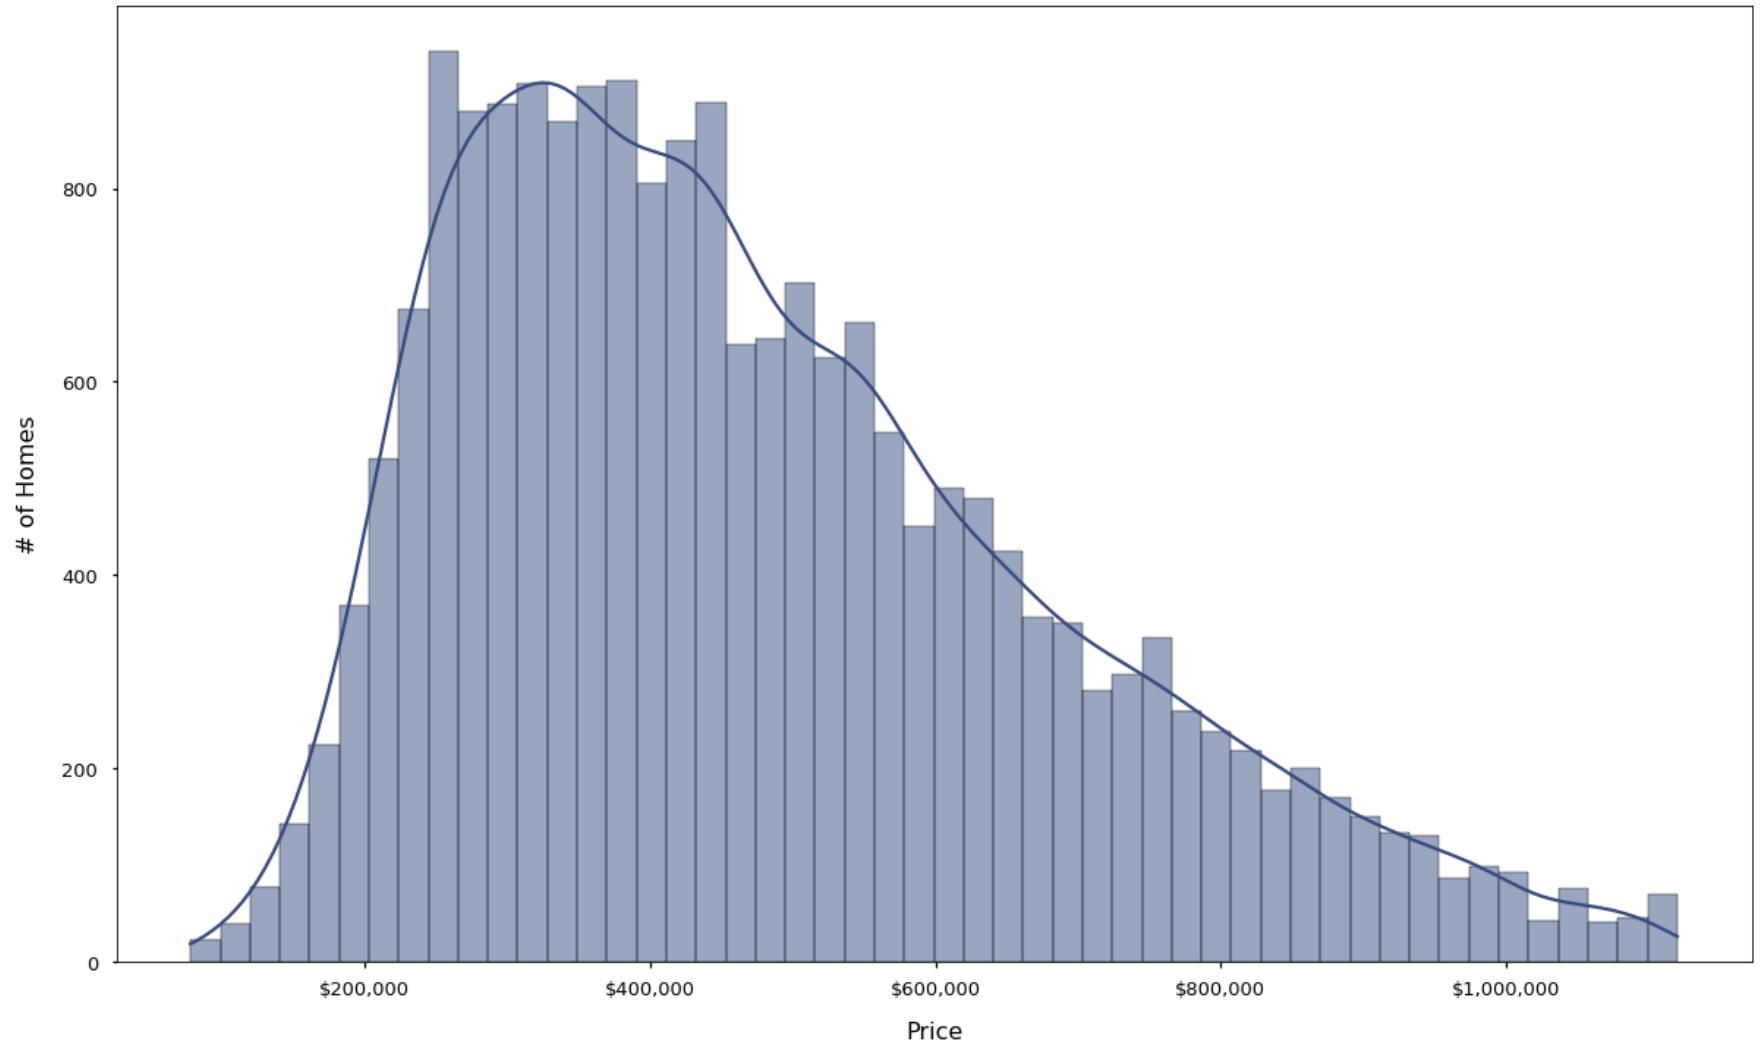
sns.histplot(data=price_check,
              kde=True,
              color=c("indigo"))

ax.set_title('# of Homes by Price - Outliers Removed', pad=15, fontsize=22)
ax.set_xlabel('Price', labelpad=15, fontsize=16)
ax.set_ylabel('# of Homes', labelpad=15, fontsize=16)

fmt = '${x:,.0f}'
tick = mtick.StrMethodFormatter(fmt)
ax.xaxis.set_major_formatter(tick)

plt.show();
```

of Homes by Price - Outliers Removed



This looks much better even without the log transformation. What we are looking for is a normal distribution of data, which will ensure our model will have the best opportunity to be objective in its results/ predictions.

```
In [34]: # Boxplots for # of bedrooms and price

fig, ax = plt.subplots(figsize=(20,12))

ax.set_title('Price Distribution by # of Bedrooms')

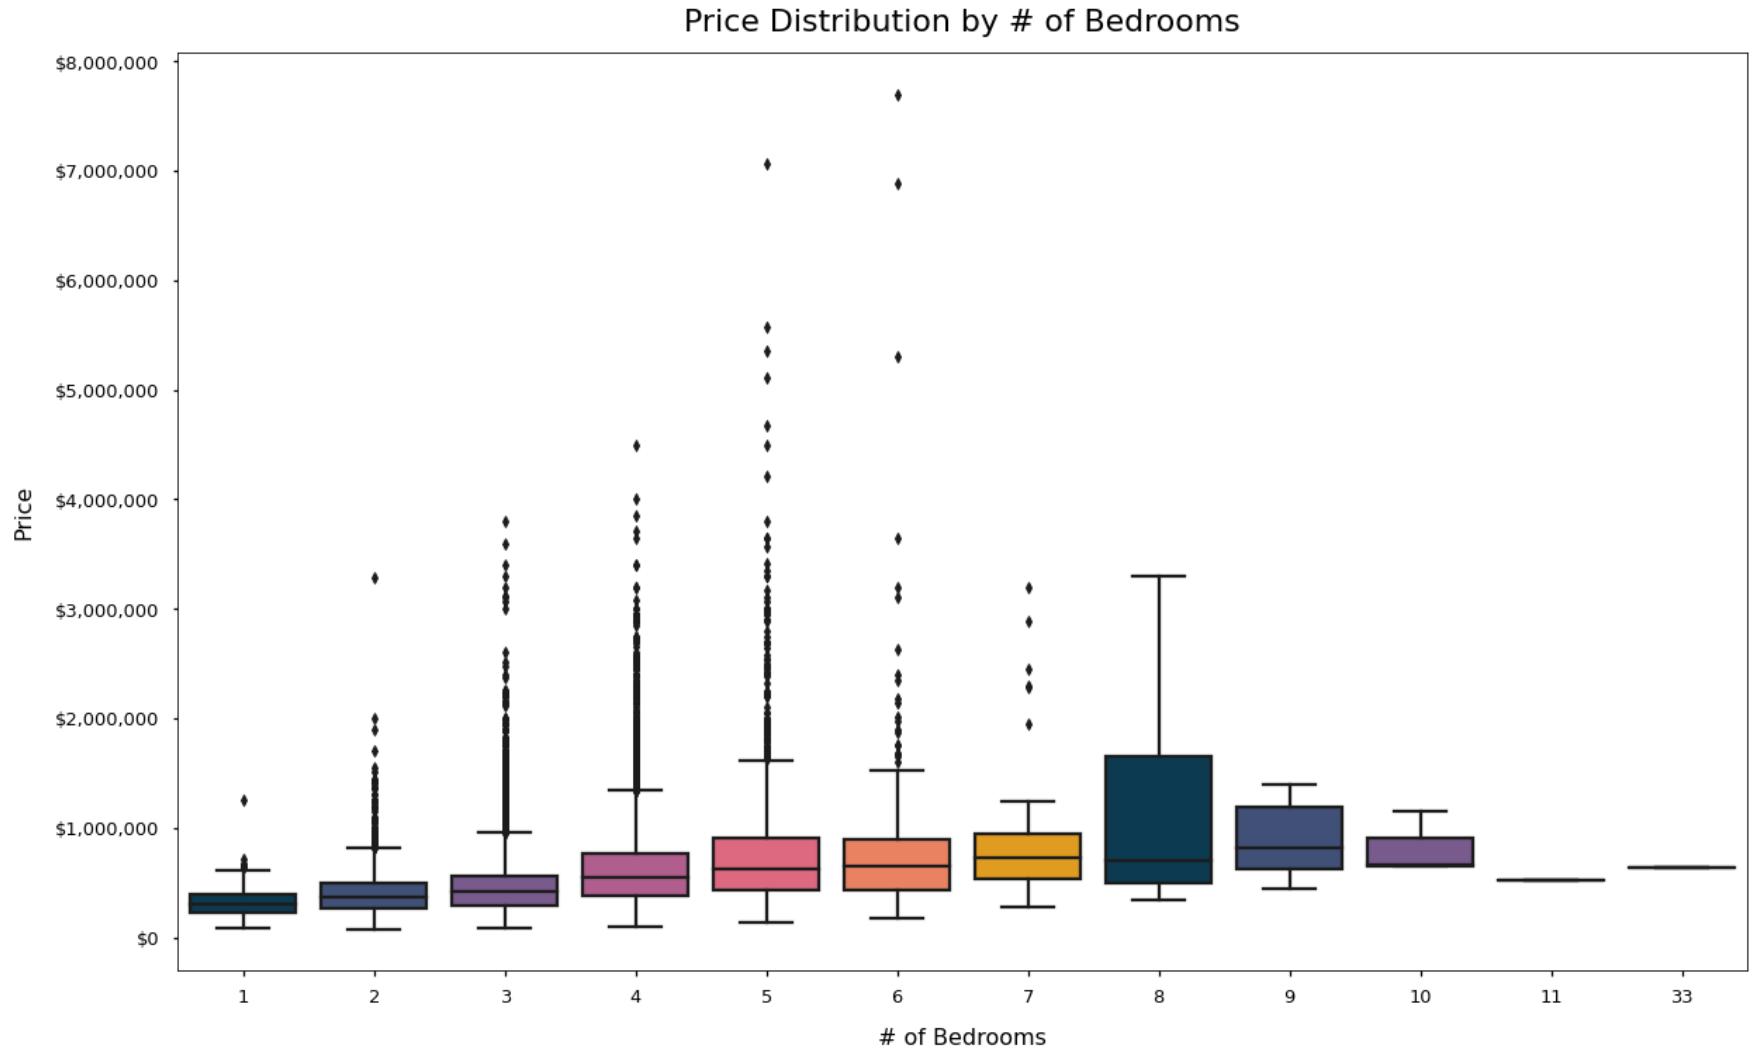
bp = sns.boxplot(data=df_clean,
```

```
y='price',
x='bedrooms',
palette=pal)

fmt = '${x:.0f}'
tick = mtick.StrMethodFormatter(fmt)
ax.yaxis.set_major_formatter(tick)

ax.set_title('Price Distribution by # of Bedrooms', pad=15, fontsize=22)
ax.set_xlabel('# of Bedrooms', labelpad=15, fontsize=16)
ax.set_ylabel('Price', labelpad=15, fontsize=16)

plt.show();
```



Removing outliers seems like it would be a good option. Also probably easier than trying to interpret log transformed coefficients. Will most likely try outlier removal first, then log transform.

```
In [35]: # Taking a look at bedrooms only  
fig, ax = plt.subplots(figsize=(20,6))  
  
ax.set_title('# of Bedrooms Distribution')  
  
bp = sns.boxplot(data=df_clean,  
                  x='bedrooms',
```

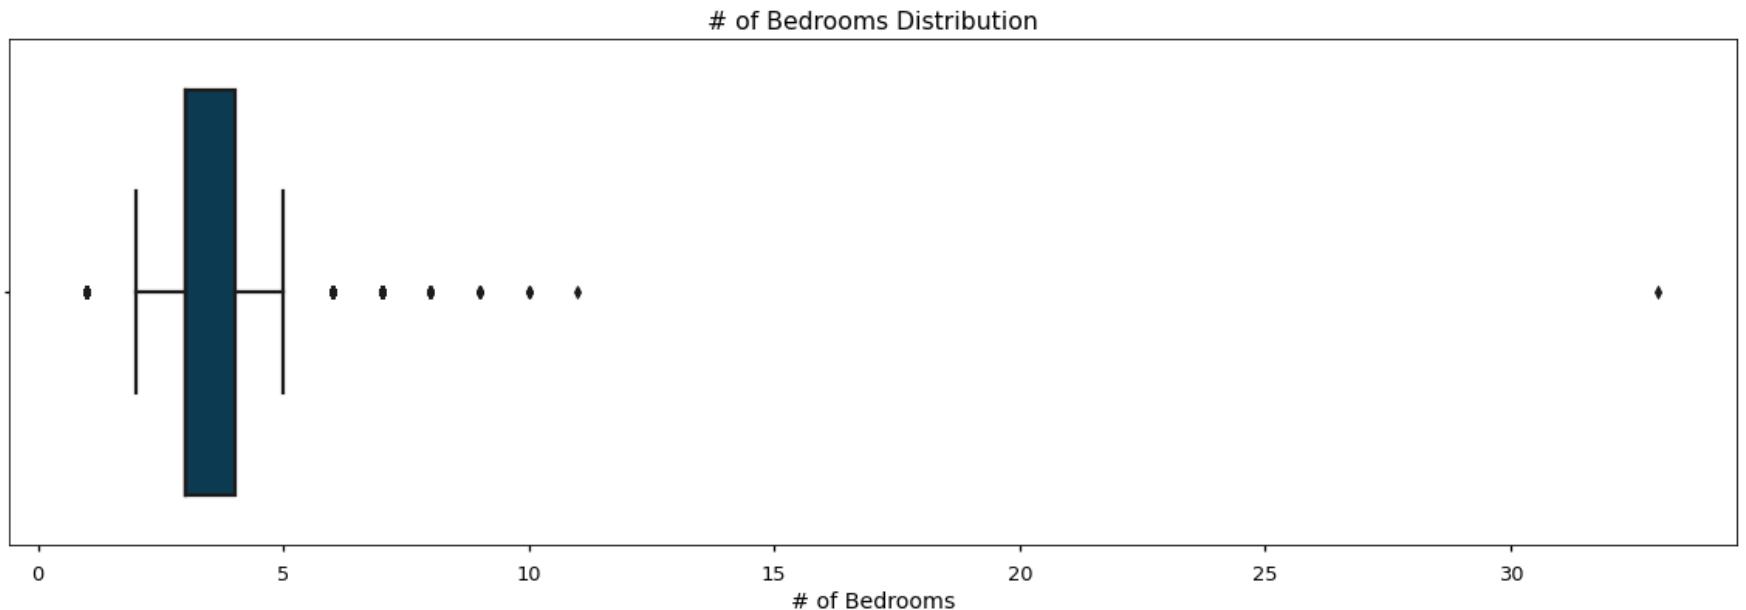
```

    palette=pal)

bp.set_ylabel('')
bp.set_xlabel('# of Bedrooms')

plt.show();

```



There doesn't seem to be a large difference between the median home price, as the number of bedrooms increase. In fact, most of the outliers for price, are situated within the non-outlier values of bedrooms. With that being said, there does seem to be a subtle increase in price, as the number of bedrooms increase. The one record with the 33 bedrooms is a bit bothersome, however it doesn't look like the house really got the benefit of having that many bedrooms because the price was just above the average, and the sqft_living is below the median. To me this looks like an error in the data in that it should be represented as 3 bedrooms, instead of 33 bedrooms. The median, and the 25th percentiles are the same for bedrooms, with 3 bedrooms.

In [36]: # Isolating the bedroom row with 33 bedrooms for further inspection

```

df_clean.loc[df_clean['bedrooms'] == 33]

```

Out[36]:

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	sqft_above	sqft_basement	yr_built	sqft_living15	sqft_lot15	i
15856	640000.0	33	1.75	1620	6000	1.0	1040	580.0	1947	1330	4700	

In [37]: # Replacing 33 bedroom value with 3 bedrooms

```
df_clean['bedrooms'] = df['bedrooms'].replace([33],3)
```

```
In [38]: # Quick view to make sure our operation worked.
```

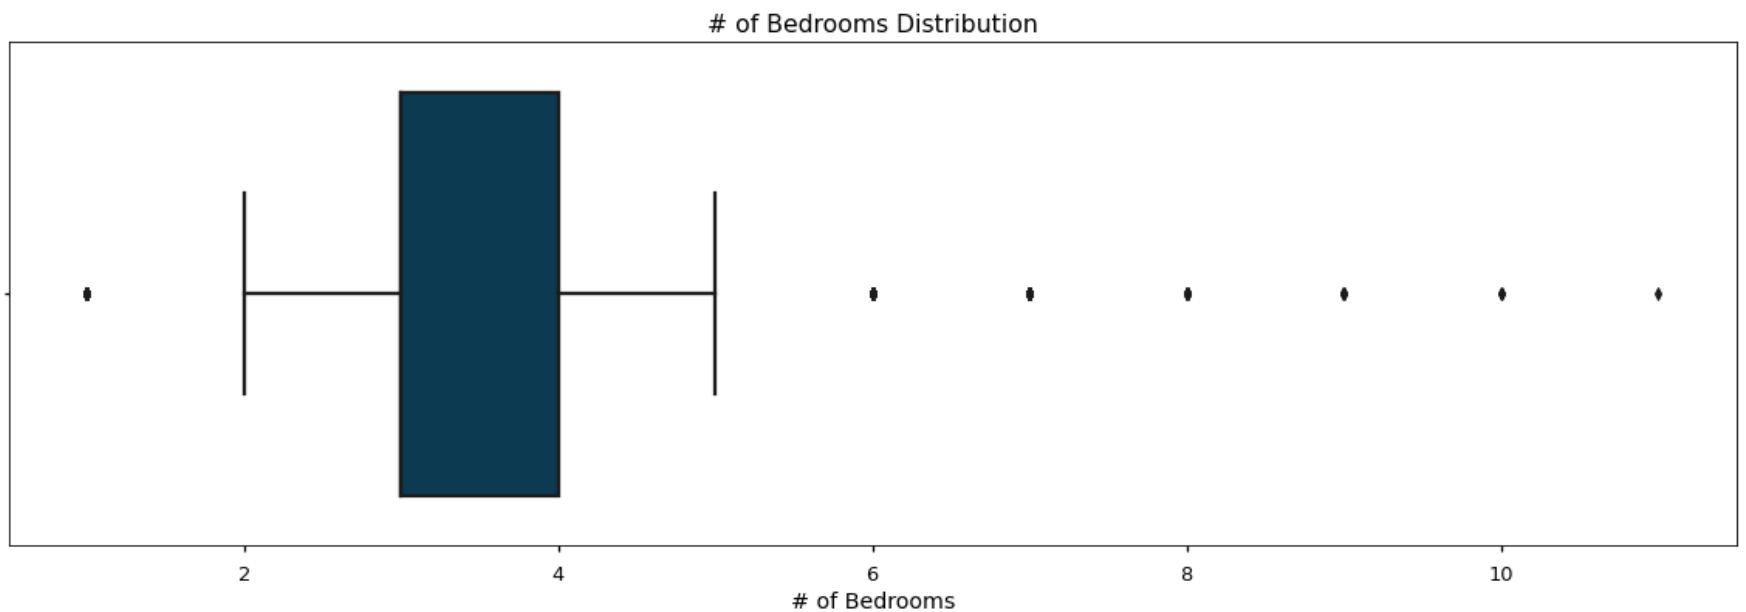
```
fig, ax = plt.subplots(figsize=(20,6))

ax.set_title('# of Bedrooms Distribution')

bp = sns.boxplot(data=df_clean,
                  x='bedrooms',
                  palette=pal)

bp.set_ylabel('')
bp.set_xlabel('# of Bedrooms')

plt.show();
```



```
In [39]: # Scatter to look at the data between our bedroom variable and price.
```

```
fig, ax = plt.subplots(figsize = (20,12))

sns.regplot(x='bedrooms',
            y='price',
            data=df_clean,
```

```

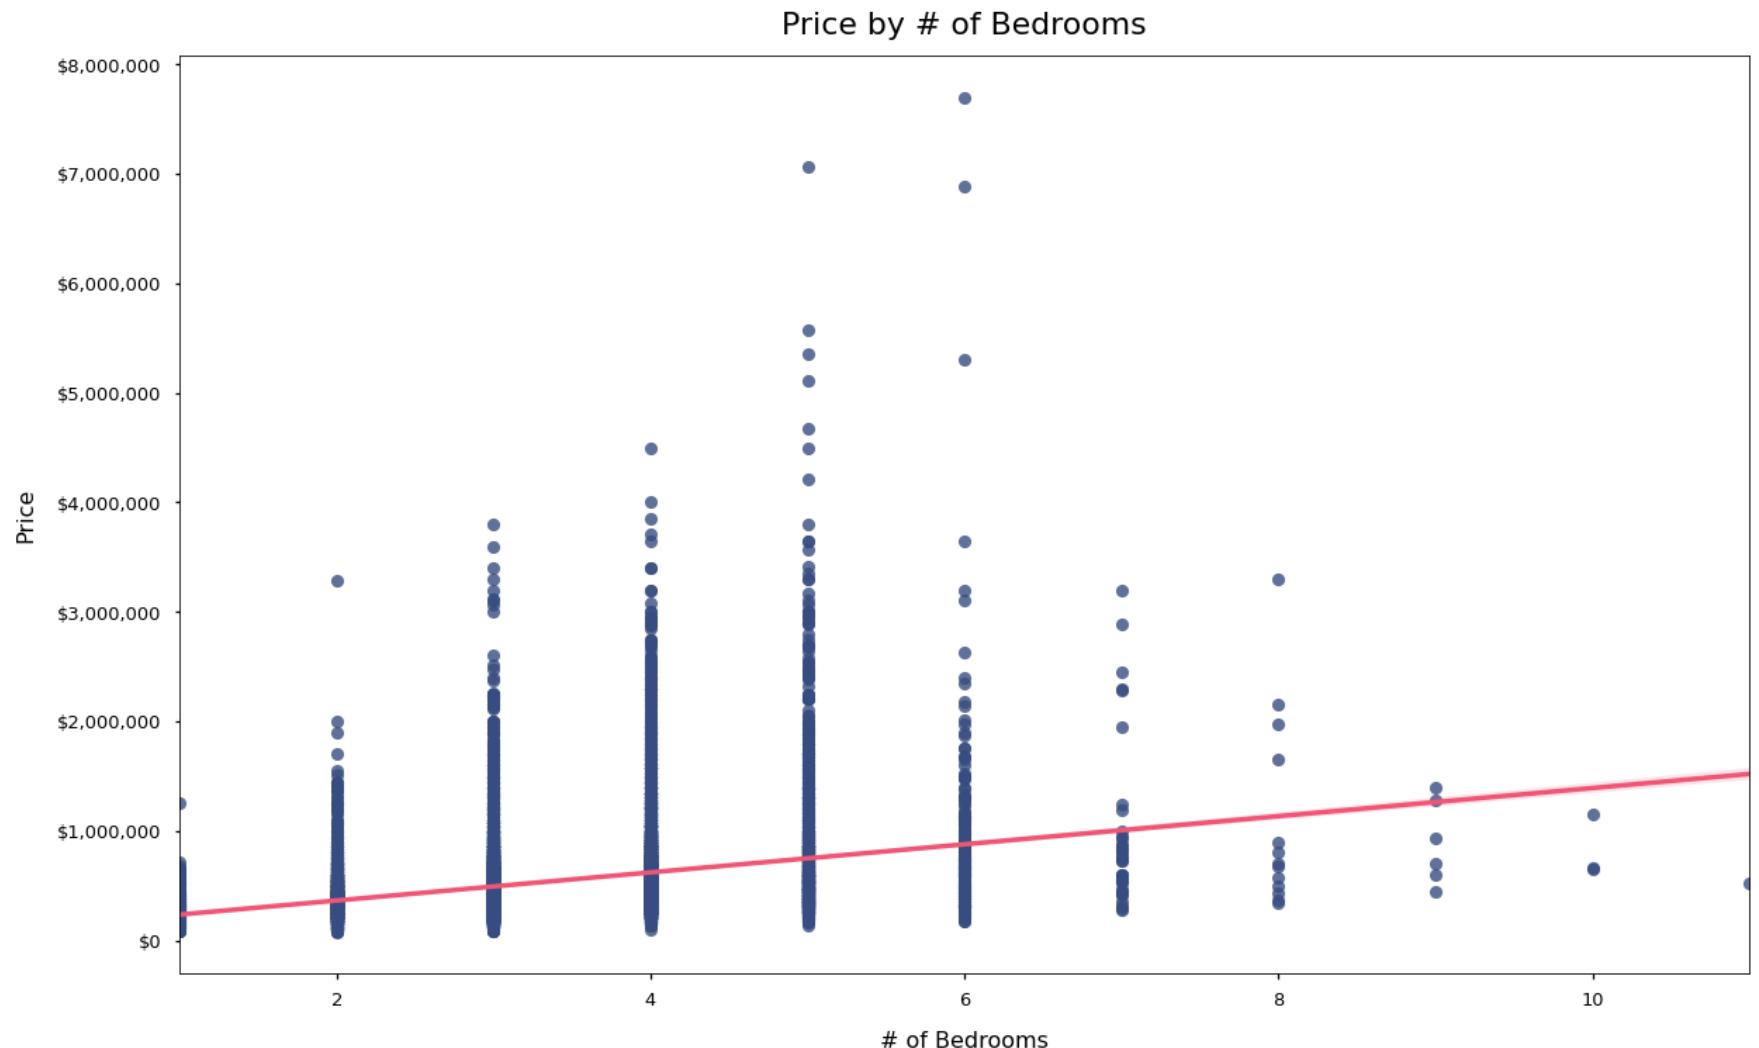
    scatter_kws={"color": c("indigo")},
    line_kws={"color": c("peach")})

fmt = '${x:,.0f}'
tick = mtick.StrMethodFormatter(fmt)
ax.yaxis.set_major_formatter(tick)

ax.set_title('Price by # of Bedrooms', pad=15, fontsize=22)
ax.set_xlabel('# of Bedrooms', labelpad=15, fontsize=16)
ax.set_ylabel('Price', labelpad=15, fontsize=16)

plt.show();

```



There seems to be a moderate positive relationship between the number of bedrooms and our price variable. The data looks much more consistent without the 33 bedrooms. I think the 33 bedrooms may have been inflating our expectations of the influence that bedrooms actually has on price. There are a few outliers that remain, we will run a regression first and then determine if removing additional will be valuable. Moving on to bathrooms next.

```
In [40]: # Quick view to make sure our operation worked.

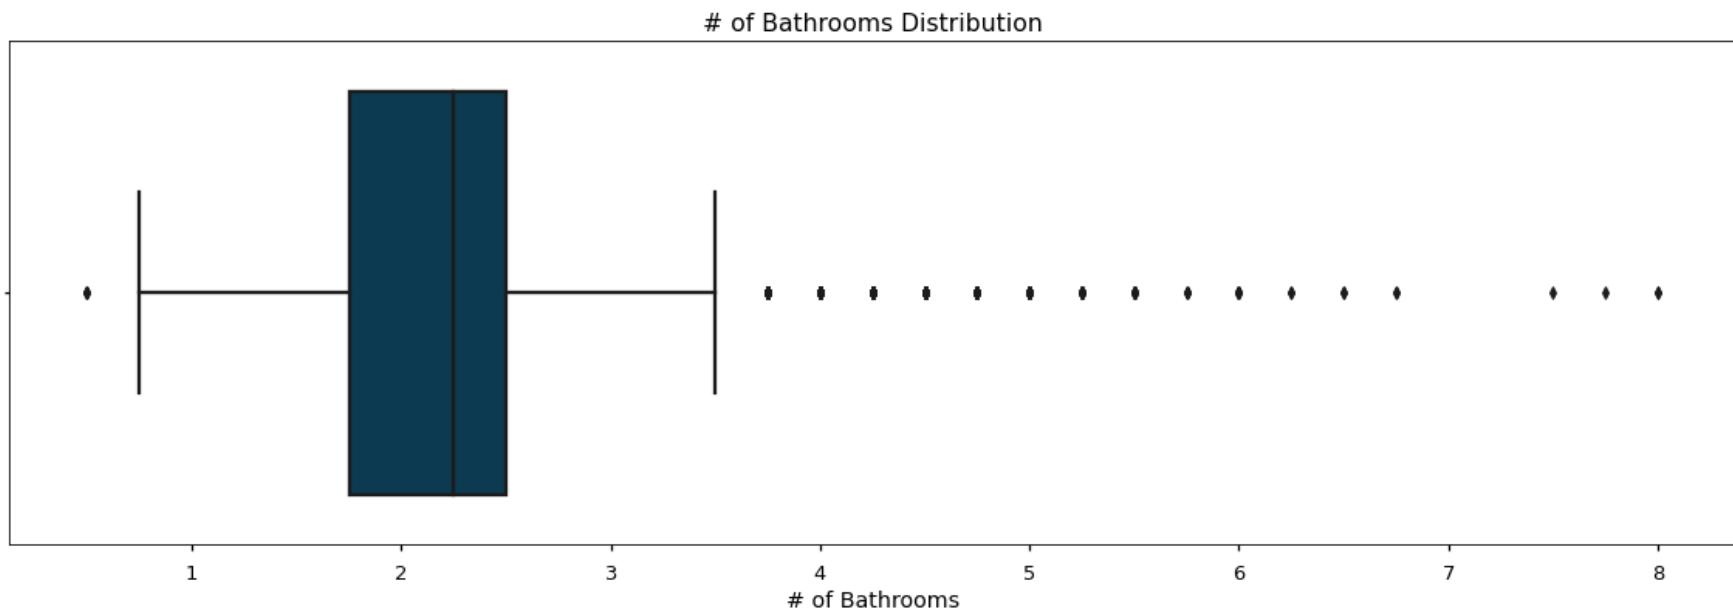
fig, ax = plt.subplots(figsize=(20,6))

ax.set_title('# of Bathrooms Distribution')

bp = sns.boxplot(data=df_clean,
                  x='bathrooms',
                  palette=pal)

bp.set_ylabel('')
bp.set_xlabel('# of Bathrooms')

plt.show();
```



A few outliers based on the boxplot, while most bathrooms are between .75 and 3.5.

```
In [41]: # Creating a chart for price distribution by number of bathrooms. Trying to understand
# distribution within each bathroom size to see if outliers driving any values in our assumed linear
```

```
# relationship.

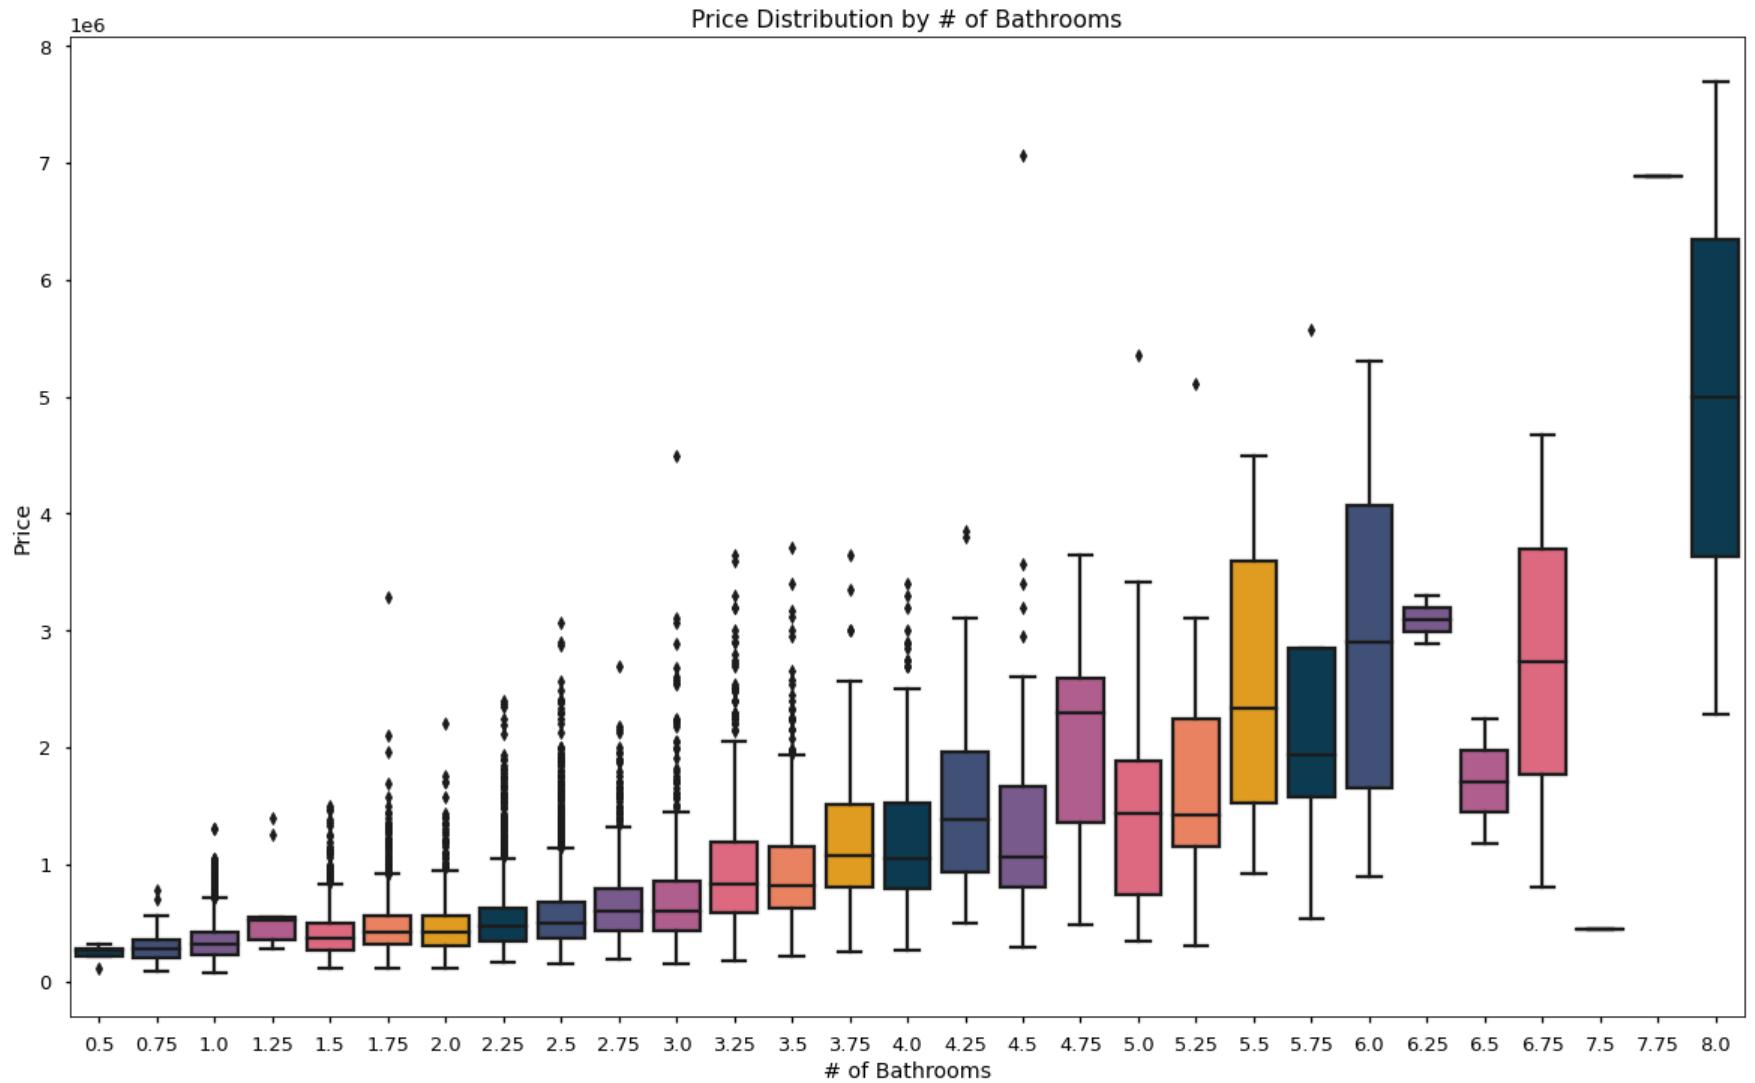
fig, ax = plt.subplots(figsize=(20,12))

ax.set_title('Price Distribution by # of Bathrooms')

bp = sns.boxplot(data=df_clean,
                  y='price',
                  x='bathrooms',
                  palette=pal)

bp.set_ylabel('Price')
bp.set_xlabel('# of Bathrooms')

plt.show();
```



Not crazy about all these .25, and .75 baths. Especially since the lower end of values have a lot of overlap in price distribution. Going to create another column that rounds up to the nearest half bath.

```
In [42]: # Create a new column for rounded bathrooms.
```

```
df_clean['baths_rnd'] = (np.floor(df_clean['bathrooms'] * 2) / 2)
```

```
In [43]: # Quick view of our columns to make sure we're transforming the way we expect.
```

```
df_clean[['baths_rnd', 'bathrooms']]
```

Out[43]:

	baths_rnd	bathrooms
0	1.0	1.00
1	2.0	2.25
2	1.0	1.00
3	3.0	3.00
4	2.0	2.00
...
21592	2.5	2.50
21593	2.5	2.50
21594	0.5	0.75
21595	2.5	2.50
21596	0.5	0.75

21597 rows × 2 columns

In [44]:

```
# Scatter plot for my rounded bathrooms and the
fig, ax = plt.subplots(figsize=(20,12))

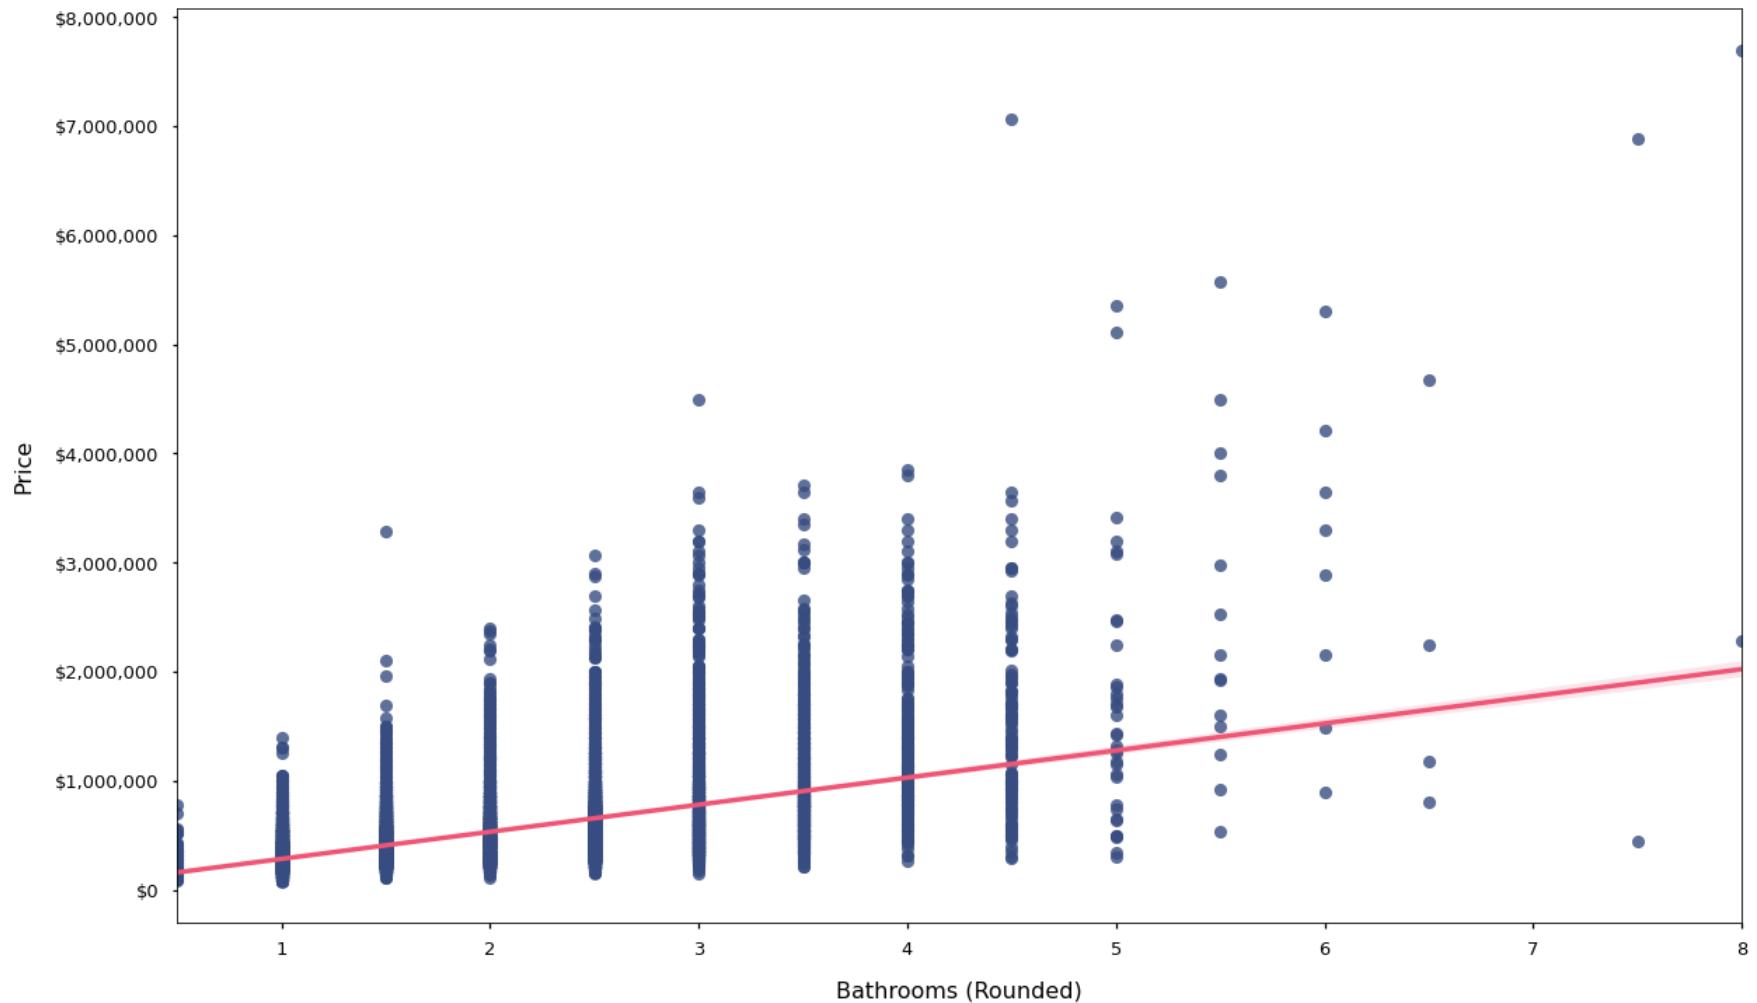
sns.regplot(x='baths_rnd',
             y='price',
             data=df_clean,
             scatter_kws={"color": c("indigo")},
             line_kws={"color": c("peach")})

fmt = '${x:,.0f}'
tick = mtick.StrMethodFormatter(fmt)
ax.yaxis.set_major_formatter(tick)

ax.set_title('Price vs Bathrooms (Rounded)', pad=15, fontsize=22)
ax.set_xlabel('Bathrooms (Rounded)', labelpad=15, fontsize=16)
ax.set_ylabel('Price', labelpad=15, fontsize=16)

plt.show();
```

Price vs Bathrooms (Rounded)



Once again a slight correlation with price. At a glance it looks like bathrooms have a stronger impact on price vs # the number of bedrooms listed on a house. If this does turn out to be the case, then suggesting bathroom upgrades and potentially rounding up existing bathrooms would be easy enhancements to increase the value of a home for our customers.

```
In [45]: # Let's take a quick look at price vs living area square footage.  
fig, ax = plt.subplots(figsize=(20,12))  
  
sns.regplot(x='sqft_living',  
            y='price',  
            data=df_clean,
```

```

        scatter_kws={"color": c("indigo")},  

        line_kws={"color": c("peach")})  
  

fmt = '${x:,.0f}'  

tick = mtick.StrMethodFormatter(fmt)  

ax.yaxis.set_major_formatter(tick)  
  

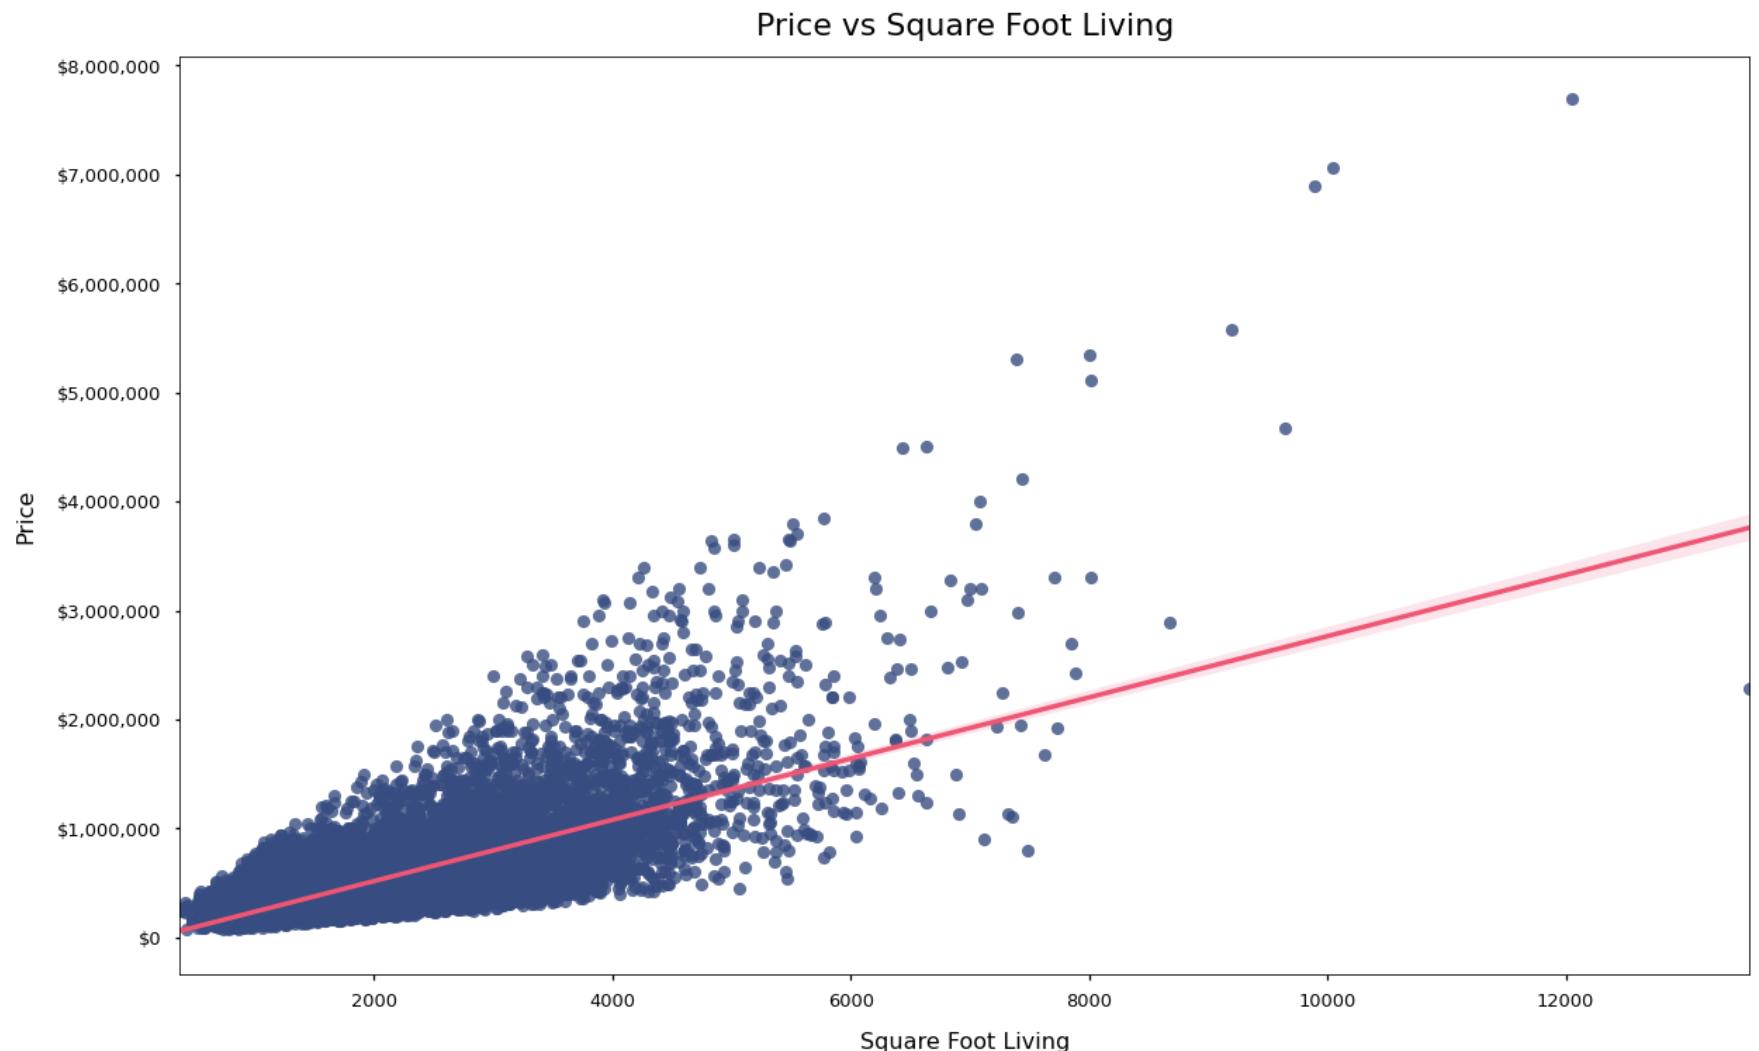
ax.set_title('Price vs Square Foot Living', pad=15, fontsize=22)  

ax.set_xlabel('Square Foot Living', labelpad=15, fontsize=16)  

ax.set_ylabel('Price', labelpad=15, fontsize=16)  
  

plt.show();

```



So at a glance, it seems as if living area square footage will have the highest impact on the sale price of a home. We will see if this holds true in our analysis/ regression. In addition, we will assume that our rounded bathrooms will also have a larger impact on the overall price.

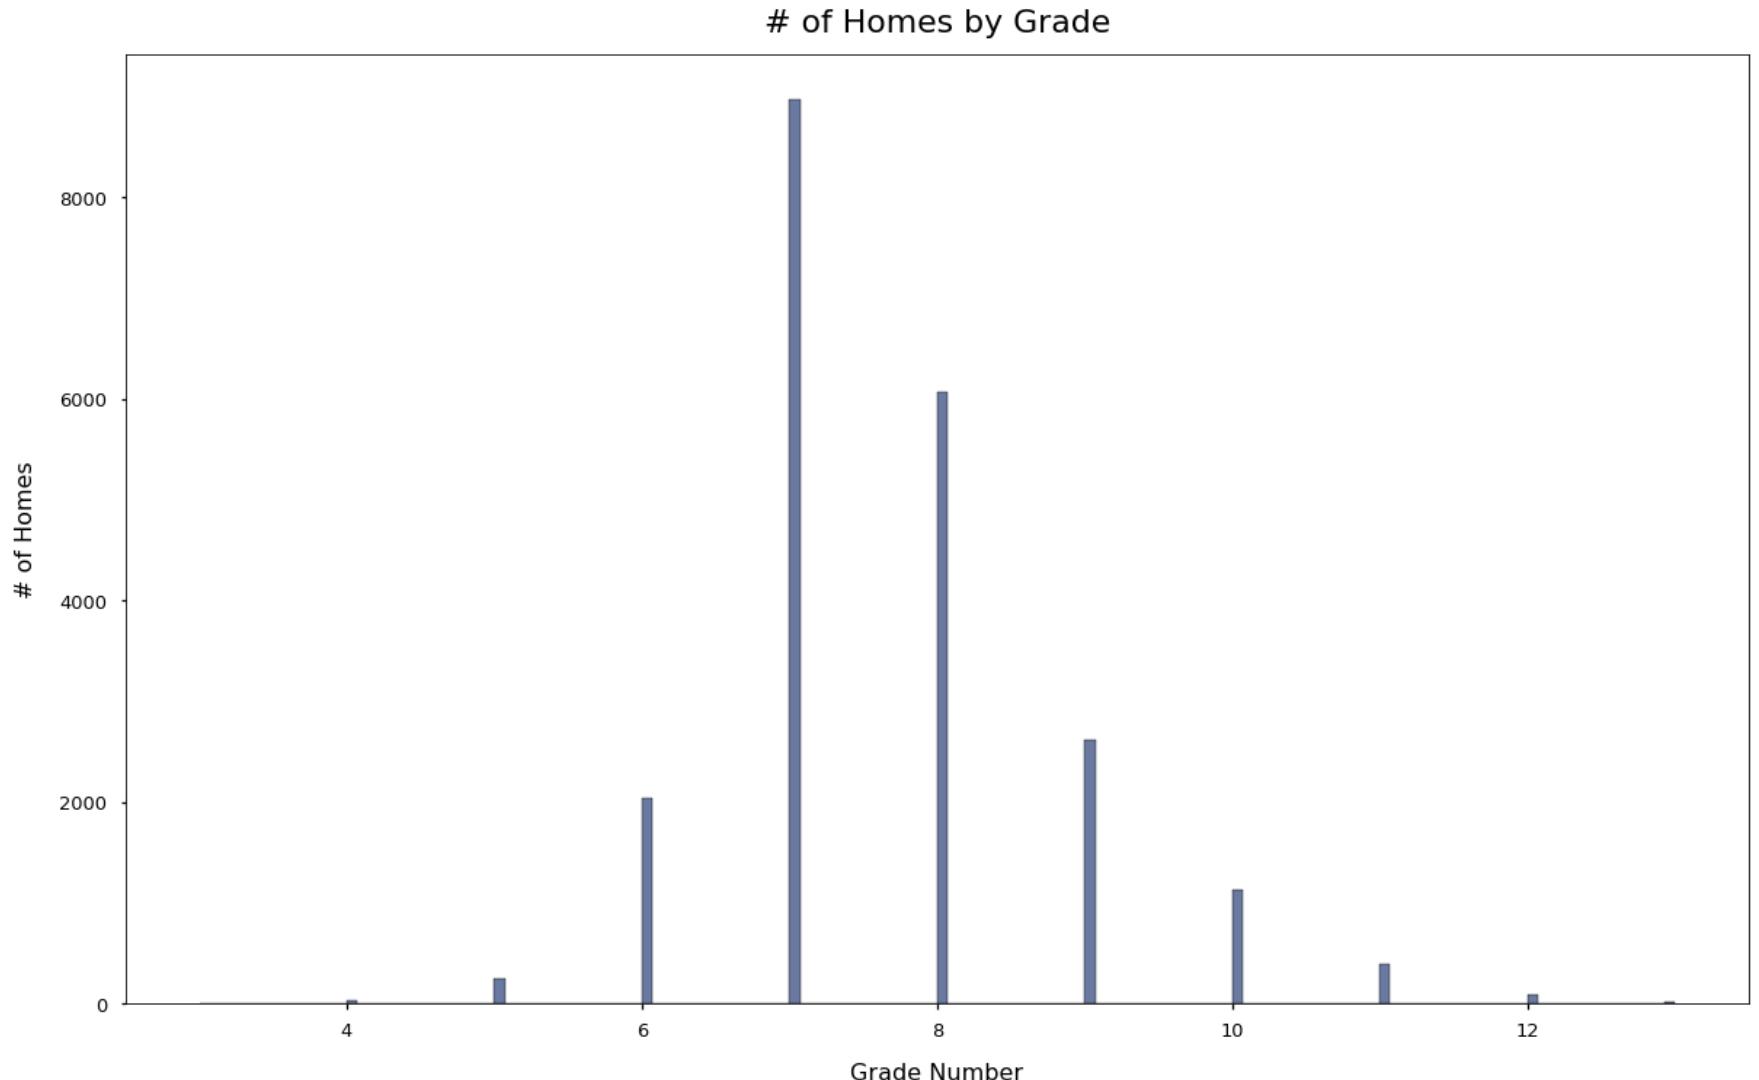
```
In [46]: # Taking a look at our grade_num feature
```

```
fig, ax = plt.subplots(figsize=(20,12))

sns.histplot(data=df_clean,
              x='grade_num',
              color=c("indigo"))

ax.set_title('# of Homes by Grade', pad=15, fontsize=22)
ax.set_xlabel('Grade Number', labelpad=15, fontsize=16)
ax.set_ylabel('# of Homes', labelpad=15, fontsize=16)

plt.show();
```



Average grade_num is 7.65 while the median is 7. Shows similar as majority of homes in this dataset have a rating of 6-9.

```
In [47]: # Let's take a quick look at price vs living area square footage.

fig, ax = plt.subplots(figsize=(20,12))

sns.regplot(x='grade_num',
            y='price',
            data=df_clean,
            scatter_kws={"color": c("indigo")},
```

```

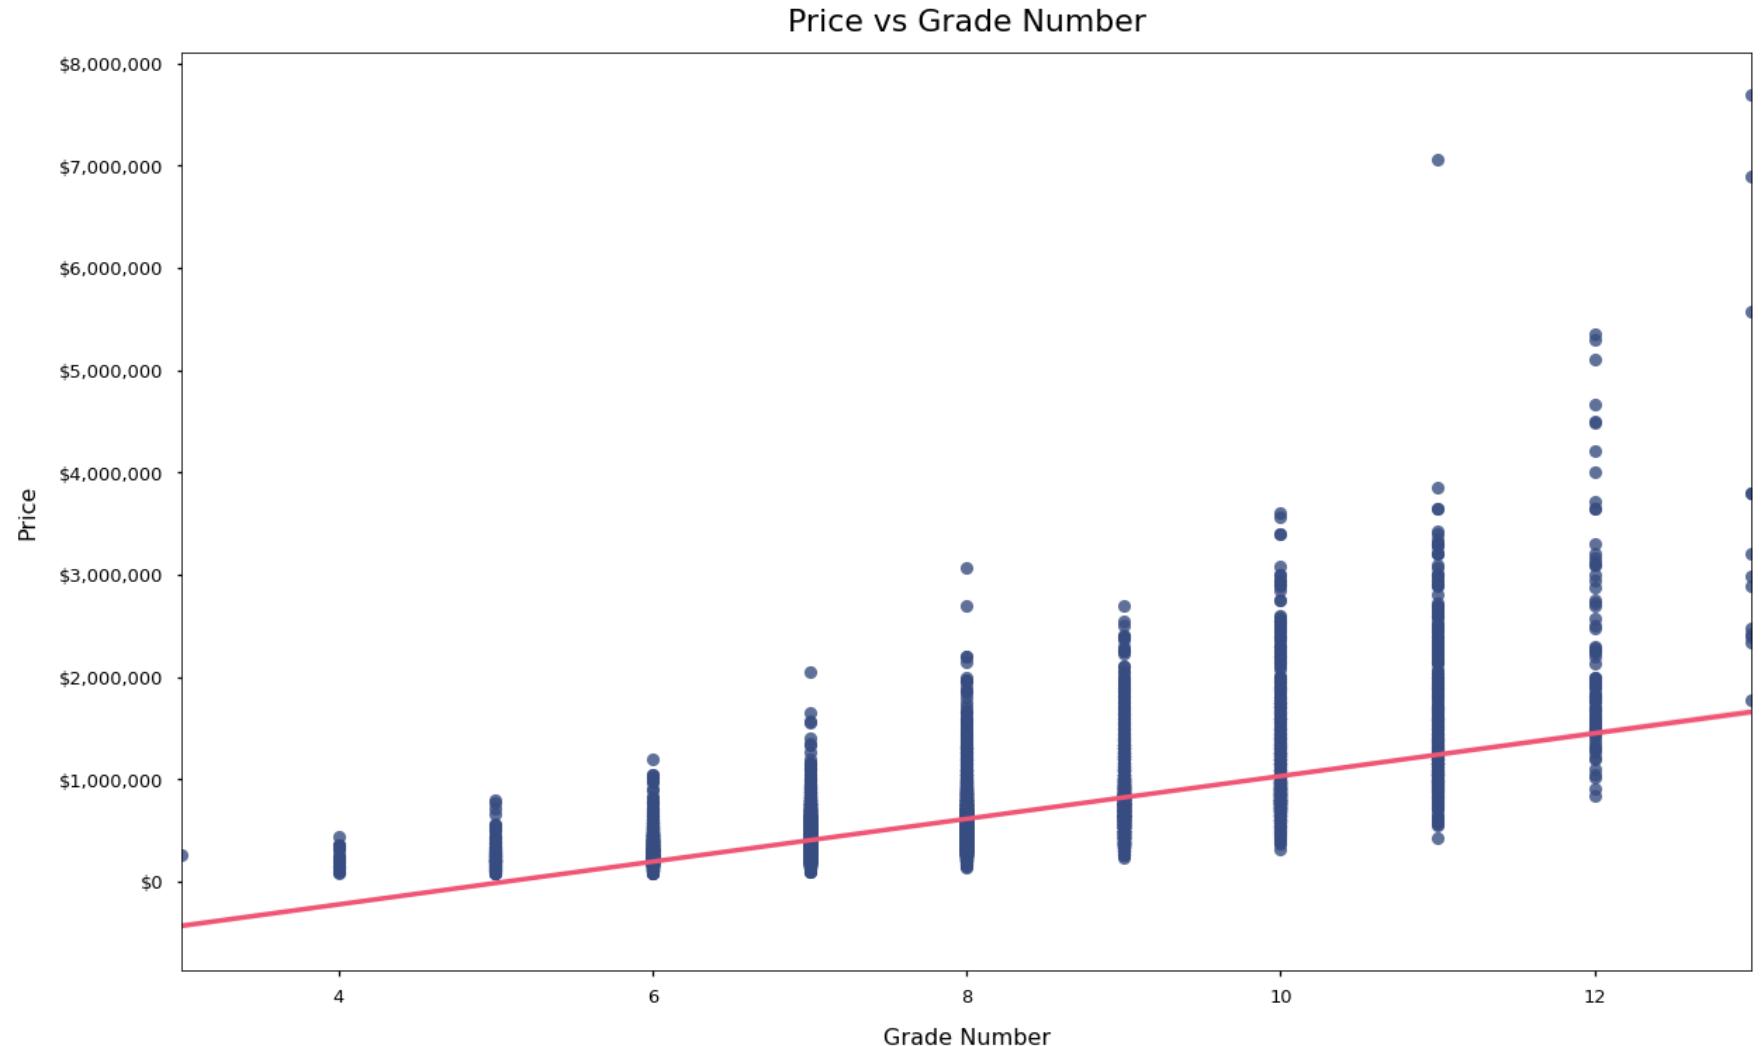
line_kws={"color": c("peach")})

fmt = '${x:.0f}'
tick = mtick.StrMethodFormatter(fmt)
ax.yaxis.set_major_formatter(tick)

ax.set_title('Price vs Grade Number', pad=15, fontsize=22)
ax.set_xlabel('Grade Number', labelpad=15, fontsize=16)
ax.set_ylabel('Price', labelpad=15, fontsize=16)

plt.show();

```



As shown in our correlation matrix, this view also shows a positive relationship between the grade number of a home, and the price. Thus indicating that if the customers's were to make upgrades (i.e. materials and builds), they could see the benefit when trying to sell their home.

Regression Modeling

Model 1

Will first create a model based on our dataframe as it sits right now. This will serve as a baseline of what we could potentially improve upon, or maybe not approve on.

```
In [48]: # Creating our first regression on basic parameters as a baseline.

# Create target
target1 = df_clean['price']

# Create predictors
predictors1 = df_clean.drop(['price'], axis=1)

# Create model intercept
predictors_int1 = sm.add_constant(predictors1)

# Fit model to data
modell = sm.OLS(df_clean['price'], predictors_int1).fit()
```

```
In [49]: modell.summary()
```

```
Out[49]: OLS Regression Results
Dep. Variable: price R-squared: 0.632
Model: OLS Adj. R-squared: 0.632
Method: Least Squares F-statistic: 2471.
Date: Thu, 27 Oct 2022 Prob (F-statistic): 0.00
Time: 20:57:03 Log-Likelihood: -2.9660e+05
No. Observations: 21597 AIC: 5.932e+05
Df Residuals: 21581 BIC: 5.934e+05
Df Model: 15
```

Covariance Type: nonrobust

	coef	std err	t	P> t	[0.025	0.975]
const	7.011e+06	1.35e+05	51.892	0.000	6.75e+06	7.28e+06
bedrooms	-1.188e+05	4754.683	-24.991	0.000	-1.28e+05	-1.1e+05
bathrooms	1.159e+05	1.35e+04	8.586	0.000	8.95e+04	1.42e+05
sqft_living	389.2911	22.329	17.434	0.000	345.524	433.058
sqft_lot	0.0210	0.053	0.396	0.692	-0.083	0.125
floors	2.966e+04	3929.084	7.549	0.000	2.2e+04	3.74e+04
sqft_above	-135.7037	22.323	-6.079	0.000	-179.458	-91.950
sqft_basement	54.7352	19.837	2.759	0.006	15.853	93.618
yr_built	-4035.6994	68.677	-58.764	0.000	-4170.311	-3901.088
sqft_living15	38.7985	3.686	10.527	0.000	31.574	46.023
sqft_lot15	-0.5063	0.081	-6.256	0.000	-0.665	-0.348
is_renovated	4.07e+04	8671.889	4.693	0.000	2.37e+04	5.77e+04
grade_num	1.258e+05	2326.271	54.093	0.000	1.21e+05	1.3e+05
sqft_a/l	3.957e+05	2.66e+04	14.886	0.000	3.44e+05	4.48e+05
sqft_l/b	-395.8944	24.269	-16.313	0.000	-443.464	-348.325
baths_rnd	-6.186e+04	1.33e+04	-4.636	0.000	-8.8e+04	-3.57e+04

Omnibus: 15863.677 **Durbin-Watson:** 1.981

Prob(Omnibus): 0.000 **Jarque-Bera (JB):** 890732.662

Skew: 2.984 **Prob(JB):** 0.00

Kurtosis: 33.891 **Cond. No.** 4.51e+06

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 4.51e+06. This might indicate that there are strong multicollinearity or other numerical problems.

Good first test. Higher r2 than what I was anticipating. I think we might have some co-linear features impacting our score and helping our model along. I am concerned that our bathrooms feature has such a high p-value. Another high p-value is sqft_lot; we can remove this feature as it will not really impact our recommendations to the customers. Being able to add square footage to your lot isn't always the easiest. With that being said, I am also going to remove both of the _15 metrics as those are also variables that will not help our customers.

```
In [50]: # Removing sqft_living15 and sqft_lot15  
df_clean = df_clean.drop(['sqft_lot15', 'sqft_living15'], axis=1)
```

Model 2

```
In [51]: # Create target  
  
target2 = df_clean['price']  
  
# Create predictors  
predictors2 = df_clean.drop(['price'], axis=1)  
  
# Create model intercept  
predictors_int2 = sm.add_constant(predictors2)  
  
# Fit model to data  
model2 = sm.OLS(df_clean['price'], predictors_int2).fit()
```

```
In [52]: model2.summary()
```

```
Out[52]: OLS Regression Results  
Dep. Variable: price R-squared: 0.630  
Model: OLS Adj. R-squared: 0.629  
Method: Least Squares F-statistic: 2822.  
Date: Thu, 27 Oct 2022 Prob (F-statistic): 0.00  
Time: 20:57:03 Log-Likelihood: -2.9667e+05  
No. Observations: 21597 AIC: 5.934e+05
```

Df Residuals:	21583	BIC:	5.935e+05			
Df Model:	13					
Covariance Type:	nonrobust					
	coef	std err	t	P> t 	[0.025	0.975]
const	6.982e+06	1.35e+05	51.615	0.000	6.72e+06	7.25e+06
bedrooms	-1.183e+05	4769.815	-24.803	0.000	-1.28e+05	-1.09e+05
bathrooms	1.156e+05	1.35e+04	8.533	0.000	8.9e+04	1.42e+05
sqft_living	390.1597	22.397	17.420	0.000	346.260	434.059
sqft_lot	-0.2115	0.038	-5.581	0.000	-0.286	-0.137
floors	2.497e+04	3891.780	6.417	0.000	1.73e+04	3.26e+04
sqft_above	-120.2947	22.350	-5.382	0.000	-164.103	-76.486
sqft_basement	59.4677	19.897	2.989	0.003	20.467	98.468
yr_built	-4020.7860	68.742	-58.491	0.000	-4155.526	-3886.046
is_renovated	3.771e+04	8696.177	4.336	0.000	2.07e+04	5.48e+04
grade_num	1.332e+05	2233.473	59.658	0.000	1.29e+05	1.38e+05
sqft_a/l	3.869e+05	2.67e+04	14.513	0.000	3.35e+05	4.39e+05
sqft_l/b	-393.8194	24.337	-16.182	0.000	-441.521	-346.118
baths_rnd	-6.222e+04	1.34e+04	-4.650	0.000	-8.85e+04	-3.6e+04
Omnibus:	15487.015	Durbin-Watson:	1.980			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	806140.880			
Skew:	2.898	Prob(JB):	0.00			
Kurtosis:	32.364	Cond. No.	3.92e+06			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 3.92e+06. This might indicate that there are strong multicollinearity or other numerical problems.

A little bit worse, but pretty consistent - .658 vs .655 r2. I am going to remove the outliers from my dependent variable price, to see if we can improve upon the r2 value. Our p-values are looking good right now as well.

```
In [53]: # Removing outliers, and checking row count compared to what we had above =20439  
df_clean_out = df_clean[df_clean['price'].between(lower_bnd, upper_bnd)]  
  
len(df_clean_out)
```

Out[53]: 20439

Model 3

```
In [54]: # Quick check for colinearity  
  
df_clean_out.corr()
```

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	sqft_above	sqft_basement	yr_builtin	is_renovated
price	1.000000	0.300832	0.454959	0.621362	0.093554	0.273360	0.529931	0.230675	0.061264	0.0812
bedrooms	0.300832	1.000000	0.504697	0.601457	0.025619	0.160827	0.478118	0.281532	0.166782	0.0022
bathrooms	0.454959	0.504697	1.000000	0.716700	0.064993	0.502707	0.638466	0.216803	0.544843	0.0231
sqft_living	0.621362	0.601457	0.716700	1.000000	0.159058	0.340518	0.853404	0.369106	0.354009	0.0190
sqft_lot	0.093554	0.025619	0.064993	0.159058	1.000000	-0.017066	0.162765	0.011572	0.042480	0.0073
floors	0.273360	0.160827	0.502707	0.340518	-0.017066	1.000000	0.528746	-0.291962	0.511811	-0.0035
sqft_above	0.529931	0.478118	0.638466	0.853404	0.162765	0.528746	1.000000	-0.160702	0.462240	-0.0022
sqft_basement	0.230675	0.281532	0.216803	0.369106	0.011572	-0.291962	-0.160702	1.000000	-0.149189	0.0396
yr_builtin	0.061264	0.166782	0.544843	0.354009	0.042480	0.511811	0.462240	-0.149189	1.000000	-0.1974
is_renovated	0.081213	0.002276	0.023136	0.019059	0.007383	-0.003562	-0.002210	0.039616	-0.197410	1.00000
grade_num	0.630980	0.329053	0.613962	0.705223	0.094244	0.457583	0.710129	0.070187	0.493127	-0.0149
sqft_a/l	-0.139485	-0.194203	-0.115714	-0.206236	0.026472	0.359244	0.308500	-0.924708	0.201746	-0.0321
sqft_l/b	0.501539	-0.090419	0.448529	0.703416	0.176788	0.269895	0.612373	0.239705	0.268447	0.0344

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	sqft_above	sqft_basement	yr_builtin	is_renovated
baths_rnd	0.443049	0.491785	0.986264	0.704719	0.061620	0.521280	0.644027	0.184412	0.547011	0.02312

```
In [55]: # Looking at absolute values above .75 for removal
```

```
abs(df_clean_out.corr()) > 0.75
```

```
Out[55]:
```

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	sqft_above	sqft_basement	yr_builtin	is_renovated	grade_ni
price	True	False	False	False	False	False	False	False	False	False	Fa
bedrooms	False	True	False	False	False	False	False	False	False	False	Fa
bathrooms	False	False	True	False	False	False	False	False	False	False	Fa
sqft_living	False	False	False	True	False	False	True	False	False	False	Fa
sqft_lot	False	False	False	False	True	False	False	False	False	False	Fa
floors	False	False	False	False	False	True	False	False	False	False	Fa
sqft_above	False	False	False	True	False	False	True	False	False	False	Fa
sqft_basement	False	False	False	False	False	False	False	True	False	False	Fa
yr_builtin	False	False	False	False	False	False	False	False	True	False	Fa
is_renovated	False	False	False	False	False	False	False	False	False	True	Fa
grade_num	False	False	False	False	False	False	False	False	False	False	Tr
sqft_a/l	False	False	False	False	False	False	False	True	False	False	Fa
sqft_l/b	False	False	False	False	False	False	False	False	False	False	Fa
baths_rnd	False	False	True	False	False	False	False	False	False	False	Fa

```
In [56]: # Save absolute value of correlation matrix as a data frame
```

```
df = df_clean_out.corr().abs().stack().reset_index().sort_values(0, ascending=False)

# Zip the variable name columns
df['pairs'] = list(zip(df.level_0, df.level_1))

# Set index to pairs
df.set_index(['pairs'], inplace = True)
```

```

# Drop level columns
df.drop(columns=['level_1', 'level_0'], inplace = True)

# Rename correlation column as cc rather than 0
df.columns = ['cc']

# Drop duplicates
df.drop_duplicates(inplace=True)

```

In [57]: df[(df.cc >.75) & (df.cc <1)]

Out[57]:

	cc
	pairs
(bathrooms, baths_rnd)	0.986264
(sqft_basement, sqft_a/l)	0.924708
(sqft_above, sqft_living)	0.853404

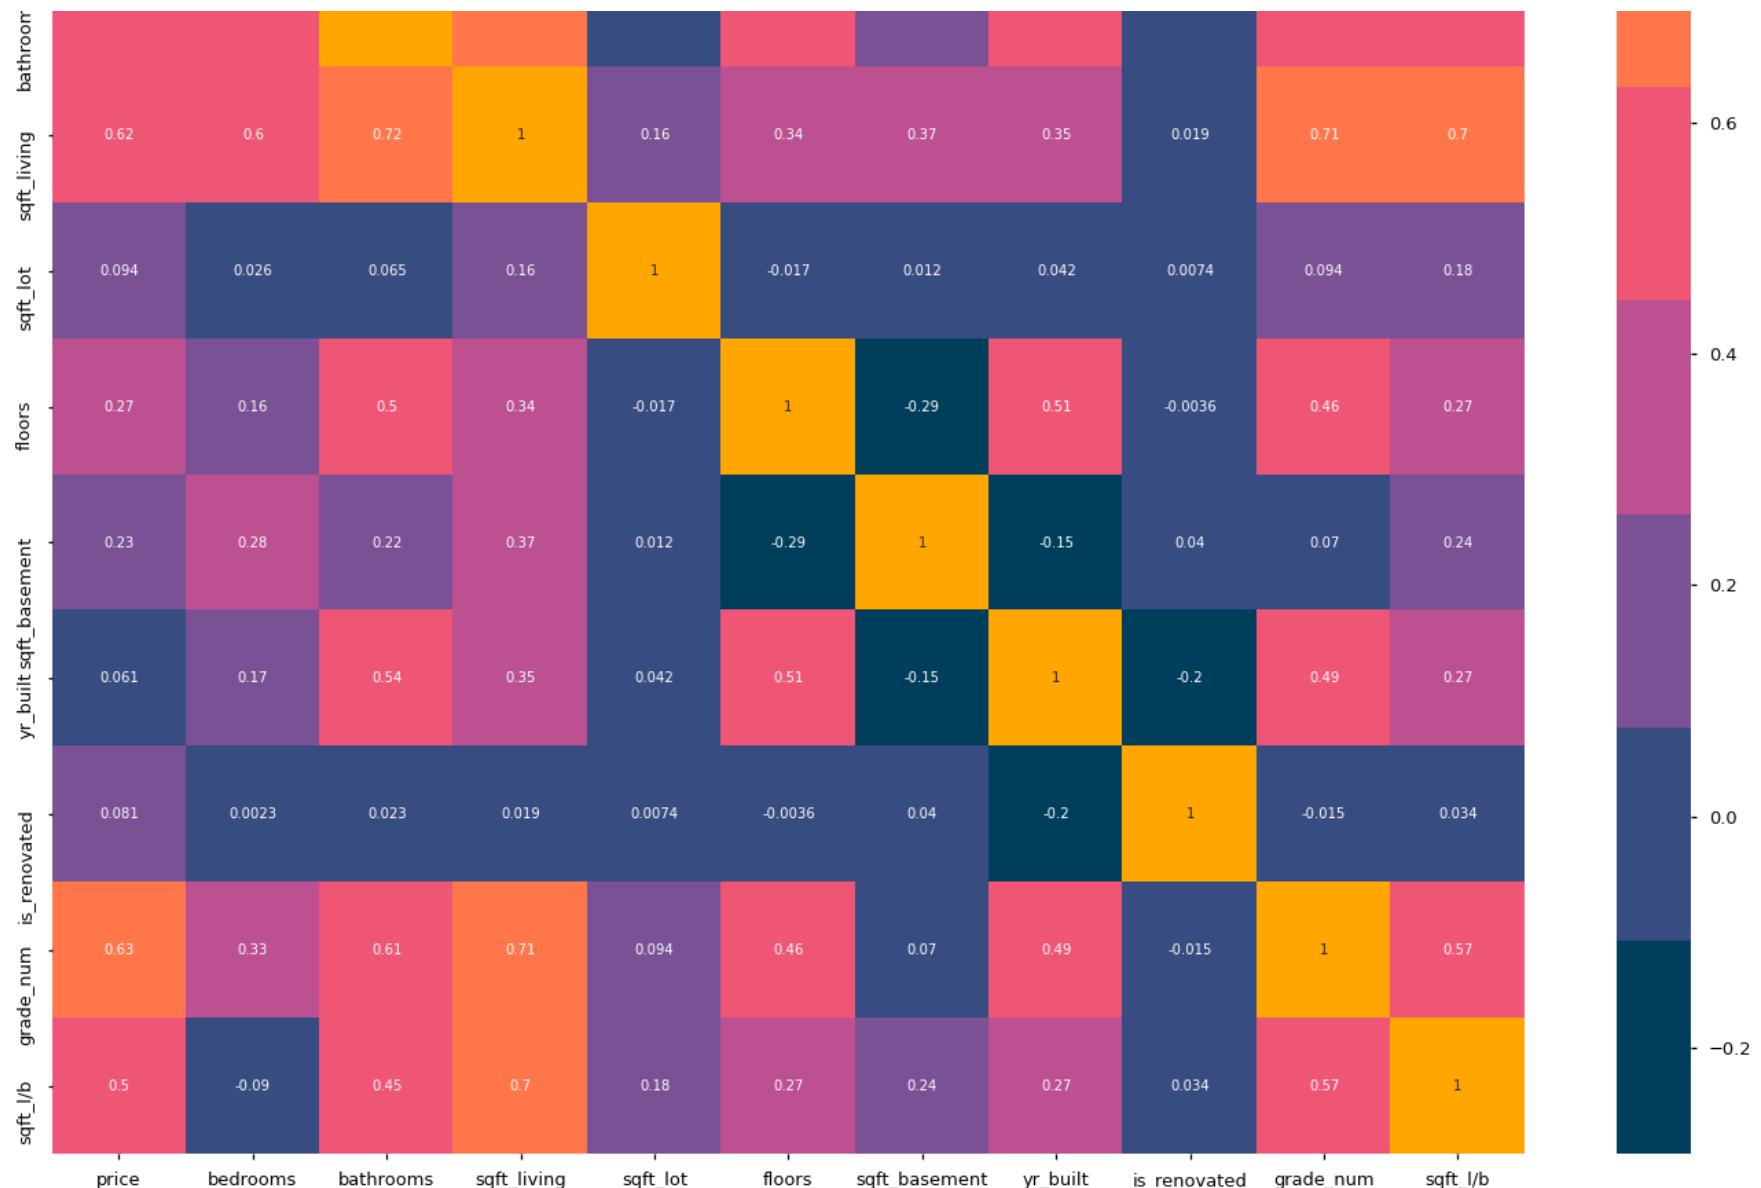
In [58]: final_df = df_clean_out.drop(['sqft_above',
 'sqft_a/l',
 'baths_rnd'], axis=1)

Out[58]: (20439, 11)

In [59]: plt.figure(figsize = (24,20))
corrM = final_df.corr()

sns.heatmap(corrM, annot=True, cmap=pal)
plt.show()





Looks like with the changes we have made we are still maintaining our relationships with our dependent variable. The baths_rnd, sqft_living, and grade_num are all consistent with our initial correlation matrix.

```
In [60]: # Model 3 includes same variables excluding price outliers
```

```
# create target
target3 = final_df['price']
```

```

# create predictors
predictors3 = final_df.drop(['price'], axis=1)

# create model intercept
predictors_int3 = sm.add_constant(predictors3)

# fit model to data
model3 = sm.OLS(final_df['price'], predictors_int3).fit()

```

In [61]: `model3.summary()`

Out[61]: OLS Regression Results

Dep. Variable:	price	R-squared:	0.560			
Model:	OLS	Adj. R-squared:	0.560			
Method:	Least Squares	F-statistic:	2602.			
Date:	Thu, 27 Oct 2022	Prob (F-statistic):	0.00			
Time:	20:57:05	Log-Likelihood:	-2.7087e+05			
No. Observations:	20439	AIC:	5.418e+05			
Df Residuals:	20428	BIC:	5.418e+05			
Df Model:	10					
Covariance Type:	nonrobust					
	coef	std err	t	P> t 	[0.025	0.975]
const	4.975e+06	8.72e+04	57.079	0.000	4.8e+06	5.15e+06
bedrooms	-1.469e+04	3256.618	-4.512	0.000	-2.11e+04	-8309.609
bathrooms	2.911e+04	2354.778	12.364	0.000	2.45e+04	3.37e+04
sqft_living	80.3567	5.431	14.797	0.000	69.712	91.001
sqft_lot	0.0436	0.025	1.764	0.078	-0.005	0.092
floors	4.258e+04	2471.739	17.229	0.000	3.77e+04	4.74e+04
sqft_basement	25.7348	3.052	8.432	0.000	19.752	31.717
yr_builtin	-2799.9482	45.365	-61.720	0.000	-2888.867	-2711.029
is_renovated	1536.3157	5810.485	0.264	0.791	-9852.700	1.29e+04

grade_num	1.024e+05	1472.237	69.522	0.000	9.95e+04	1.05e+05	
sqft_l/b	15.7406	16.702	0.942	0.346	-16.996	48.477	
Omnibus: 902.165 Durbin-Watson: 1.967							
Prob(Omnibus):	0.000	Jarque-Bera (JB):	1277.633				
Skew:	0.427	Prob(JB):	3.68e-278				
Kurtosis:	3.877	Cond. No.	3.85e+06				

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 3.85e+06. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [62]: # Removing columns with > .05 p values and going to rerun model 3
```

```
final_df = final_df.drop(['sqft_l/b',
                         'sqft_lot'],
                        axis=1)
```

```
final_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 20439 entries, 0 to 21596
Data columns (total 9 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   price              20439 non-null   float64
 1   bedrooms           20439 non-null   int64  
 2   bathrooms          20439 non-null   float64
 3   sqft_living        20439 non-null   int64  
 4   floors              20439 non-null   float64
 5   sqft_basement      20439 non-null   float64
 6   yr_built            20439 non-null   int64  
 7   is_renovated       20439 non-null   int64  
 8   grade_num           20439 non-null   int64  
dtypes: float64(4), int64(5)
memory usage: 1.6 MB
```

Model 3b

In [63]:

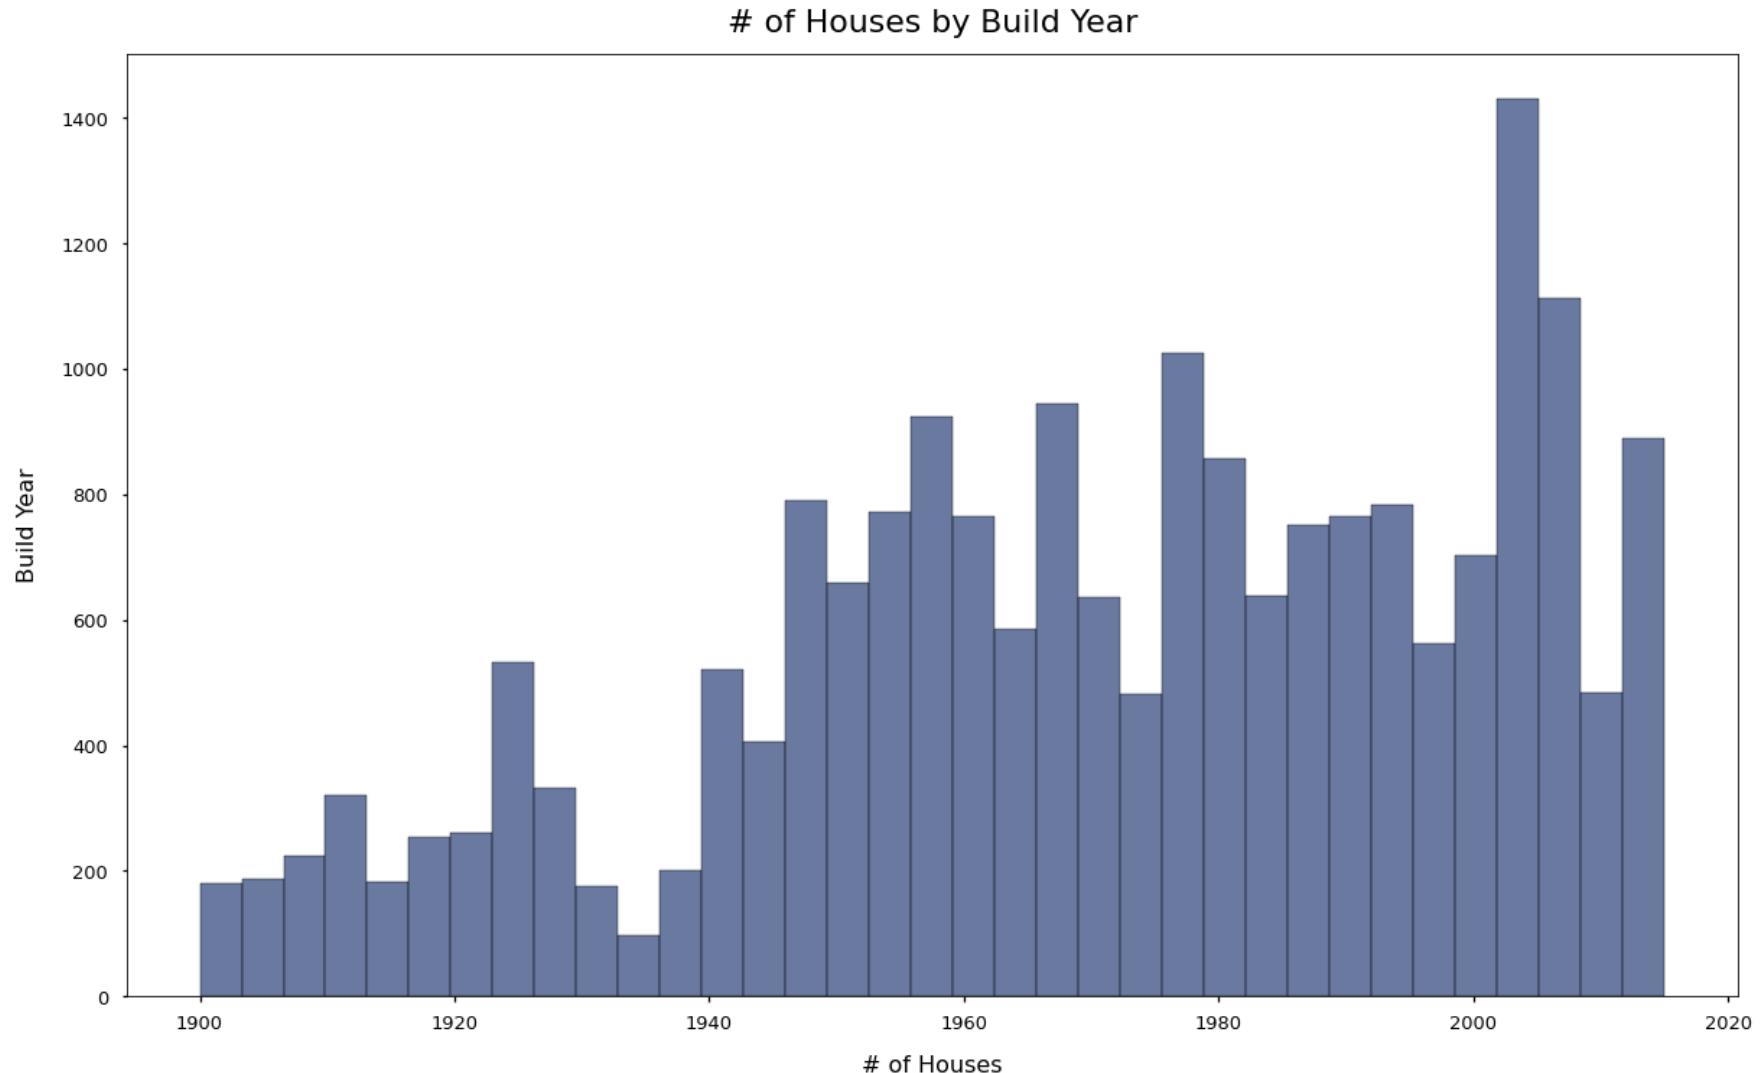
```
# taking a look at our frame once again for final summary stats in presentation.  
# Still not satisfied with r2, going to look at yr_built to see if there are any  
# additional insights to remove potential outliers.  
  
final_df.describe()
```

Out[63]:

	price	bedrooms	bathrooms	sqft_living	floors	sqft_basement	yr_built	is_renovated	grade
count	2.043900e+04	20439.000000	20439.000000	20439.000000	20439.000000	20439.000000	20439.000000	20439.000000	20439.000000
mean	4.767026e+05	3.329468	2.051886	1975.189931	1.475561	261.425950	1970.805079	0.030579	8.0
std	2.077903e+05	0.884165	0.710196	773.797924	0.536629	406.064503	29.159010	0.172178	8.0
min	7.800000e+04	1.000000	0.500000	370.000000	1.000000	0.000000	1900.000000	0.000000	8.0
25%	3.150000e+05	3.000000	1.500000	1400.000000	1.000000	0.000000	1951.000000	0.000000	8.0
50%	4.375000e+05	3.000000	2.000000	1860.000000	1.000000	0.000000	1974.000000	0.000000	8.0
75%	6.000000e+05	4.000000	2.500000	2430.000000	2.000000	500.000000	1996.000000	0.000000	8.0
max	1.120000e+06	11.000000	7.500000	7480.000000	3.500000	2720.000000	2015.000000	1.000000	11.0

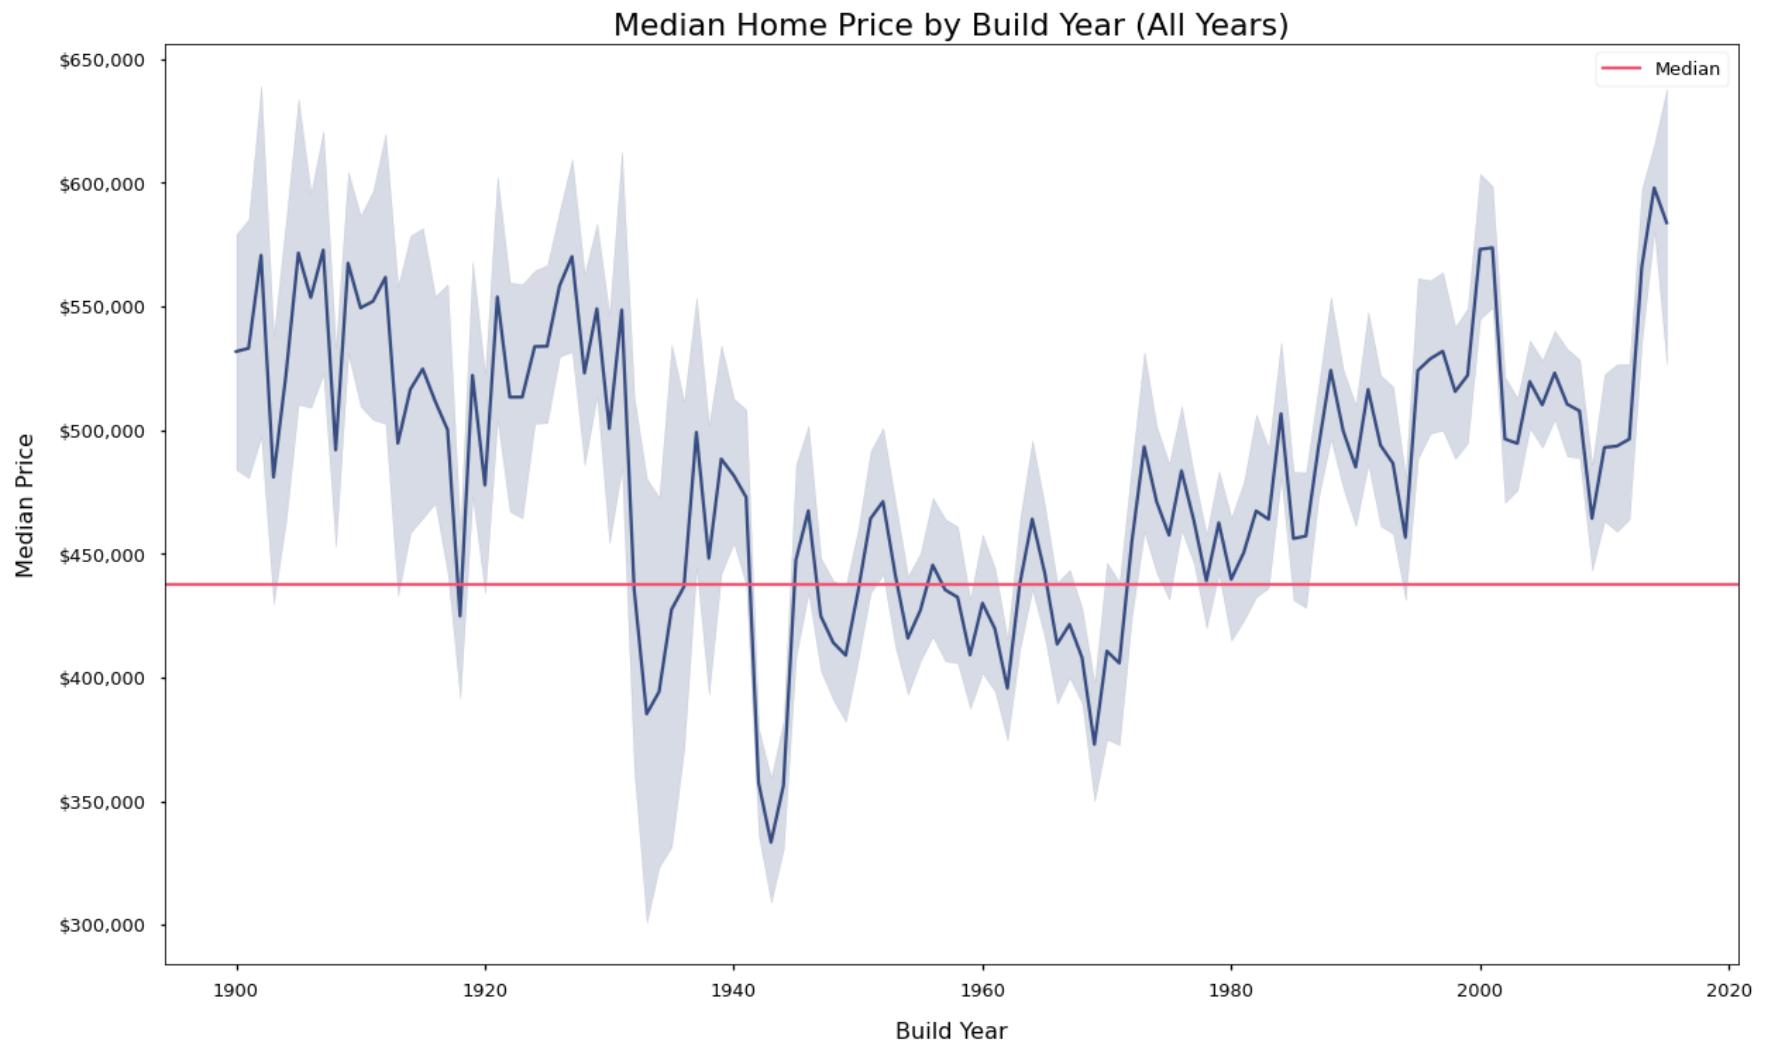
In [64]:

```
# Taking a look at the number of houses by build year.  
  
fig, ax = plt.subplots(figsize = (20,12))  
  
sns.histplot(data=final_df, x='yr_built', color = c("indigo"))  
  
ax.set_title('# of Houses by Build Year', pad = 15, fontsize = 22)  
ax.set_xlabel('# of Houses', labelpad = 15, fontsize = 16)  
ax.set_ylabel('Build Year', labelpad = 15, fontsize = 16)  
plt.show();
```



```
In [65]: # Now that we know the value counts by year, I am going to take a look at the median  
# values by year to look how that is changing over time. The tricky part here is that we don't  
# know necessarily when the price value is represented.  
  
fig, ax = plt.subplots(figsize = (20,12))  
  
sns.lineplot(x='yr_built',  
             y='price',  
             data=final_df,  
             ax=ax,
```

```
    color=c('indigo'))\n\nplt.axhline(final_df['price'].median(), color = c("peach"), label = 'Median')\n\nfmt = '${x:.0f}'\ntick = mtick.StrMethodFormatter(fmt)\nax.yaxis.set_major_formatter(tick)\n\nax.set_title('Median Home Price by Build Year (All Years)', fontsize = 22)\nax.set_xlabel('Build Year', labelpad = 15, fontsize = 16)\nax.set_ylabel('Median Price', labelpad = 15, fontsize = 16)\nplt.legend()\n\nplt.show();
```



What is interesting is that median values are similar from years 1900 - 1935ish, and from 1990 to the end of our dataset in 2015. Going to look at removing the values before 1990 to see if we can increase the relationship between price and yr_built, without sacrificing any sort of colinear relationship. There is probably also something to be said around how houses were constructed in these different time periods.

```
In [66]: # Looks like price bounces around quite a bit when comparing records, and yr_built
# Going to adjust my dataframe to include more recent housing prices

final_df_recent = final_df.loc[final_df['yr_built'] >= 1990]
```

```
In [67]: final_df_recent.describe()
```

```
Out[67]:
```

	price	bedrooms	bathrooms	sqft_living	floors	sqft_basement	yr_builtin	is_renovated	grade_num
count	6.463000e+03	6463.000000	6463.000000	6463.000000	6463.000000	6463.000000	6463.000000	6463.000000	6463.000000
mean	5.187825e+05	3.449636	2.565720	2321.461396	1.992109	121.717314	2002.830419	0.000464	8.183197
std	2.054968e+05	0.798226	0.510816	820.258728	0.432650	292.102191	7.017290	0.021542	0.985589
min	1.540000e+05	1.000000	0.500000	550.000000	1.000000	0.000000	1990.000000	0.000000	5.000000
25%	3.500000e+05	3.000000	2.500000	1670.000000	2.000000	0.000000	1997.000000	0.000000	7.000000
50%	4.750000e+05	3.000000	2.500000	2240.000000	2.000000	0.000000	2004.000000	0.000000	8.000000
75%	6.525500e+05	4.000000	2.750000	2860.000000	2.000000	0.000000	2008.000000	0.000000	9.000000
max	1.120000e+06	9.000000	7.500000	7350.000000	3.500000	2600.000000	2015.000000	1.000000	12.000000

```
In [68]: # Creating a data frame for a chart below
```

```
yr_chrt = (
    final_df_recent
    .groupby('yr_builtin')['price']
    .median()
    .reset_index()
)
```

```
In [69]: # creating visual for presentation. Showing median home values over time for
# reference to show that they really don't fluctuate much above or below the median
# for the entire time frame.
```

```
fig, ax = plt.subplots(figsize = (20,12))

sns.barplot(x='yr_builtin',
            y='price',
            data=yr_chrt,
            color = c("indigo"),
            ax=ax)

plt.axhline(final_df_recent['price'].median(), color = c("peach"), label = 'Median')

fmt = '${x:,.0f}'
```

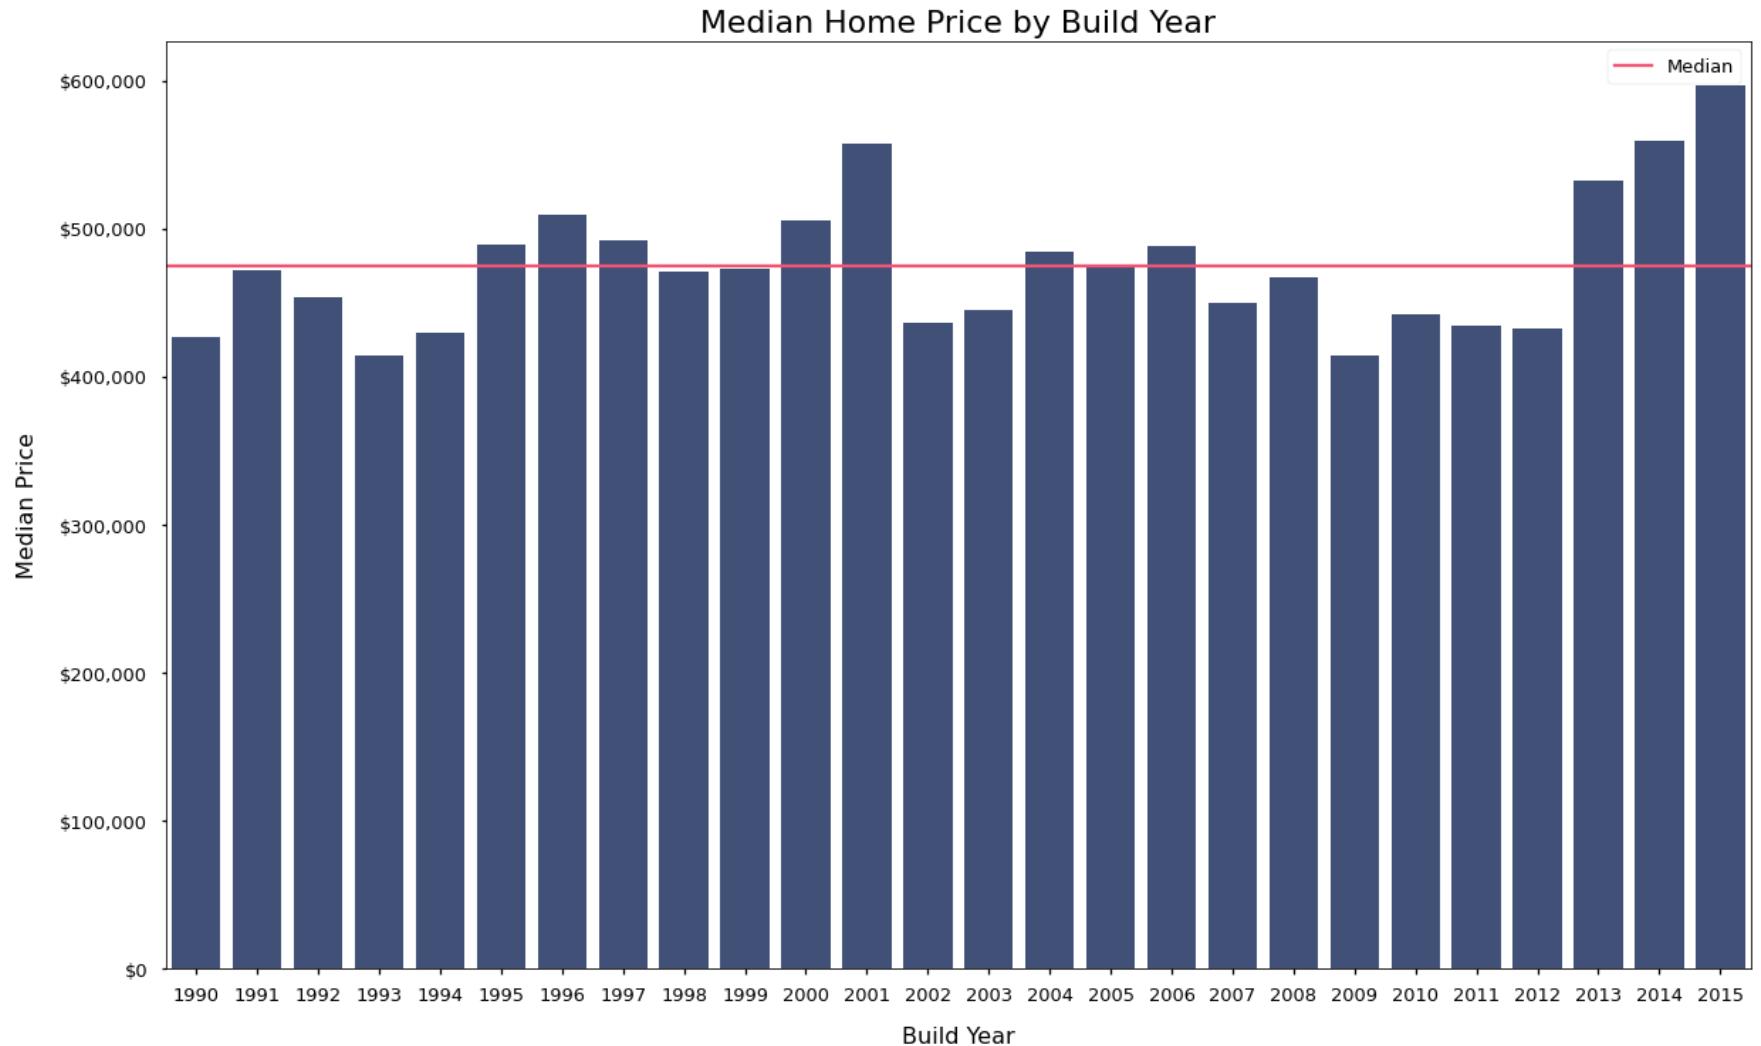
```

tick = mtick.StrMethodFormatter(fmt)
ax.yaxis.set_major_formatter(tick)

ax.set_title('Median Home Price by Build Year', fontsize = 22)
ax.set_xlabel('Build Year', labelpad = 15, fontsize = 16)
ax.set_ylabel('Median Price', labelpad = 15, fontsize = 16)
plt.legend()

plt.show();

```



In [70]: # Taking a look at our median home prices by build year

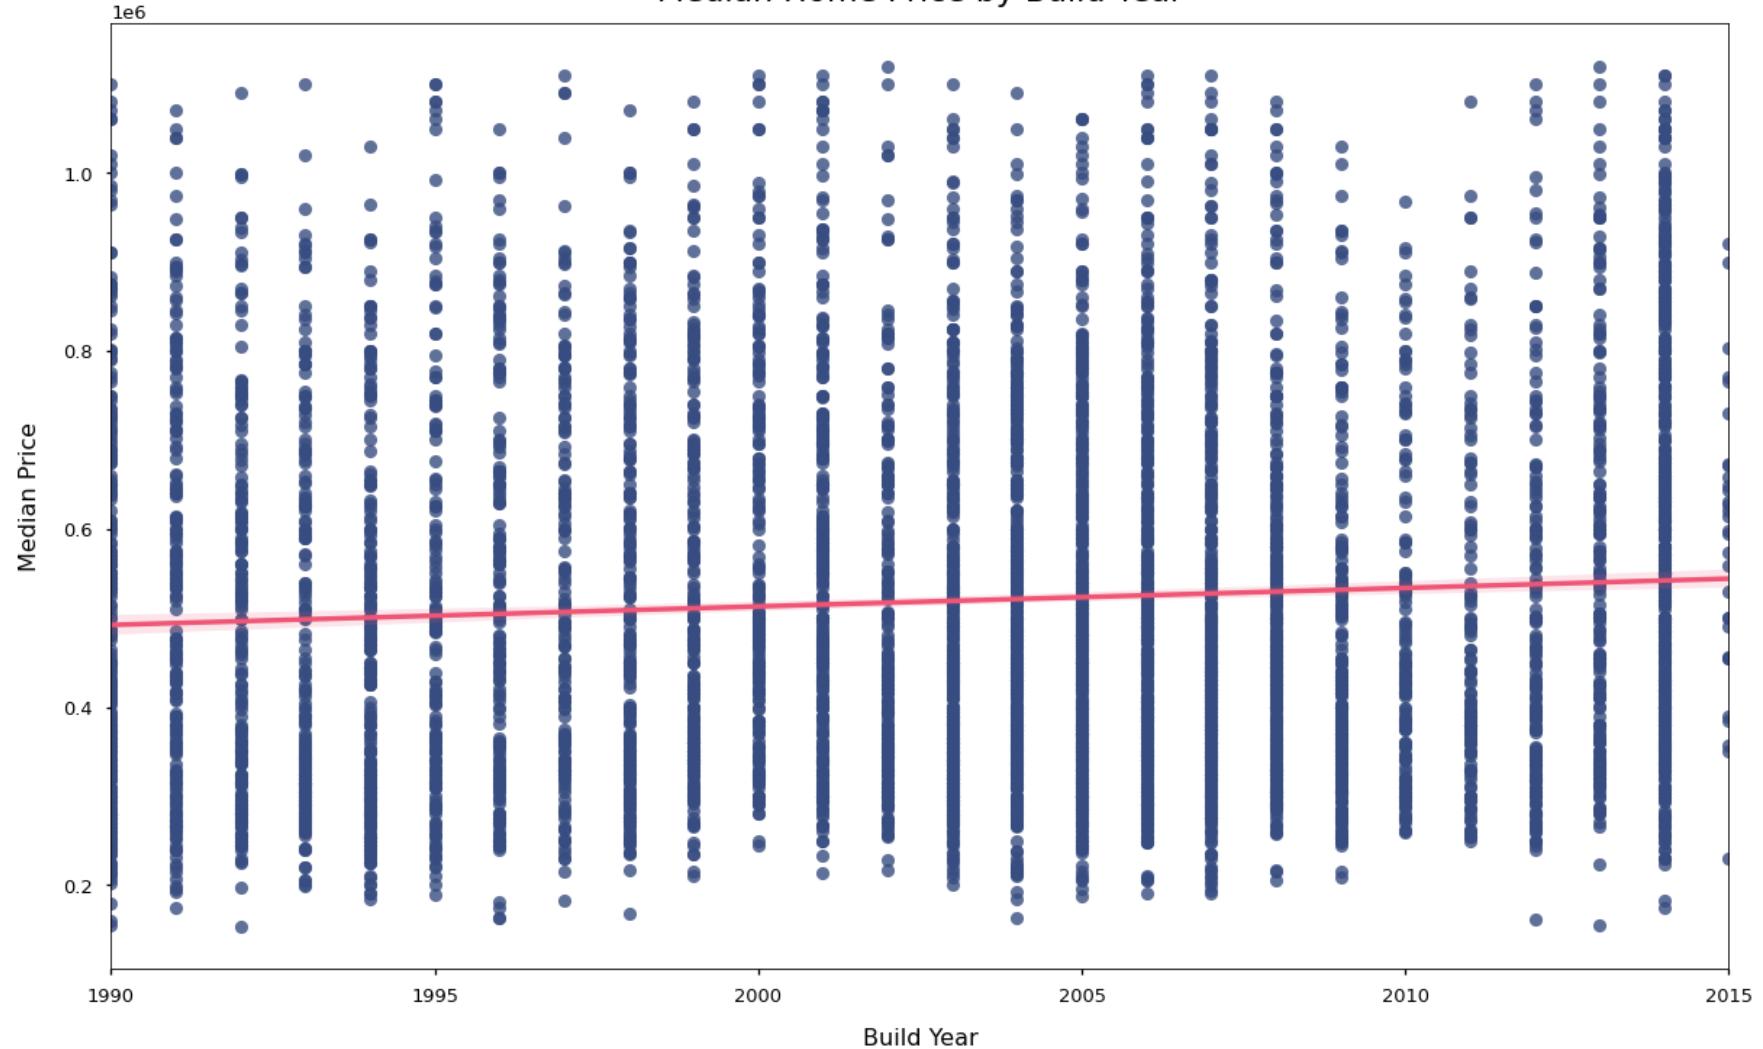
```
fig, ax = plt.subplots(figsize = (20,12))

sns.regplot(x='yr_built',
             y='price',
             data=final_df_recent,
             scatter_kws={"color": c("indigo")},
             line_kws={"color": c("peach")})

ax.set_title('Median Home Price by Build Year', pad = 15, fontsize = 22)
ax.set_xlabel('Build Year', labelpad = 15, fontsize = 16)
ax.set_ylabel('Median Price', labelpad = 15, fontsize = 16)

plt.show();
```

Median Home Price by Build Year



```
In [71]: # Looking at the relationship between price and square foot living. This will
# be a good visual for our presentation.

fig, ax = plt.subplots(figsize = (20,12))

sns.regplot(x='sqft_living',
            y='price',
            data=final_df_recent,
            scatter_kws={"color": c("indigo")},
            line_kws={"color": c("peach")})
```

```
fmt = '${x:.0f}'
tick = mtick.StrMethodFormatter(fmt)
ax.yaxis.set_major_formatter(tick)

ax.set_title('Price vs. Square Foot Living', pad = 15, fontsize = 22)
ax.set_xlabel('Square Foot Living', labelpad = 15, fontsize = 16)
ax.set_ylabel('Price', labelpad = 15, fontsize = 16)
plt.show();
```



In [74]:

```
baths_chrt = (
    final_df_recent
```

```
.groupby('bathrooms')['price']
    .median()
    .reset_index()
)
```

In [75]: # Quick peek

```
baths_chrt
```

Out[75]:

	bathrooms	price
0	0.50	312500.0
1	0.75	363550.5
2	1.00	318250.0
3	1.25	528250.0
4	1.50	358500.0
5	1.75	358500.0
6	2.00	369725.0
7	2.25	400000.0
8	2.50	461000.0
9	2.75	599000.0
10	3.00	510000.0
11	3.25	650000.0
12	3.50	712500.0
13	3.75	912000.0
14	4.00	816675.0
15	4.25	858450.0
16	4.50	932808.0
17	4.75	1034495.0
18	5.00	492500.0
19	5.25	1060000.0

	bathrooms	price
20	5.50	925000.0
21	6.00	900000.0
22	7.50	450000.0

```
In [76]: # Visual for presentation -- visualizing data from above data frame.

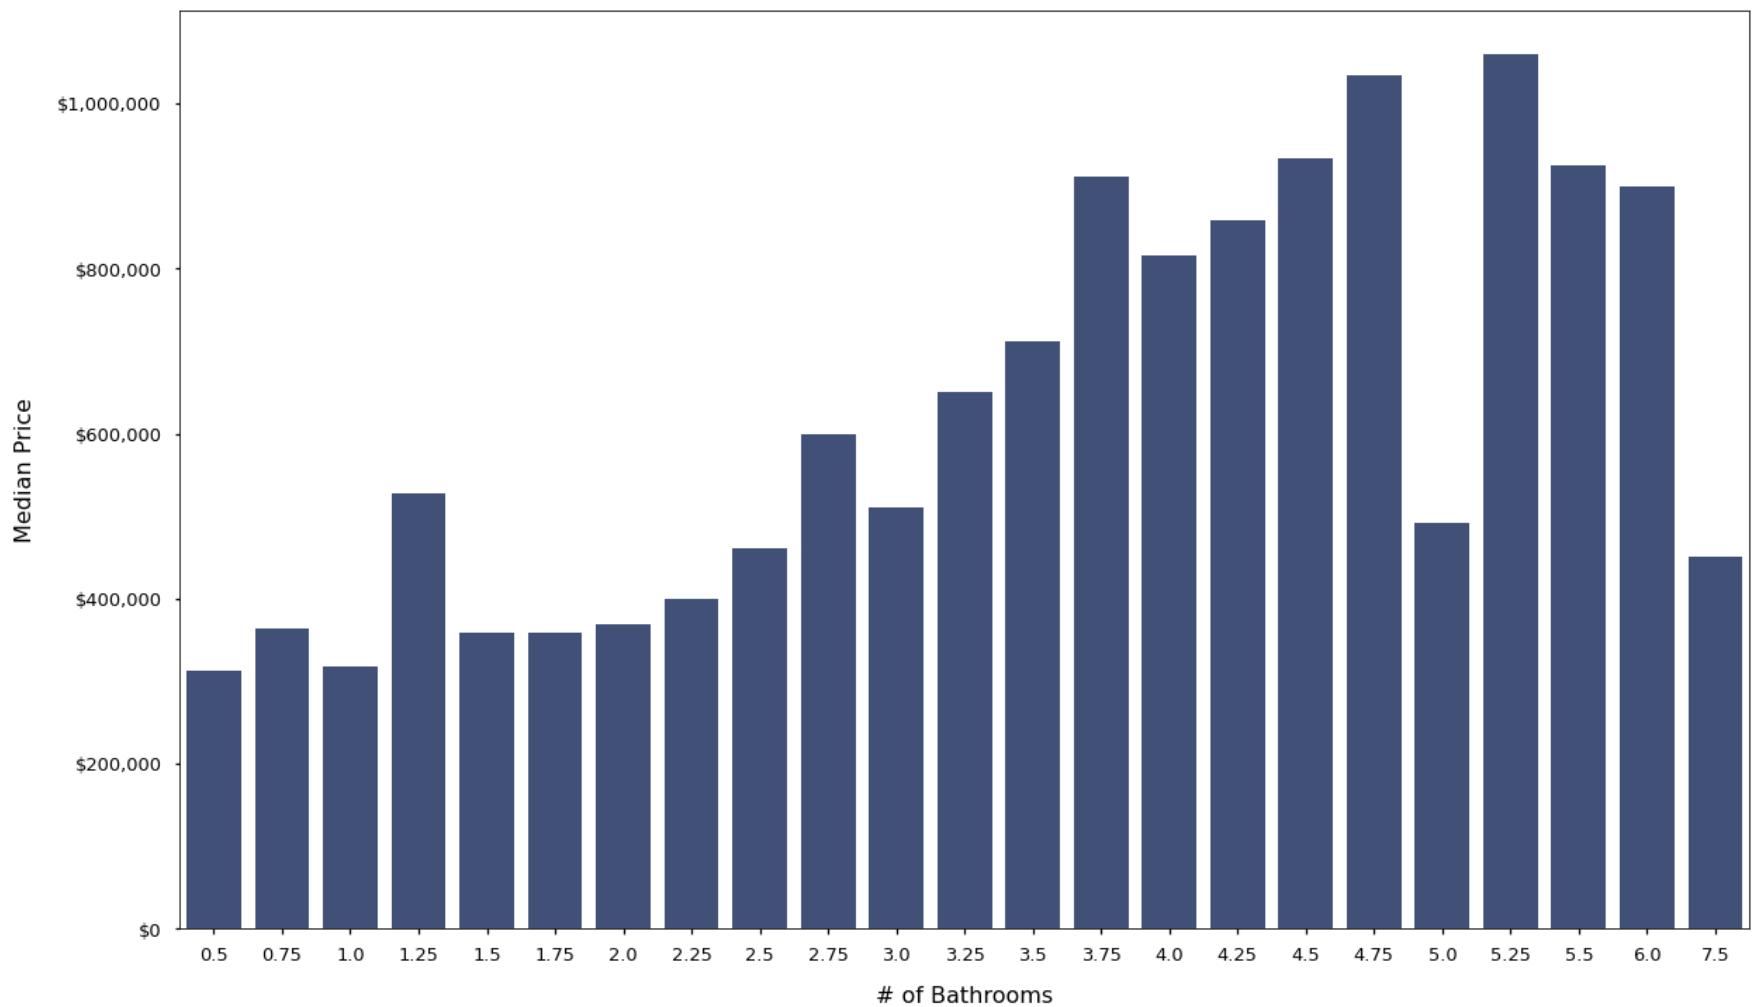
fig, ax = plt.subplots(figsize = (20,12))

sns.barplot(x='bathrooms',
             y='price',
             data=baths_chrt,
             color=c('indigo')
            )

fmt = '${x:,.0f}'
tick = mtick.StrMethodFormatter(fmt)
ax.yaxis.set_major_formatter(tick)

ax.set_title('Median Price vs. Bathrooms', pad = 15, fontsize = 22)
ax.set_xlabel('# of Bathrooms', labelpad = 15, fontsize = 16)
ax.set_ylabel('Median Price', labelpad = 15, fontsize = 16)
plt.show();
```

Median Price vs. Bathrooms



```
In [78]: # Model 3b - removed outliers, combined with a round of feature removal, and final dataset of  
# > 1990 values.  
  
# create target  
target3b = final_df_recent['price']  
  
# create predictors  
predictors3b = final_df_recent.drop(['price'], axis=1)  
  
# create model intercept  
predictors_int3b = sm.add_constant(predictors3b)
```

```
# fit model to data
model3b = sm.OLS(final_df_recent['price'],predictors_int3b).fit()
```

```
In [79]: model3b.summary()
```

```
Out[79]: OLS Regression Results
```

Dep. Variable:	price	R-squared:	0.616			
Model:	OLS	Adj. R-squared:	0.615			
Method:	Least Squares	F-statistic:	1293.			
Date:	Thu, 27 Oct 2022	Prob (F-statistic):	0.00			
Time:	21:00:02	Log-Likelihood:	-85142.			
No. Observations:	6463	AIC:	1.703e+05			
Df Residuals:	6454	BIC:	1.704e+05			
Df Model:	8					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
const	-4.055e+06	4.68e+05	-8.673	0.000	-4.97e+06	-3.14e+06
bedrooms	-3.222e+04	2744.816	-11.737	0.000	-3.76e+04	-2.68e+04
bathrooms	2.136e+04	4251.740	5.024	0.000	1.3e+04	2.97e+04
sqft_living	101.2899	3.516	28.808	0.000	94.397	108.183
floors	2.723e+04	3984.626	6.833	0.000	1.94e+04	3.5e+04
sqft_basement	15.5551	5.993	2.595	0.009	3.806	27.304
yr_built	1763.6396	234.373	7.525	0.000	1304.191	2223.088
is_renovated	6.755e+04	7.37e+04	0.917	0.359	-7.68e+04	2.12e+05
grade_num	9.862e+04	2234.949	44.126	0.000	9.42e+04	1.03e+05
Omnibus:	165.038	Durbin-Watson:	1.987			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	230.519			
Skew:	0.287	Prob(JB):	8.78e-51			

Kurtosis: 3.726

Cond. No. 9.24e+05

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 9.24e+05. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [80]: final_df_recent = final_df_recent.drop(['is_renovated'], axis=1)
```

```
In [81]: # Model 3b - removed outliers, combined with features with p values > .05. What I am
# assuming should be our final model. Finally removed is_renovated

# create target
target3b = final_df_recent['price']

# create predictors
predictors3b = final_df_recent.drop(['price'], axis=1)

# create model intercept
predictors_int3b = sm.add_constant(predictors3b)

# fit model to data
model3b = sm.OLS(target3b,predictors_int3b).fit()
```

```
In [82]: model3b.summary()
```

```
Out[82]: OLS Regression Results
```

Dep. Variable:	price	R-squared:	0.616
Model:	OLS	Adj. R-squared:	0.615
Method:	Least Squares	F-statistic:	1478.
Date:	Thu, 27 Oct 2022	Prob (F-statistic):	0.00
Time:	21:00:04	Log-Likelihood:	-85142.
No. Observations:	6463	AIC:	1.703e+05
Df Residuals:	6455	BIC:	1.704e+05
Df Model:	7		

Covariance Type: nonrobust

	coef	std err	t	P> t	[0.025	0.975]
const	-4.046e+06	4.67e+05	-8.655	0.000	-4.96e+06	-3.13e+06
bedrooms	-3.222e+04	2744.767	-11.740	0.000	-3.76e+04	-2.68e+04
bathrooms	2.142e+04	4251.229	5.039	0.000	1.31e+04	2.98e+04
sqft_living	101.3310	3.516	28.822	0.000	94.439	108.223
floors	2.73e+04	3983.875	6.851	0.000	1.95e+04	3.51e+04
sqft_basement	15.4769	5.993	2.583	0.010	3.729	27.224
yr_built	1758.7396	234.309	7.506	0.000	1299.416	2218.063
grade_num	9.86e+04	2234.853	44.121	0.000	9.42e+04	1.03e+05
Omnibus: 165.612		Durbin-Watson: 1.986				
Prob(Omnibus): 0.000		Jarque-Bera (JB): 231.242				
Skew: 0.288		Prob(JB): 6.12e-51				
Kurtosis: 3.726		Cond. No. 9.24e+05				

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 9.24e+05. This might indicate that there are strong multicollinearity or other numerical problems.

Here we go -- 3b is our model. It seems to have the highest r2, and our p-value meet the requirements of our testing. Even though the r2 isn't super high, I think we still know enough to make recommendations to our customers on which features within a home will most contribute to the Price. I will be adding some limitations to my presentation as well as at the end of the presentation of the dos and do nots of this final output.

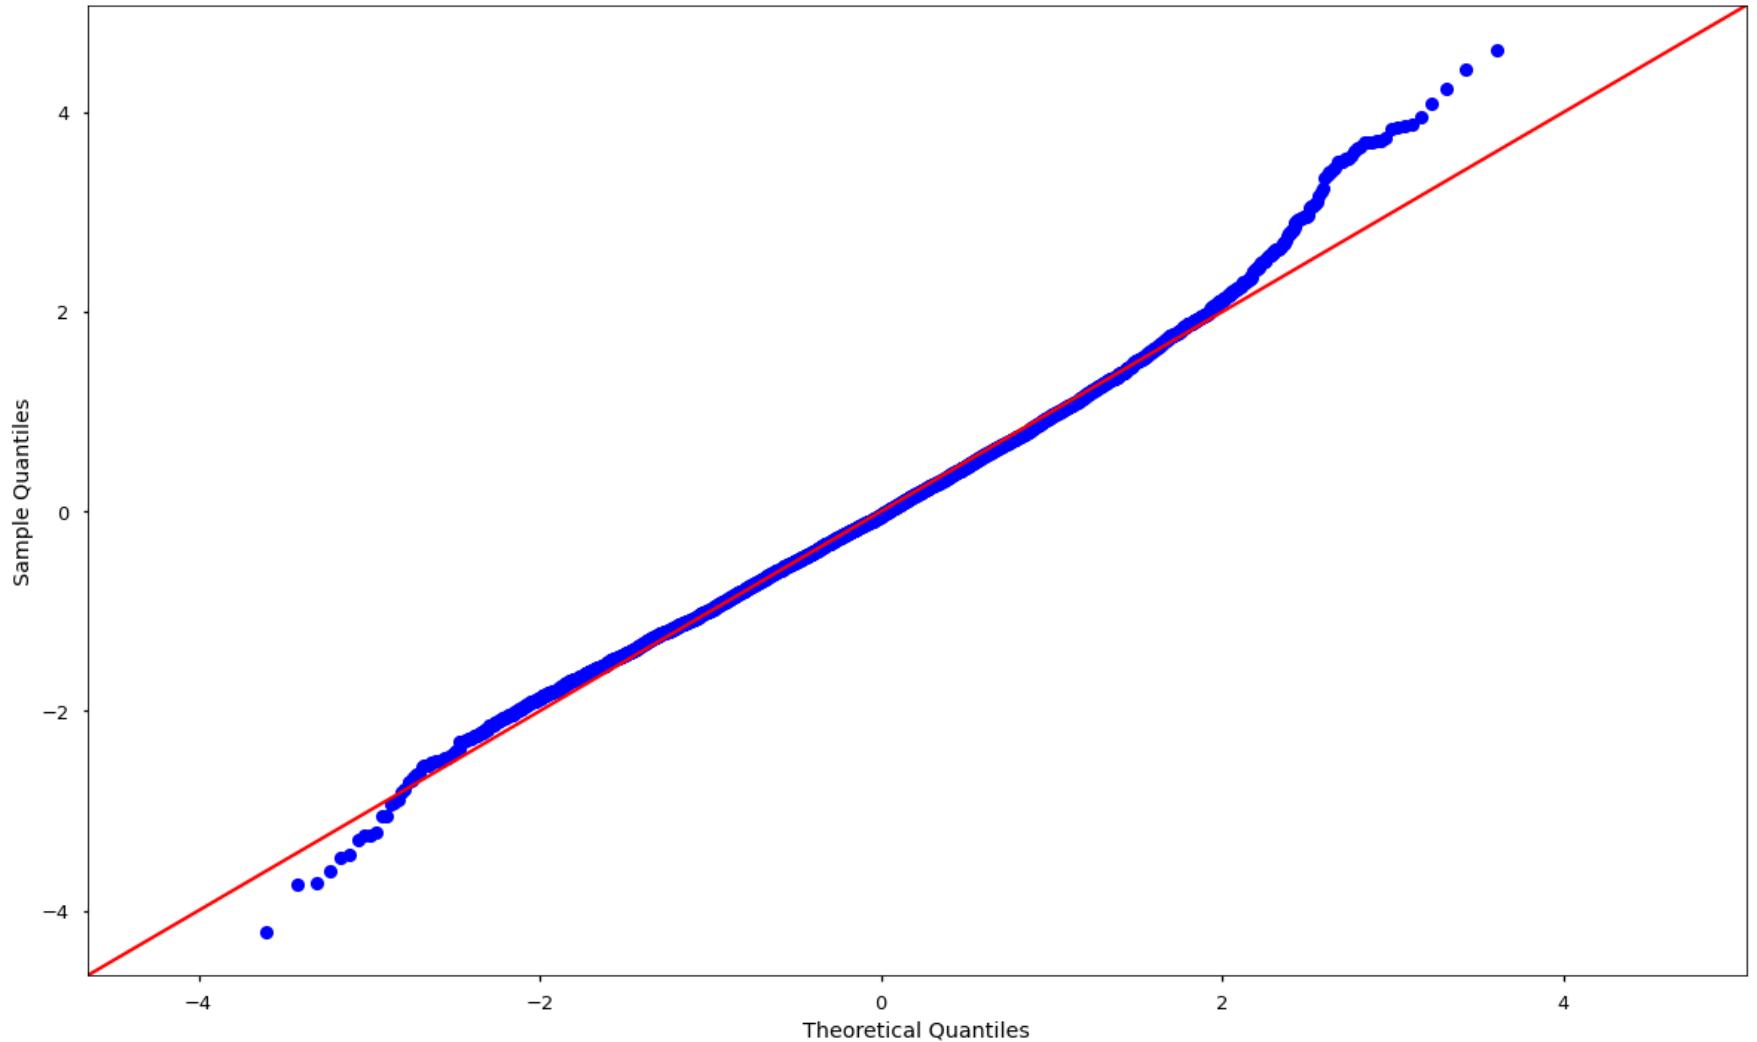
Insights for presentation

- bathrooms increase the value of the home by \$21,420
- sqft_living increases the value of a home by \$101 per square foot

- increasing the grade of your home by 1, increases the value by \$99k
- bedrooms decrease the value of a home by \$32,220
- adding floors to your home adds \$ 27,300 per floor
- newer homes generally get more money than the older in this dataset. However, this is not an impactable feature.

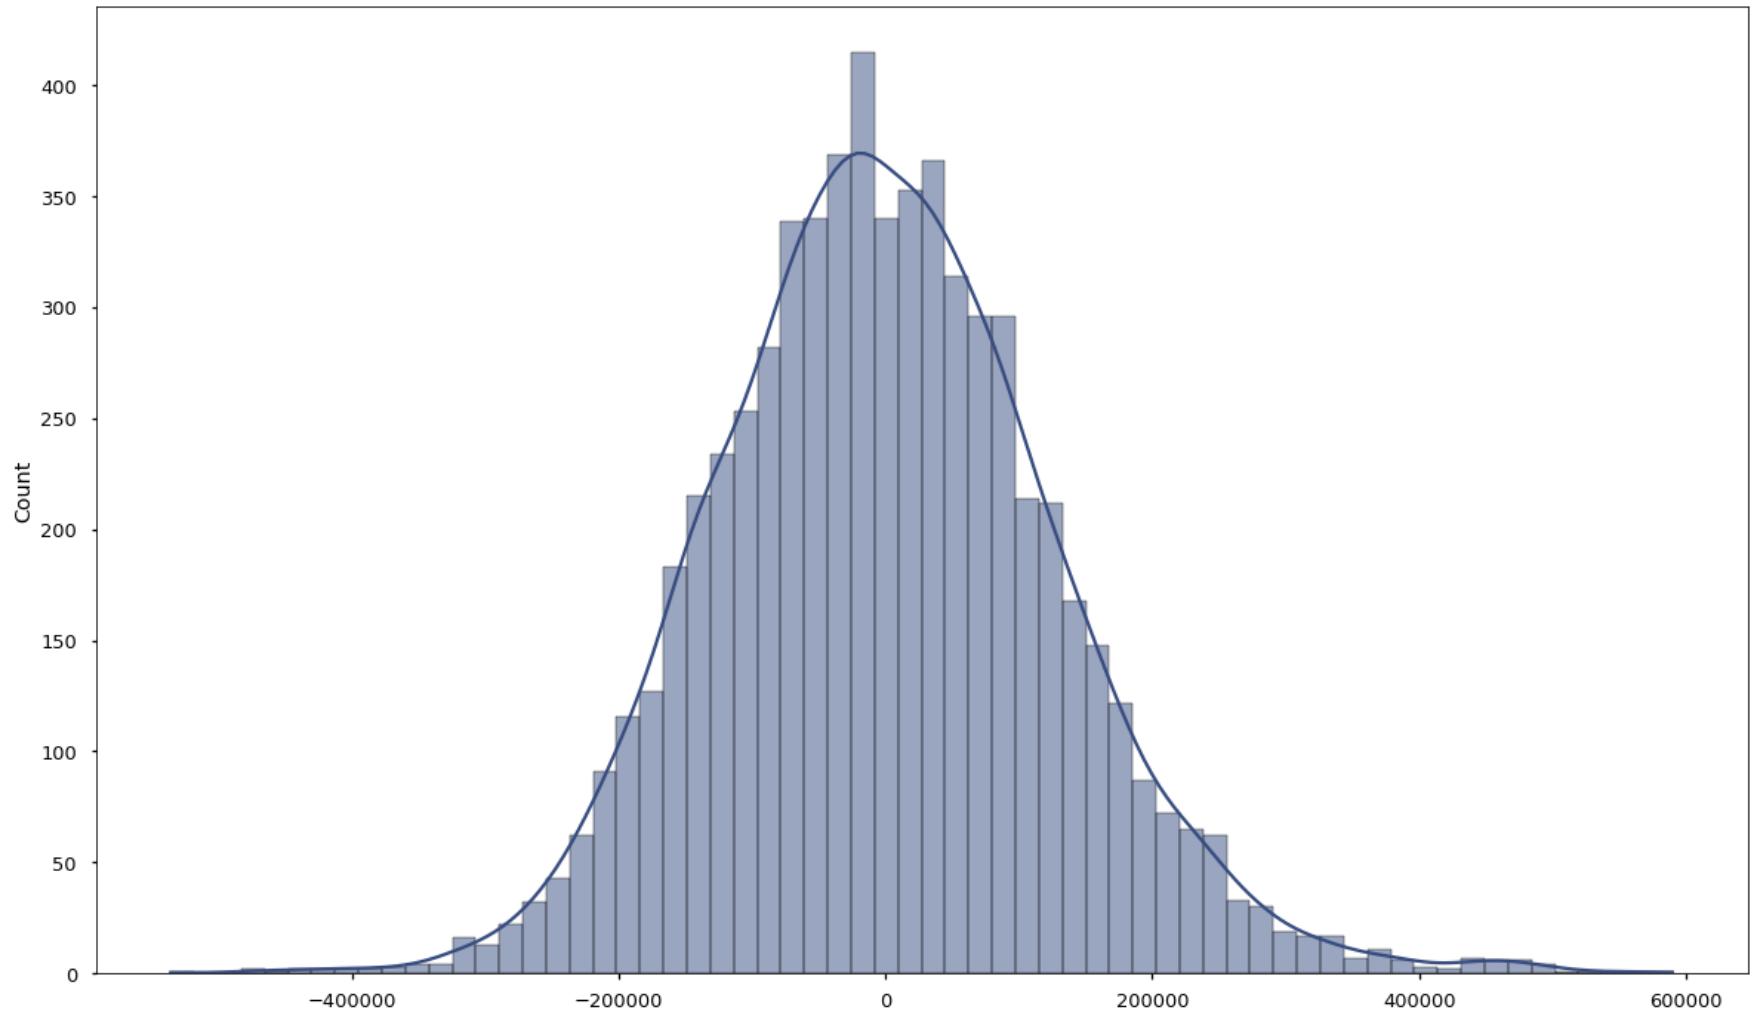
Checking Linear Regression Assumptions

```
In [83]: # Plotting residuals on qq plot to understand linearity  
  
fig, ax = plt.subplots(figsize = (20,12))  
  
res = model3b.resid # residuals  
  
sm.qqplot(res, fit=True, line='45', ax=ax);
```



```
In [84]: # Looking at the overall distribution of residuals  
fig, ax = plt.subplots(figsize = (20,12))  
sns.histplot(data=res, kde=True, color=c("indigo"))
```

```
Out[84]: <AxesSubplot:ylabel='Count'>
```



```
In [85]: # fitting our model using ski-kit learn here to do create some predictions  
lr.fit(predictors3b, target3b)
```

```
Out[85]: LinearRegression()
```

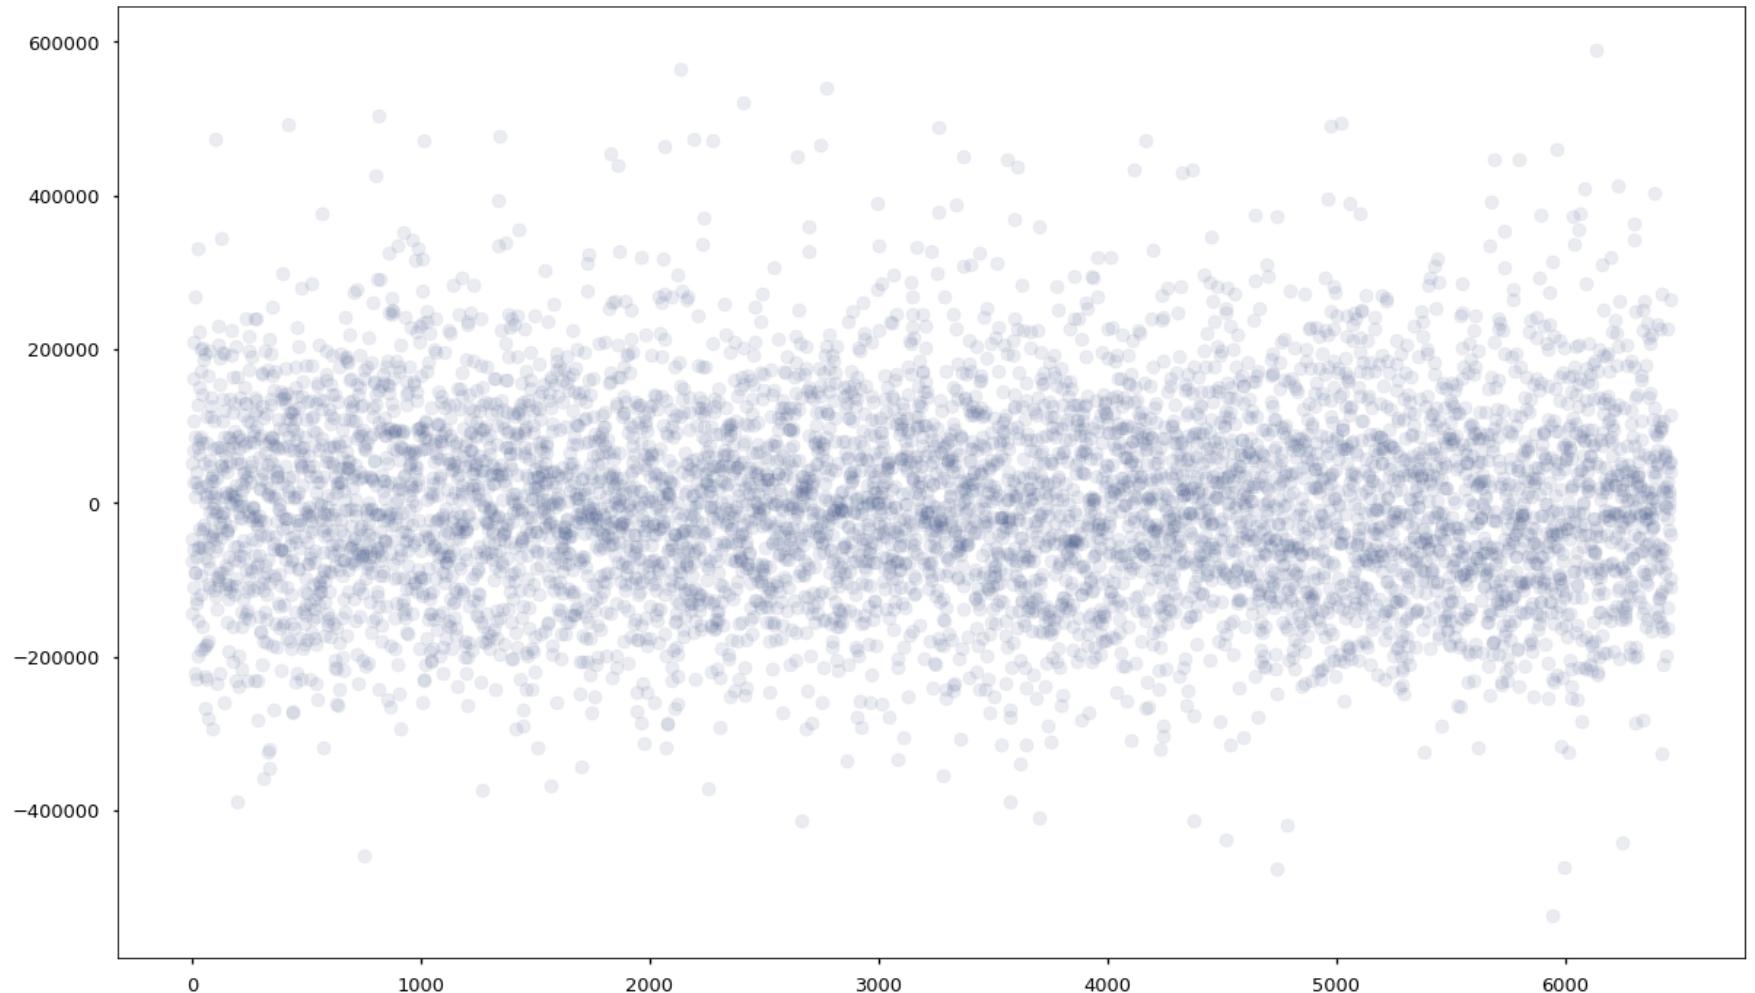
```
In [86]: # Creating a quick prediction and taking a look at the array  
y_hat = lr.predict(predictors3b)  
y_hat
```

```
Out[86]: array([332948.26833857, 370106.11254277, 343567.46280956, ...]
```

```
287240.34413748, 441082.53186347, 285481.6045327 ])
```

```
In [87]: # Looking for homoskedasticity within our residuals (difference between predicted  
# and original values)  
  
fig, ax = plt.subplots(figsize = (20,12))  
  
resid = (target3b - y_hat)  
plt.scatter(x=range(y_hat.shape[0]), y=resid, alpha=0.1, color=c("indigo"))
```

```
Out[87]: <matplotlib.collections.PathCollection at 0x7f8c5d28b5e0>
```



Colinearity Check

```
In [88]: # Checking for any colinear features after our most recent updates to the model  
final_df_recent.corr()
```

	price	bedrooms	bathrooms	sqft_living	floors	sqft_basement	yr_built	grade_num
price	1.000000	0.297334	0.436667	0.668111	0.086822	0.180284	0.070998	0.732295
bedrooms	0.297334	1.000000	0.527929	0.637277	-0.037447	0.147375	-0.006995	0.287507
bathrooms	0.436667	0.527929	1.000000	0.601500	0.090224	0.341652	0.083625	0.398573
sqft_living	0.668111	0.637277	0.601500	1.000000	-0.071633	0.265944	-0.028823	0.658424
floors	0.086822	-0.037447	0.090224	-0.071633	1.000000	-0.188913	0.245547	0.080872
sqft_basement	0.180284	0.147375	0.341652	0.265944	-0.188913	1.000000	0.016764	0.128613
yr_built	0.070998	-0.006995	0.083625	-0.028823	0.245547	0.016764	1.000000	0.005902
grade_num	0.732295	0.287507	0.398573	0.658424	0.080872	0.128613	0.005902	1.000000

```
In [89]: # Create absolute values to understand if any features are above .75 r  
abs(final_df_recent.corr()) > 0.75
```

	price	bedrooms	bathrooms	sqft_living	floors	sqft_basement	yr_built	grade_num
price	True	False	False	False	False	False	False	False
bedrooms	False	True	False	False	False	False	False	False
bathrooms	False	False	True	False	False	False	False	False
sqft_living	False	False	False	True	False	False	False	False
floors	False	False	False	False	True	False	False	False
sqft_basement	False	False	False	False	False	True	False	False
yr_built	False	False	False	False	False	False	True	False
grade_num	False	False	False	False	False	False	False	True

```
In [90]: # save absolute value as a matrix, then sorting to understand which values are needed to drop  
df=final_df_recent.corr().abs().stack().reset_index().sort_values(0, ascending=False)
```

```
# zip the variable name columns
df['pairs'] = list(zip(df.level_0, df.level_1))

# set index to pairs
df.set_index(['pairs'], inplace = True)

# drop level columns
df.drop(columns=['level_1', 'level_0'], inplace = True)

# rename correlation column as cc rather than 0
df.columns = ['cc']

# drop duplicates
df.drop_duplicates(inplace=True)
```

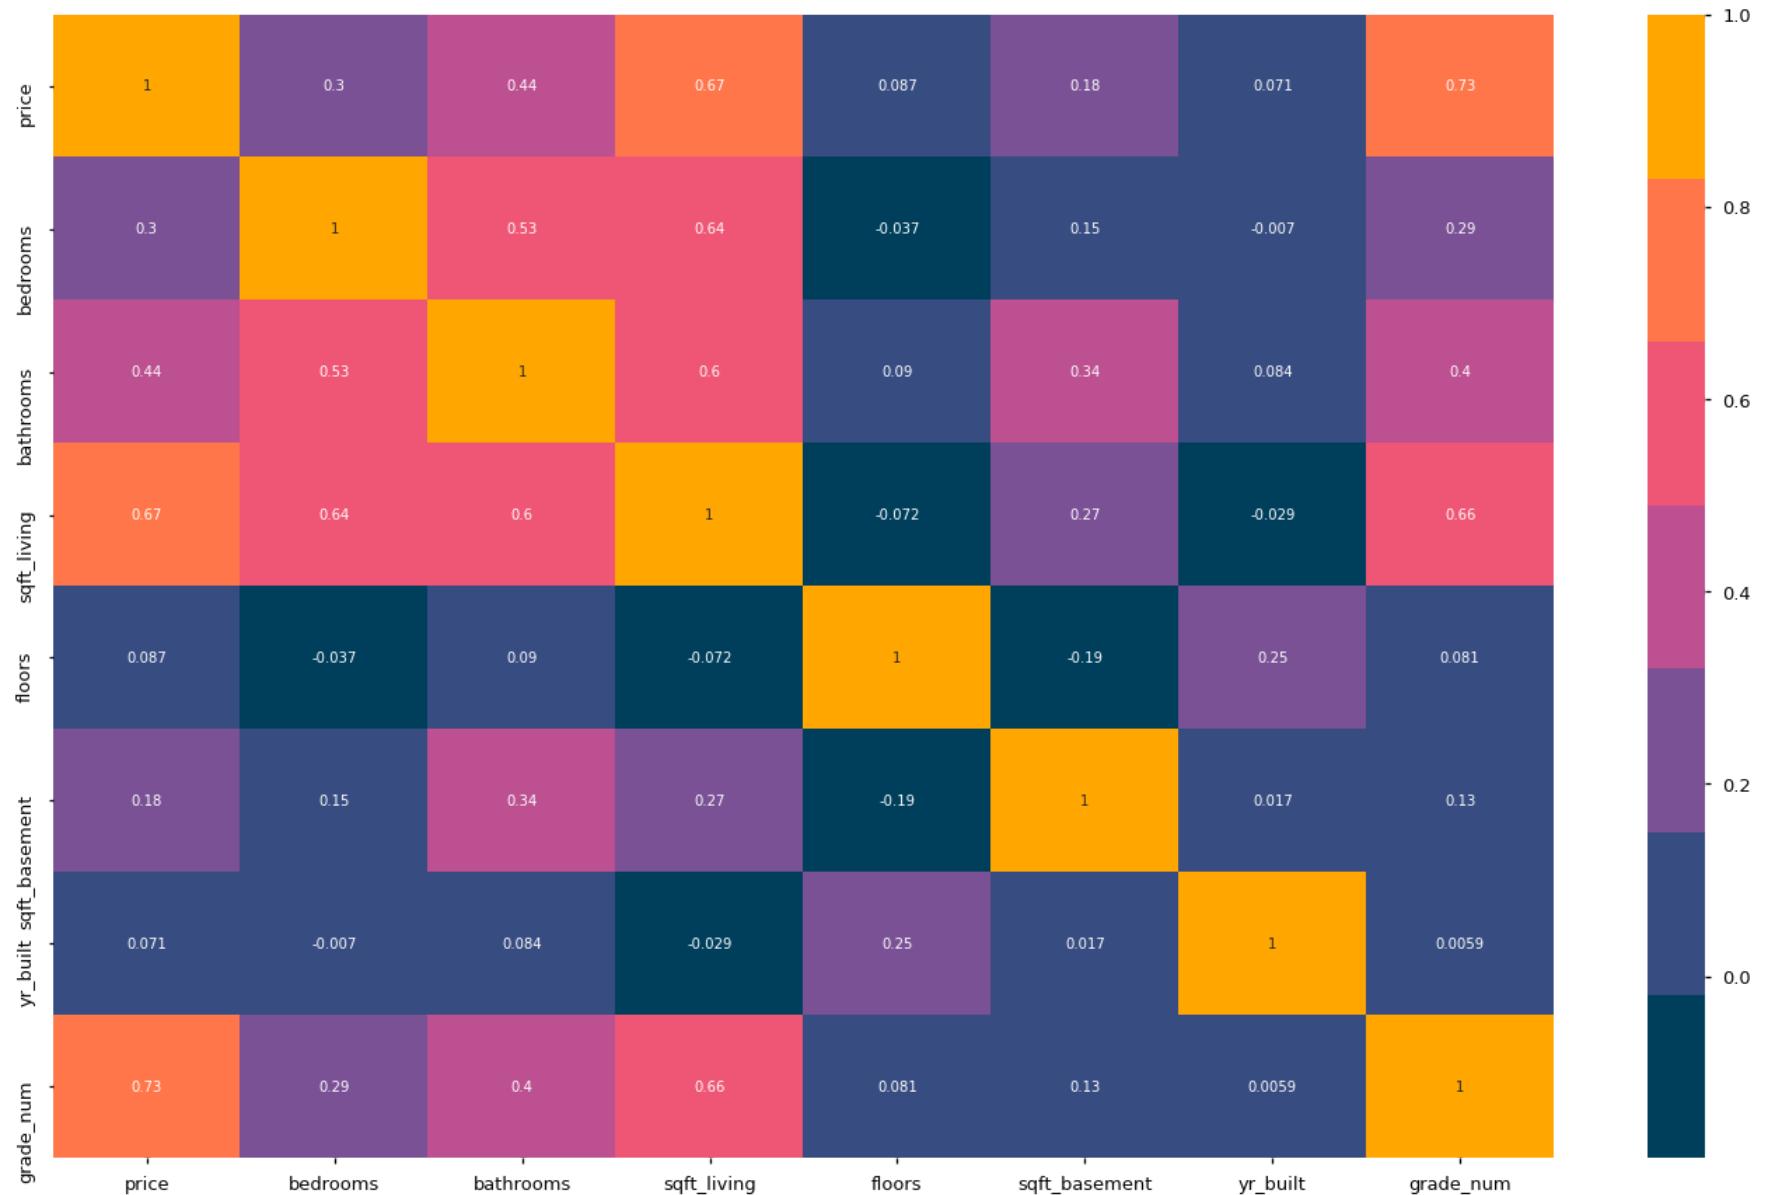
```
In [91]: df[(df.cc >.75) & (df.cc <1)]
```

```
Out[91]: cc
```

```
pairs
```

```
In [93]: plt.figure(figsize = (24,15))
corr = final_df_recent.corr()

sns.heatmap(corr, annot=True, cmap=pal)
plt.show()
```



That is a beautiful looking heatmap. Independent variables identified earlier are still holding strong.

Model3c

One last feature log transforming our price dependent variable to try and normalize one last time.

```
In [94]: # Log transforming the price column
```

```
final_df_recent['price_log'] = np.log(final_df_recent['price'])
final_df_recent = final_df_recent.drop(['price'], axis=1)
final_df_recent.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 6463 entries, 6 to 21596
Data columns (total 8 columns):
 #   Column           Non-Null Count  Dtype  
---  --  
 0   bedrooms        6463 non-null    int64  
 1   bathrooms       6463 non-null    float64 
 2   sqft_living     6463 non-null    int64  
 3   floors          6463 non-null    float64 
 4   sqft_basement   6463 non-null    float64 
 5   yr_built        6463 non-null    int64  
 6   grade_num       6463 non-null    int64  
 7   price_log       6463 non-null    float64 
dtypes: float64(4), int64(4)
memory usage: 774.4 KB
```

```
In [95]: # Model 3b - removed outliers, combined with features with p values > .05
```

```
# create target
target3c = final_df_recent['price_log']

# create predictors
predictors3c = final_df_recent.drop(['price_log'], axis=1)

# create model intercept
predictors_int3c = sm.add_constant(predictors3c)

# fit model to data
model3c = sm.OLS(target3c,predictors_int3c).fit()
```

```
In [96]: model3c.summary()
```

```
Out[96]: OLS Regression Results
Dep. Variable: price_log R-squared: 0.598
Model: OLS Adj. R-squared: 0.598
Method: Least Squares F-statistic: 1372.
Date: Thu, 27 Oct 2022 Prob (F-statistic): 0.00
```

Time: 21:00:39 **Log-Likelihood:** -203.56
No. Observations: 6463 **AIC:** 423.1
Df Residuals: 6455 **BIC:** 477.3
Df Model: 7
Covariance Type: nonrobust

	coef	std err	t	P> t 	[0.025	0.975]
const	2.9664	0.917	3.237	0.001	1.170	4.763
bedrooms	-0.0581	0.005	-10.794	0.000	-0.069	-0.048
bathrooms	0.0340	0.008	4.084	0.000	0.018	0.050
sqft_living	0.0002	6.89e-06	27.088	0.000	0.000	0.000
floors	0.0767	0.008	9.815	0.000	0.061	0.092
sqft_basement	3.902e-05	1.17e-05	3.321	0.001	1.6e-05	6.2e-05
yr_built	0.0040	0.000	8.803	0.000	0.003	0.005
grade_num	0.1880	0.004	42.910	0.000	0.179	0.197

Omnibus: 7.973 **Durbin-Watson:** 1.990
Prob(Omnibus): 0.019 **Jarque-Bera (JB):** 7.748
Skew: -0.065 **Prob(JB):** 0.0208
Kurtosis: 2.891 **Cond. No.** 9.24e+05

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 9.24e+05. This might indicate that there are strong multicollinearity or other numerical problems.

Predictions

Going back to our final model 3b to understand mean errors, and any train/ test information. This is not as important for the analysis

that we completed. However, I am executing here in the event there are questions around capability and model performance.

```
In [97]: # Looking at mean absolute error
```

```
metrics.mean_absolute_error(target3b, y_hat)
```

```
Out[97]: 99629.81033569487
```

```
In [98]: # Looking at mean squared error
```

```
metrics.mean_squared_error(target3b, y_hat, squared=False)
```

```
Out[98]: 127365.31991408017
```

```
In [99]: # Creating a test/ train split of 80/20 for both our predictors and dependent features
```

```
x_train, x_test, y_train, y_test = train_test_split(predictors3b, target3b, test_size=0.20)
print(len(x_train), len(x_test), len(y_train), len(y_test))
```

```
5170 1293 5170 1293
```

```
In [100...]: # Fitting our train/ test splits
```

```
lr.fit(x_train, y_train)
```

```
# Create price predictions on train and test data from the independant variables
y_hat_train = lr.predict(x_train)
y_hat_test = lr.predict(x_test)
```

```
In [101...]: # Create train & test residuals
```

```
train_residuals = y_hat_train - y_train
test_residuals = y_hat_test - y_test
```

```
In [102...]: # Compute MSE for train and test set
```

```
mse_train = np.sum((y_train-y_hat_train)**2)/len(y_train)
mse_test = np.sum((y_test-y_hat_test)**2)/len(y_test)
```

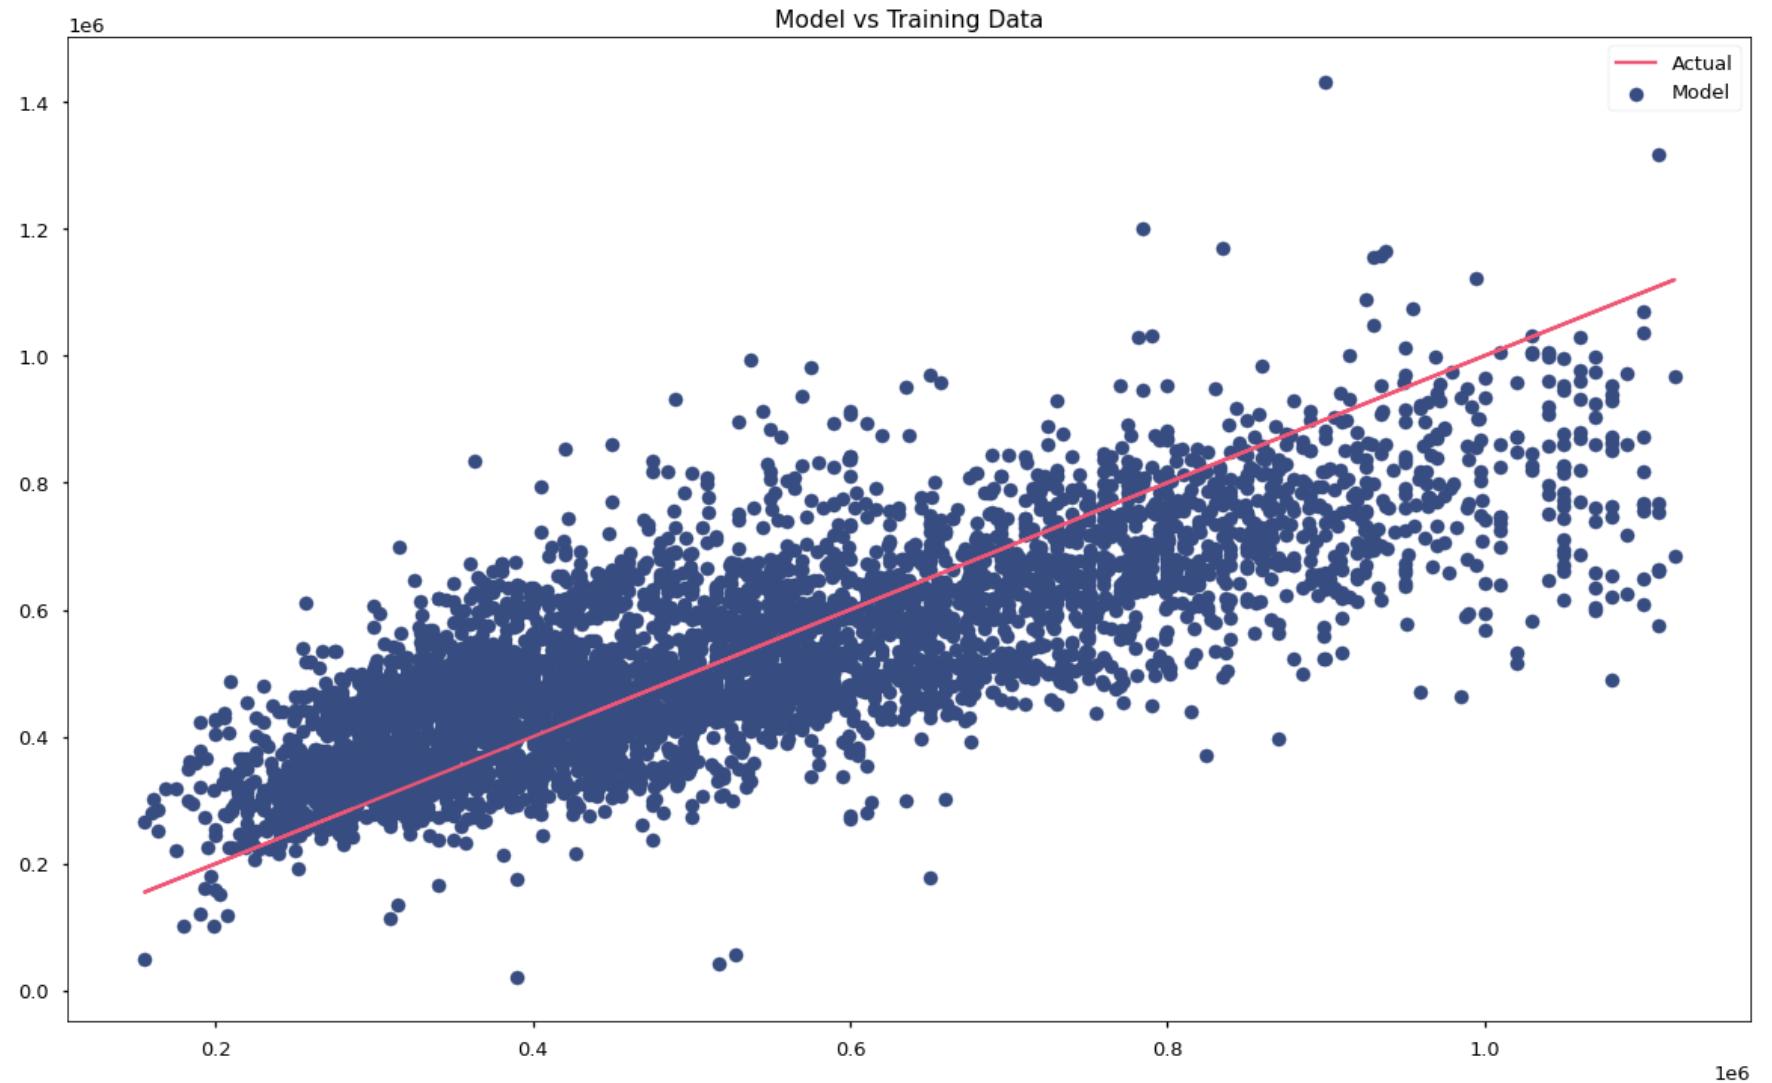
```
print('Train Mean Squared Error: ', mse_train)
print('Test Mean Squared Error: ', mse_test)
```

```
Train Mean Squared Error: 16264858714.254068
```

```
Test Mean Squared Error: 16082375583.891954
```

```
In [103...]: # Create a matplotlib figure for train vs actual  
  
plt.figure(figsize=(20, 12))  
  
plt.scatter(y_train, y_hat_train, label='Model', color=c("indigo"))  
  
plt.plot(y_train, y_train, label='Actual', color=c("peach"))  
plt.title('Model vs Training Data')  
plt.legend()
```

```
Out[103...]: <matplotlib.legend.Legend at 0x7f8c5e7481c0>
```



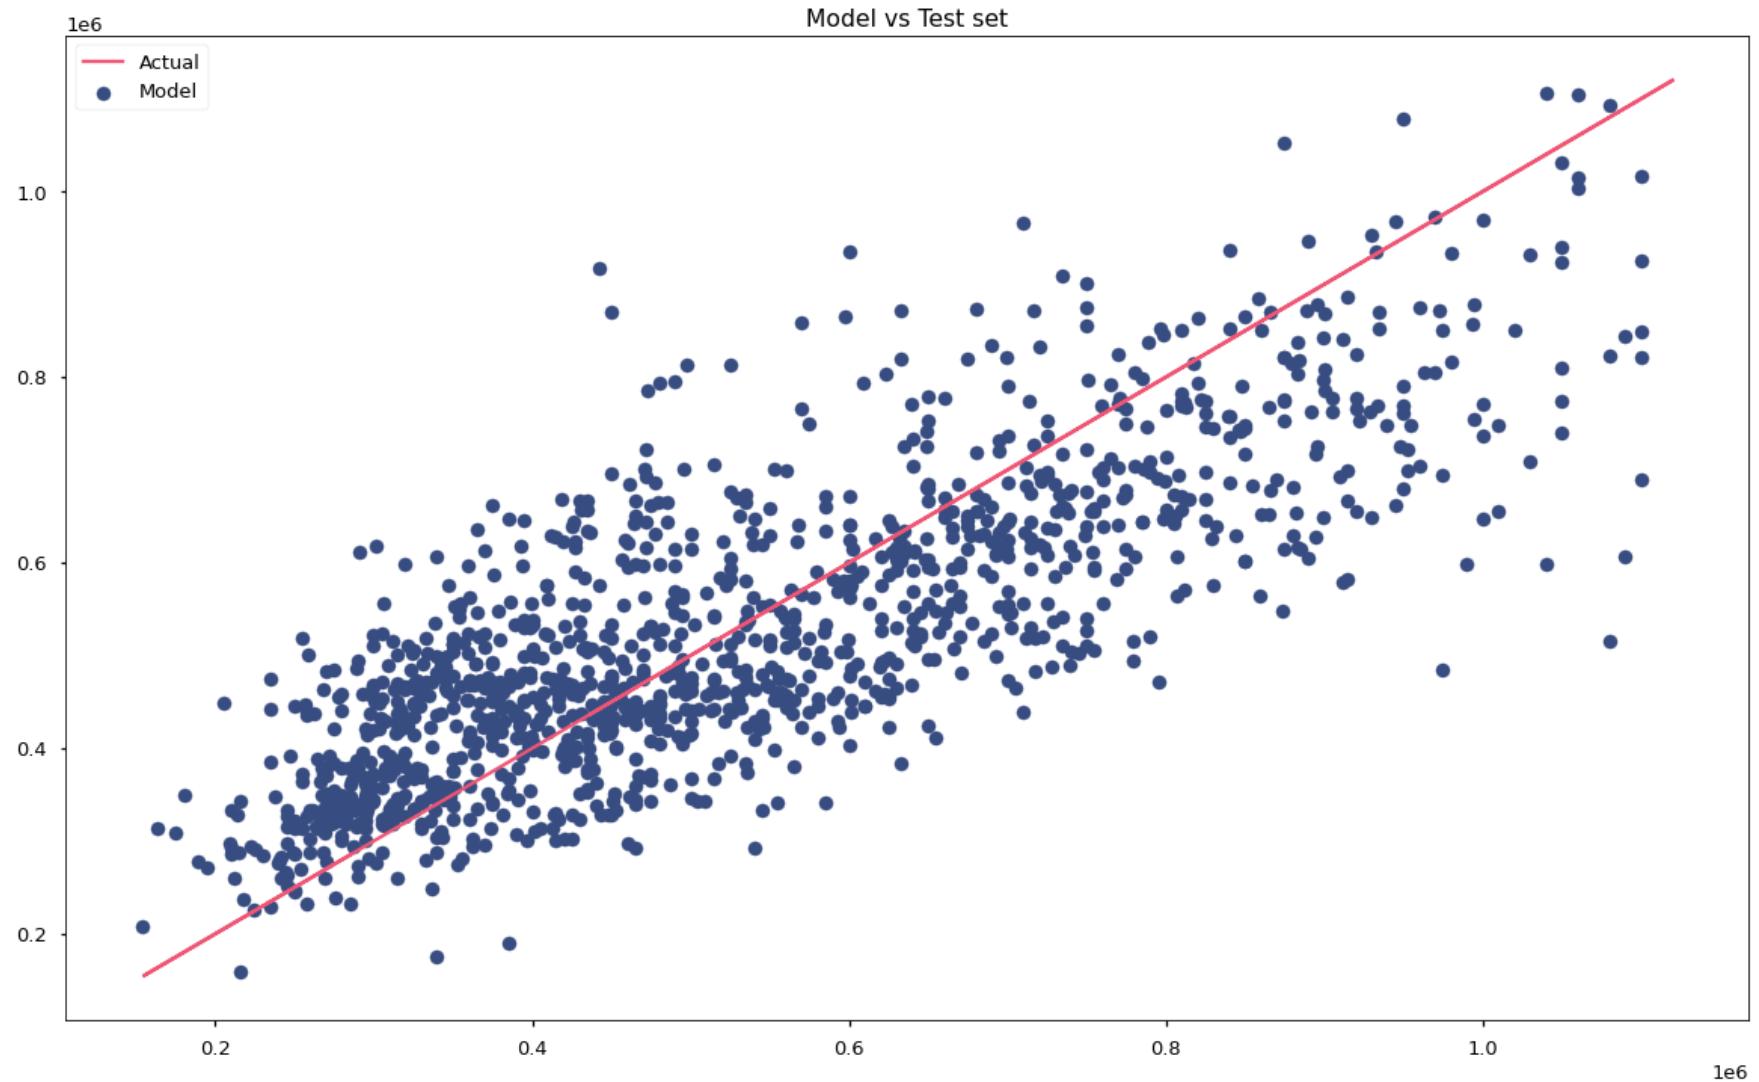
```
In [104]: # Create a matplotlib figure for test vs actual
```

```
plt.figure(figsize=(20,12))

plt.scatter(y_test, y_hat_test, label='Model', color=c("indigo"))

plt.plot(y_train, y_train, label='Actual', color=c("peach"))
plt.title('Model vs Test set')
plt.legend()
```

Out[104... <matplotlib.legend.Legend at 0x7f8c42a22160>



Additional Models

3b - Cohort 1 Regression (< 1932)

In [105... # Looking at values prior to 1932. Trying to understand the data sets as we have defined
the splits for the final dataset.

```
final_df_old = final_df.loc[final_df['yr_built'] < 1932]
final_df_old.describe()
```

Out[105...]

	price	bedrooms	bathrooms	sqft_living	floors	sqft_basement	yr_built	is_renovated	grade_num
count	2.611000e+03	2611.000000	2611.000000	2611.000000	2611.000000	2611.000000	2611.000000	2611.000000	2611.000000
mean	5.290214e+05	3.091153	1.622654	1650.523171	1.370931	298.385676	1917.659900	0.100345	6.930678
std	2.119122e+05	1.064276	0.668541	632.095906	0.378039	380.835432	8.818302	0.300517	0.862690
min	8.000000e+04	1.000000	0.500000	370.000000	1.000000	0.000000	1900.000000	0.000000	3.000000
25%	3.765000e+05	2.000000	1.000000	1160.000000	1.000000	0.000000	1910.000000	0.000000	6.000000
50%	5.162500e+05	3.000000	1.500000	1560.000000	1.500000	0.000000	1919.000000	0.000000	7.000000
75%	6.675000e+05	4.000000	2.000000	2060.000000	1.500000	620.000000	1925.000000	0.000000	7.000000
max	1.120000e+06	11.000000	4.500000	4530.000000	3.000000	1570.000000	1931.000000	1.000000	10.000000

In [106...]

```
# This is the "Old Dataset/ cohort" that has the final independent variables
# as 3b.

# create target
target5old = final_df_old['price']

# create predictors
predictors5old = final_df_old.drop(['price'], axis=1)

# create model intercept
predictors_int5old = sm.add_constant(predictors5old)

# fit model to data
model5old = sm.OLS(target5old,predictors_int5old).fit()
```

In [107...]

```
# Summary of our regression

model5old.summary()
```

Out[107...]

OLS Regression Results

Dep. Variable:	price	R-squared:	0.567
Model:	OLS	Adj. R-squared:	0.566
Method:	Least Squares	F-statistic:	426.2

Date: Thu, 27 Oct 2022 **Prob (F-statistic):** 0.00
Time: 21:00:50 **Log-Likelihood:** -34632.
No. Observations: 2611 **AIC:** 6.928e+04
Df Residuals: 2602 **BIC:** 6.934e+04
Df Model: 8
Covariance Type: nonrobust

	coef	std err	t	P> t 	[0.025	0.975]
const	1.369e+06	6.04e+05	2.265	0.024	1.84e+05	2.55e+06
bedrooms	-7399.1077	3430.514	-2.157	0.031	-1.41e+04	-672.294
bathrooms	2.348e+04	6076.634	3.864	0.000	1.16e+04	3.54e+04
sqft_living	85.0675	9.848	8.638	0.000	65.758	104.377
floors	2.115e+04	1.06e+04	1.989	0.047	294.960	4.2e+04
sqft_basement	12.9589	11.299	1.147	0.252	-9.197	35.115
yr_built	-995.2438	315.632	-3.153	0.002	-1614.159	-376.329
is_renovated	-4.064e+04	9402.776	-4.322	0.000	-5.91e+04	-2.22e+04
grade_num	1.276e+05	4164.805	30.630	0.000	1.19e+05	1.36e+05

Omnibus: 39.458 **Durbin-Watson:** 2.033
Prob(Omnibus): 0.000 **Jarque-Bera (JB):** 62.254
Skew: 0.135 **Prob(JB):** 3.03e-14
Kurtosis: 3.707 **Cond. No.** 5.72e+05

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 5.72e+05. This might indicate that there are strong multicollinearity or other numerical problems.

3b - Cohort 2 Regression (1932-1990)

In [108...]

```
# Looking at values from 1932 to 1990. This is my second of the 3 cohorts to understand  
# what data is really hurting my model.  
  
final_df_middle = final_df.loc[(final_df['yr_built'] >= 1932) & (final_df['yr_built'] < 1990)]  
final_df_middle.describe()
```

Out[108...]

	price	bedrooms	bathrooms	sqft_living	floors	sqft_basement	yr_built	is_renovated	grade_
count	1.136500e+04	11365.000000	11365.000000	11365.000000	11365.000000	11365.000000	11365.000000	11365.000000	11365.000000
mean	4.407531e+05	3.315882	1.858293	1852.862824	1.205851	332.383722	1964.802640	0.031676	7.297
std	2.009916e+05	0.872957	0.652777	703.656975	0.389401	445.275806	14.663887	0.175144	0.909
min	7.800000e+04	1.000000	0.500000	380.000000	1.000000	0.000000	1932.000000	0.000000	4.000
25%	2.850000e+05	3.000000	1.500000	1330.000000	1.000000	0.000000	1953.000000	0.000000	7.000
50%	4.000000e+05	3.000000	1.750000	1764.000000	1.000000	0.000000	1965.000000	0.000000	7.000
75%	5.525000e+05	4.000000	2.250000	2250.000000	1.000000	650.000000	1978.000000	0.000000	8.000
max	1.120000e+06	10.000000	6.750000	7480.000000	3.000000	2720.000000	1989.000000	1.000000	11.000

In [109...]

```
# model 3b - removed outliers, combined with features with p values > .05  
  
# create target  
target5mid = final_df_middle['price']  
  
# create predictors  
predictors5mid = final_df_middle.drop(['price'], axis=1)  
  
# create model intercept  
predictors_int5mid = sm.add_constant(predictors5mid)  
  
# fit model to data  
model5mid = sm.OLS(target5mid,predictors_int5mid).fit()
```

In [110...]

```
model5mid.summary()
```

Out[110...]

OLS Regression Results

Dep. Variable:	price	R-squared:	0.519
Model:	OLS	Adj. R-squared:	0.519

Method: Least Squares **F-statistic:** 1534.
Date: Thu, 27 Oct 2022 **Prob (F-statistic):** 0.00
Time: 21:00:52 **Log-Likelihood:** -1.5074e+05
No. Observations: 11365 **AIC:** 3.015e+05
Df Residuals: 11356 **BIC:** 3.016e+05
Df Model: 8
Covariance Type: nonrobust

	coef	std err	t	P> t	[0.025]	0.975]
const	6.995e+06	2.09e+05	33.516	0.000	6.59e+06	7.4e+06
bedrooms	-1.522e+04	1912.443	-7.961	0.000	-1.9e+04	-1.15e+04
bathrooms	4.041e+04	3267.397	12.369	0.000	3.4e+04	4.68e+04
sqft_living	85.4102	4.053	21.075	0.000	77.466	93.354
floors	2.745e+04	4400.961	6.237	0.000	1.88e+04	3.61e+04
sqft_basement	18.1956	4.396	4.139	0.000	9.579	26.812
yr_built	-3827.1125	108.805	-35.174	0.000	-4040.389	-3613.836
is_renovated	1.527e+04	7715.681	1.979	0.048	145.563	3.04e+04
grade_num	1.018e+05	2167.400	46.967	0.000	9.75e+04	1.06e+05

Omnibus: 717.407 **Durbin-Watson:** 1.987
Prob(Omnibus): 0.000 **Jarque-Bera (JB):** 1066.119
Skew: 0.536 **Prob(JB):** 3.13e-232
Kurtosis: 4.050 **Cond. No.** 4.42e+05

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 4.42e+05. This might indicate that there are strong multicollinearity or other numerical problems.

Conclusion

In conclusion, we were able to evaluate and identify a few different options for our customer and ultimately some recommendations that they can make to the clients about increasing the value of their home.

- **The median home price in King County, WA is 475k dollars**
- **By increasing the living square footage of your home, you can increase the value by 99 dollars per square foot.**
- **By adding bathrooms, and adding 1/4 bath features, you can increase the value of your home by 22 dollars for each 1/4 bath you add.**

There are a few additional analyses that we could do moving forward to identify additional features, and those are:

- **Address limitations/ concerns - add additional pricing data, adjust for inflation, etc.**
- **Look at additional zip codes to understand values more specifically (or not).**
- **Evaluate additional variables/ features such as environmental, or other factors that may have an impact on a home's value.**
- **Refresh the analysis regularly with new data to understand how the market is evolving over time.**